

# Databricks- CI/CD & DevOps for Data Engineering

Shivam Panicker,  
Senior Specialist Solutions Architect, APJ  
Databricks Inc.

# Agenda

01

About CI &  
CD

02

Why do we  
need it?

03

About  
DevOps

04

DE with  
Databricks

05

Notebook  
based  
CICD

06

IDE Based  
CICD

07

Lab

08

Summary

09

Appendix

# Objectives

- At the end of the course, you will be familiar with the concepts of CI/CD & DevOps.
- You will have the understanding of how Databricks can be leveraged to build scalable solutions.
- Will be able to apply knowledge hands-on in developing CI/CD & DevOps solutions with Databricks.

# Definitions

- CI- Continuous Integration is the practice of merging all developers' working copies to a shared mainline several times a day that triggers automated build with testing.
- CD- Continuous Delivery is an approach in which teams produce software in short cycles, ensuring that the software can be reliably released at any time and, following a pipeline through a "production-like environment", without doing so manually.
- DevOps- is a set of practices that combines software development (Dev) and IT operations (Ops). It aims to shorten the systems development life cycle and provide continuous delivery with high software quality.

# About CI & CD

**Continuous Integration** is the process of setting up **automated practices write from the scratch to develop code with the focus on testing**. It intends to automate test from the lowest level of individual code development (unit testing) to integration of various modules (integration testing) to the final product (system testing)

**Continuous Delivery** is the process of **automating & building error free artefacts and deploying them** into different environments relevant to the stage of the development life cycle. It uses the modern generation of tools which simplifies code management, running different cycles of testing.

# Why do we need it?

## **Continuous Integration**

- Improved quality of code.
- Efficient bug tracking and fixes.
- Streamlined and expedited review process.
- Effective communication, collaboration and feedback loops.

## **Continuous Delivery**

- Measurable progress.
- Scalable code deployed into production.
- Reduced manual overheads and thus, lesser human error prone deployments.

# Overview of a typical Databricks CI/CD pipeline

**Continuous  
integration**

**Continuous  
delivery**

**Code**



**Build**



**Release**




**Deploy**



**Test**



**Operate**




# About DevOps

**Devops** is the process of combining the practices of software engineering and IT operations together to deliver a tangible solution in a faster iteration of time. It enables different personas to work together and automate the software development life cycle with the use of the right tools that can enable Continuous Integration & Continuous development.



# DevOps Lifecycle



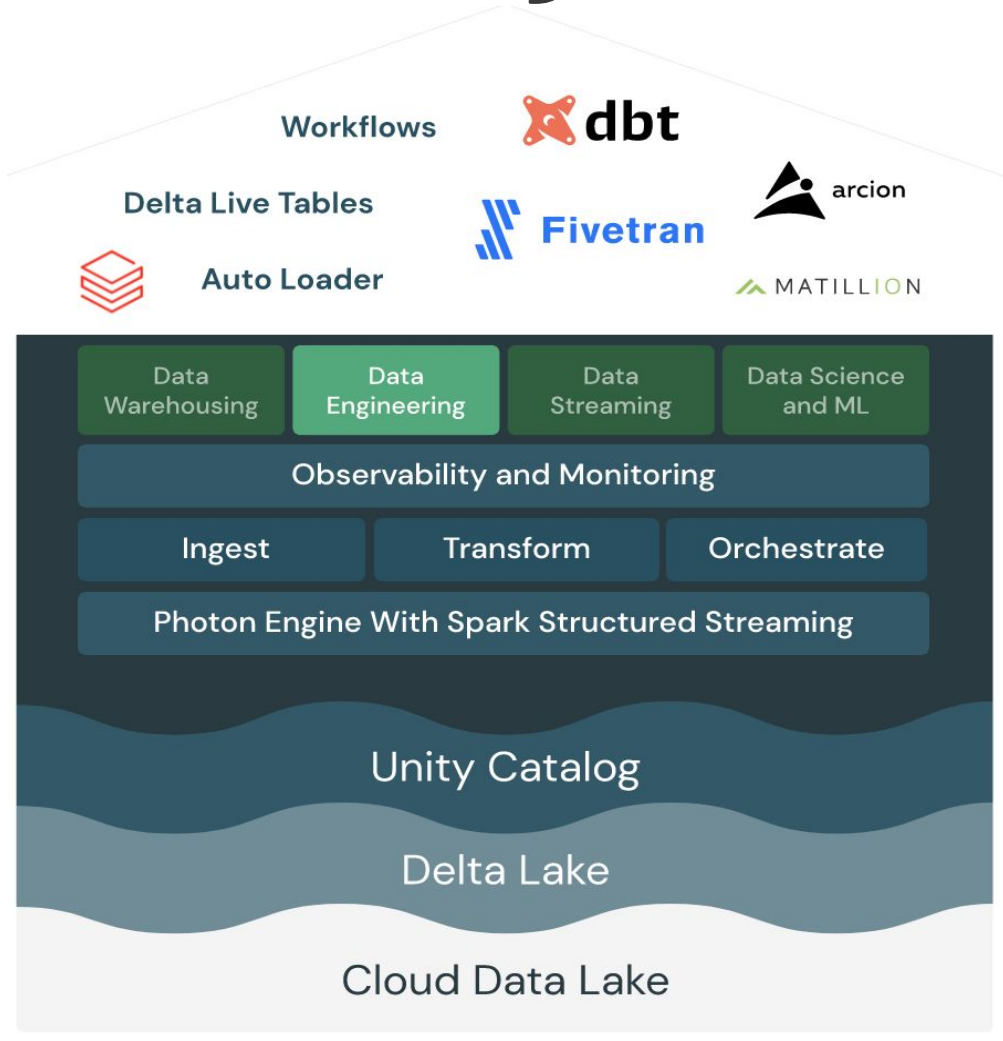
# Data Engineering with Databricks

**Databricks** allows customers to engineer the data pipelines in various ways. Databricks powers the data engineers to develop data pipelines using sql, python, scala and R.

It allows following options to develop the pipelines:

- Notebook driven development,
- IDE based development, and
- Databricks SQL

# Data Engineering with Databricks



## Simple

Unify your data warehousing and AI use cases on a single platform

## Multicloud


One consistent data platform across clouds

## Open

Built on open source and open standards

# Workloads on Databricks

- Data orchestration through Databricks Workflows
- Delta Live Tables manage your full data pipelines
- Simplifies data engineering with a curated data lake approach through Delta Lake



# Thrives within your modern data stack

## BI and Dashboards

Power BI +tableau Looker

MicroStrategy ThoughtSpot Qlik

## Machine Learning

MathWorks Labelbox John Snow LABS

Azure Machine Learning H2O.ai Amazon SageMaker

## Data Science

PyCharm Jupyter

HEX R Studio

## Data Governance

Collibra MMUTA PRIVACERA

Quest Alation AZURE PURVIEW

## Data Pipelines

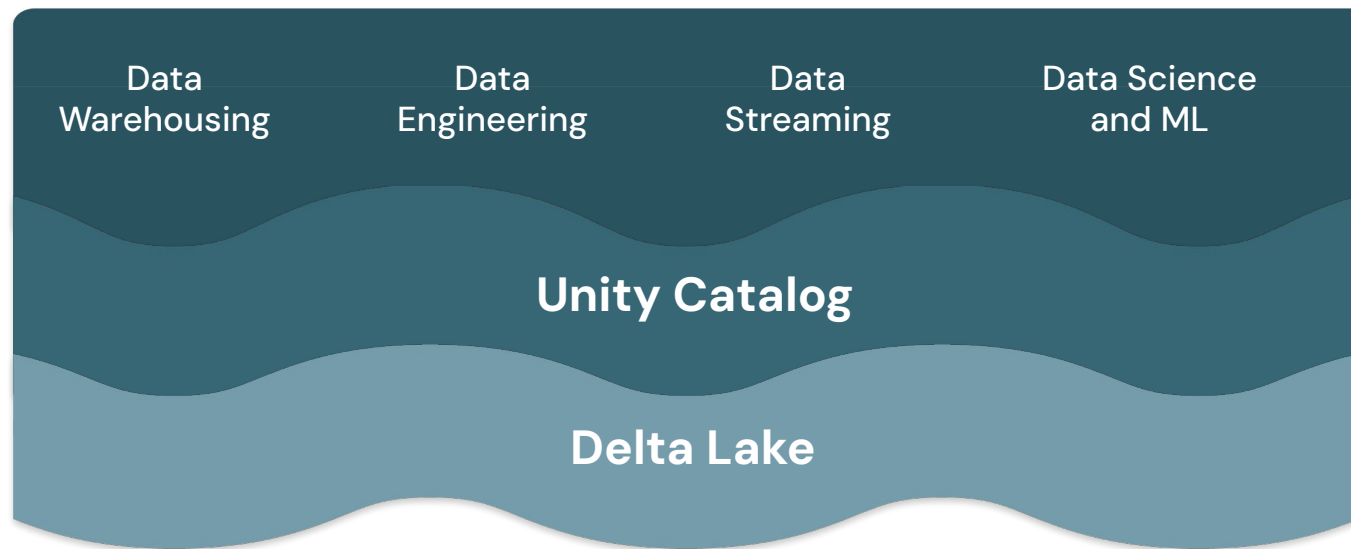
dbt Labs MATILLION Azure Data Factory

Informatica Prophecy

## Data Ingestion

Fivetran arcion CONFLUENT

Rivery Airbyte Qlik



Cloud Data Lake

Microsoft Azure

aws

Google Cloud

## Consulting & SI Partners

accenture

avanade

Capgemini


Cognizant

Deloitte

slalom

# SQL workloads on Databricks

- Great performance and concurrency for BI and SQL workloads on Delta Lake
- Native SQL interface for analysts
- Support for BI tools to directly query your most recent data in Delta Lake



# Development workflows

- Notebooks only:
  - Faster experimentation / feedback
  - Harder to automate, challenges with testing, chaining jobs
- Code only:
  - Develop Python / Scala code in IDE, build, execute as jobs
  - Better automation, more tooling, delayed feedback loop
- Mixed (typical usage by our customers):
  - Code developed in IDE, packaged & deployed as libraries
  - Notebooks are used for configuration & orchestration of calls to libraries

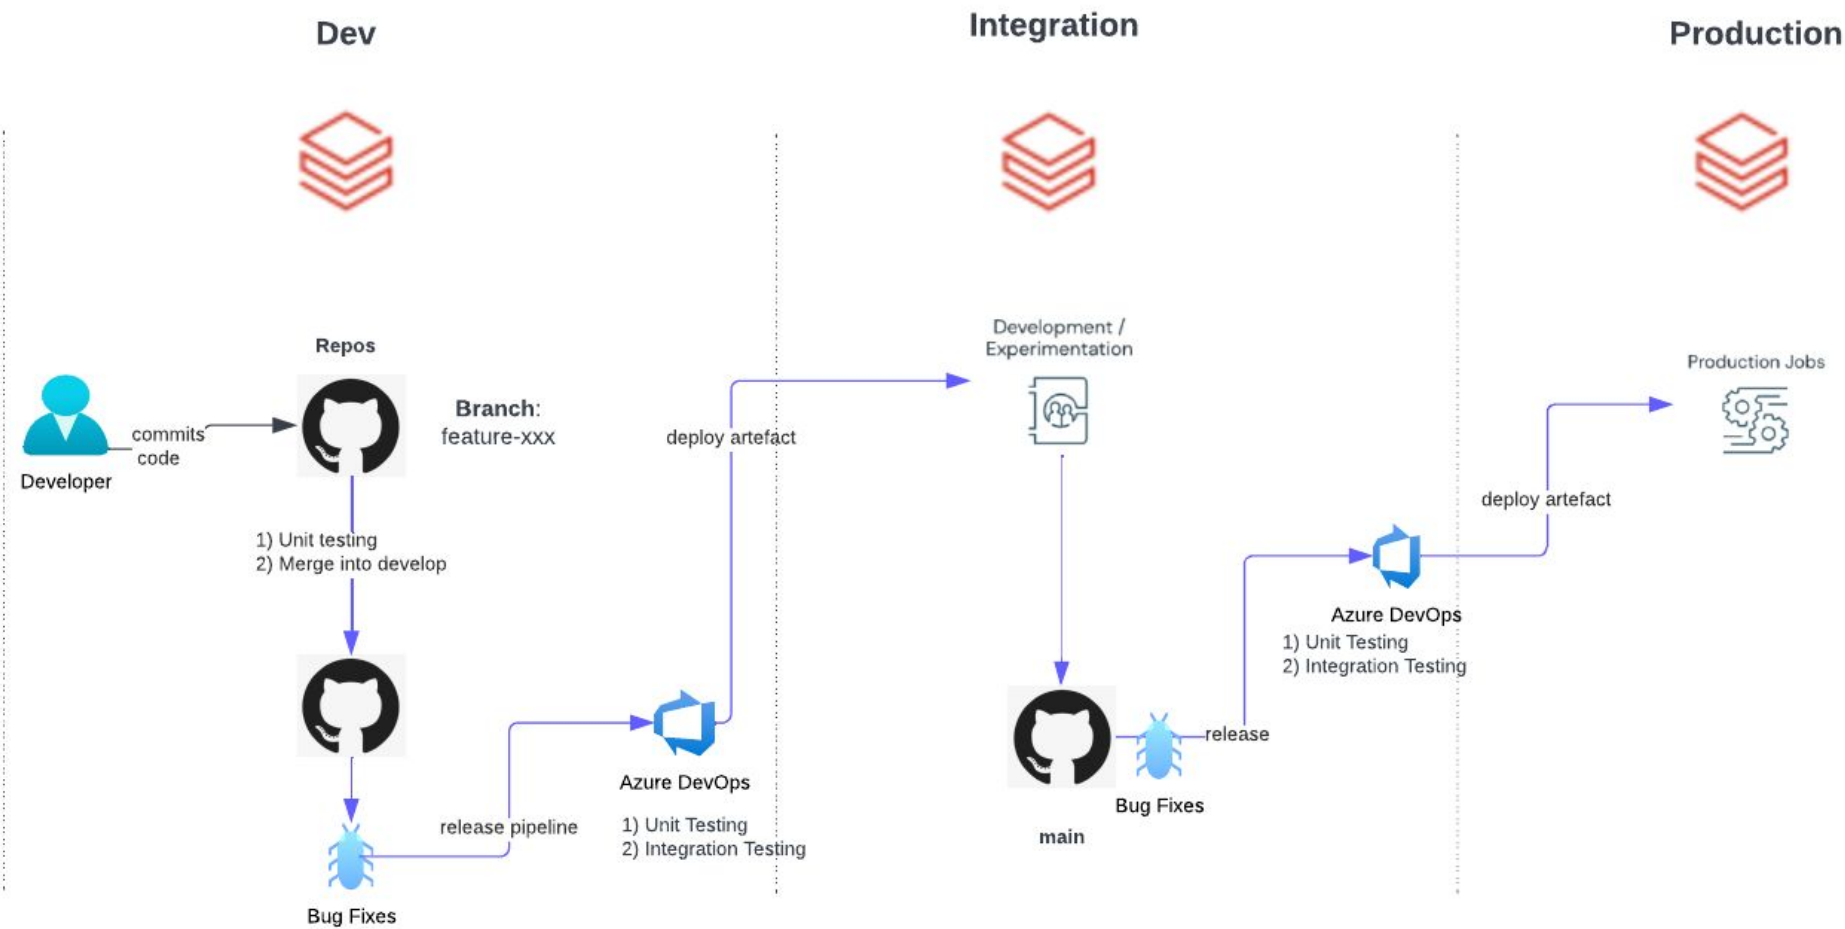
# CI/CD for Databricks using Notebooks



# General CI/CD workflow on Databricks

- Develop the code using one of the workflows (notebooks/IDE/mixed)
- Pull notebooks & commit them (dbx/databricks-cli/stack-cli)
- Build assets via build pipeline & perform unit testing (Github Actions, Azure Devops, Jenkins, etc)
- Push notebooks to staging (databricks-cli/stack-cli)
- Push assets to staging (databricks-cli, terraform)
- Run tests (unit/integration/e2e) on staging (dbx, databricks-cli, schedulers)
- If successful, push notebooks & assets to production (databricks-cli/stack-cli)
- Reconfigure jobs / cluster to use new notebooks/assets, if necessary (databricks-cli, terraform)

# Notebooks based CI/CD- Developer Workflow



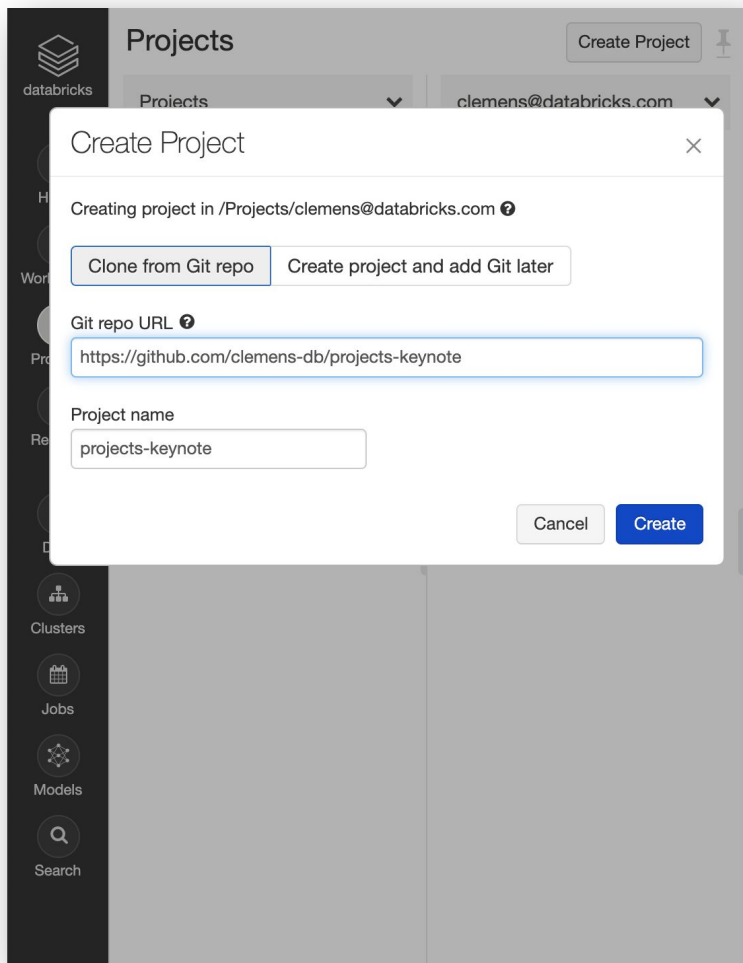
# Writing the testable code

- Split code into testable chunks
  - Individual functions: DataFrame+Configuration in, DataFrame out
  - Business logic – functions, calling other functions
  - No dependency onto the global state/external systems
- Unit tests
  - Test only one specific function, on small amount of data
  - Fast, able to run locally
- Integration tests
  - Test business logic on limited amount of data
  - Usually run in a separate environment, as part of CI implementation
- Acceptance/End-to-End tests
  - Run on data, close to production (volume/velocity/...)
  - Run in environment, close to production

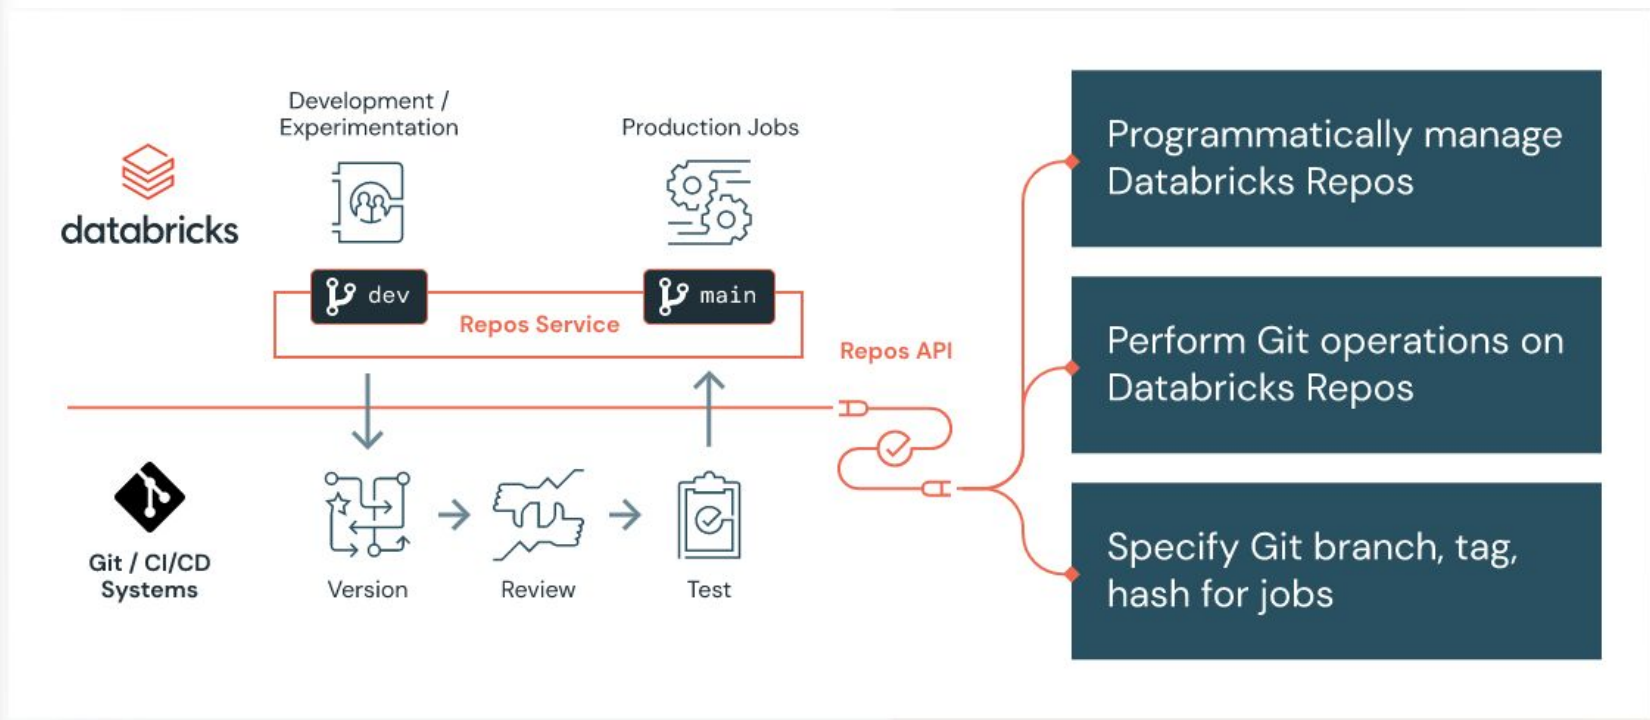
# Testing libraries for Spark

- Built-in Spark test suite
  - Designed to test all parts of Spark
  - Supports RDD, Dataframe/Dataset, Streaming APIs
- spark-testing-base:
  - Scala & Python support
  - Supports RDD, Dataframe/Dataset, Streaming APIs
- spark-fast-tests – Scala, Spark 2 & 3
- chispa – Python version of spark-fast-tests
- pytest-spark – Python, native integration with pytest
  
- Code samples for all libraries in one place

# Git-based Repos in Databricks



## CI/CD Integration

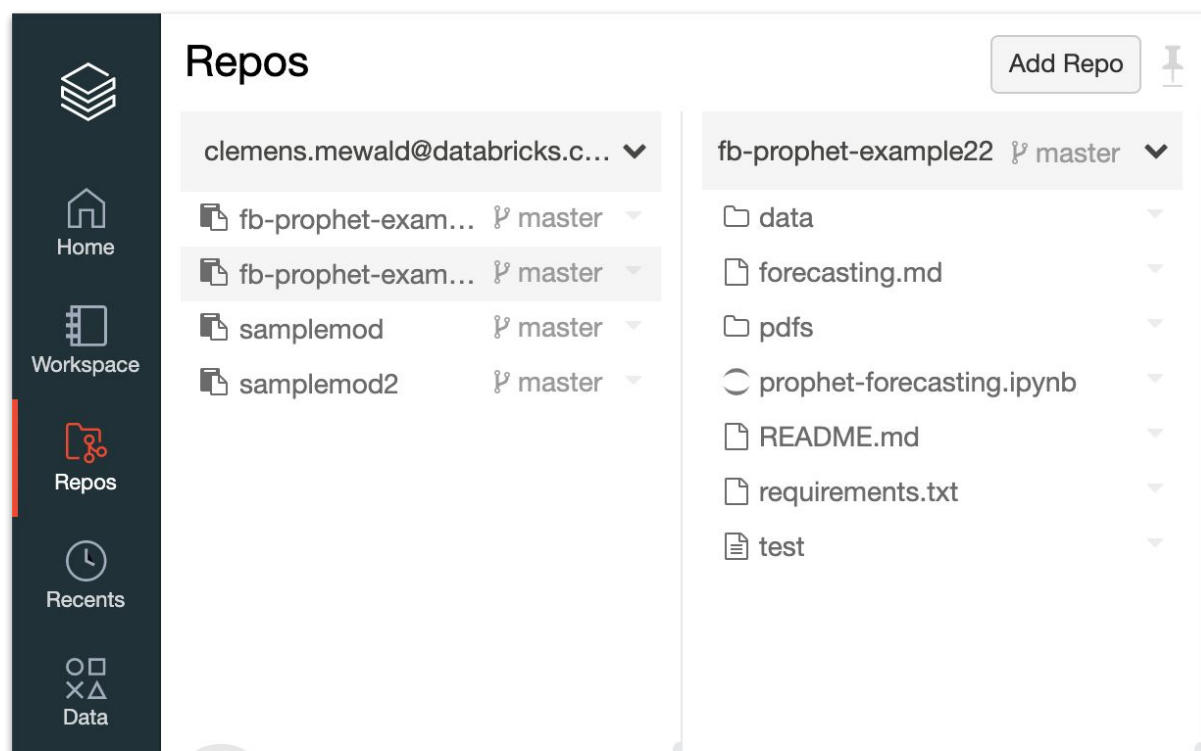


## Supported Git Providers



# Arbitrary files support in Repos

Seamlessly bring any type of workload to Databricks



- Portability of code
- Library files – use Python/R files as packages



- Environment specification portability
- Build packages from the same repo



- Small data ease of use
- Relative imports



- Whatever you can do with files “just works”

# Using arbitrary files

Package with library functions

Automatically import changes

Import library functions

It's possible to build packages from the same source

File Explorer:

- dlt-best-practices main
- .gitignore
- conftest.py
- dlt\_package
- my\_package
- notebooks
- pipelines
- pytest.ini
- README.md
- requirements.txt
- scripts
- setup.py
- tests
- unit-requirements.txt

my\_package

- \_\_init\_\_.py
- \_\_version\_\_.py
- code1.py
- functions.py
- code2.py

Cmd 2

```
1 %load_ext autoreload
2 %autoreload 2
```

Command took 0.20 seconds -- by alexey.ott@databricks.com at 27/01/2022, 17:46:05 on NutterDemo

Cmd 3

```
1 from my_package.code1 import * # instead of %run ./Code1
2 from my_package.code2 import * # instead of %run ./Code2
```

Command took 1.14 seconds -- by alexey.ott@databricks.com at 27/01/2022, 17:46:17 on NutterDemo

# Testing notebooks: using %run

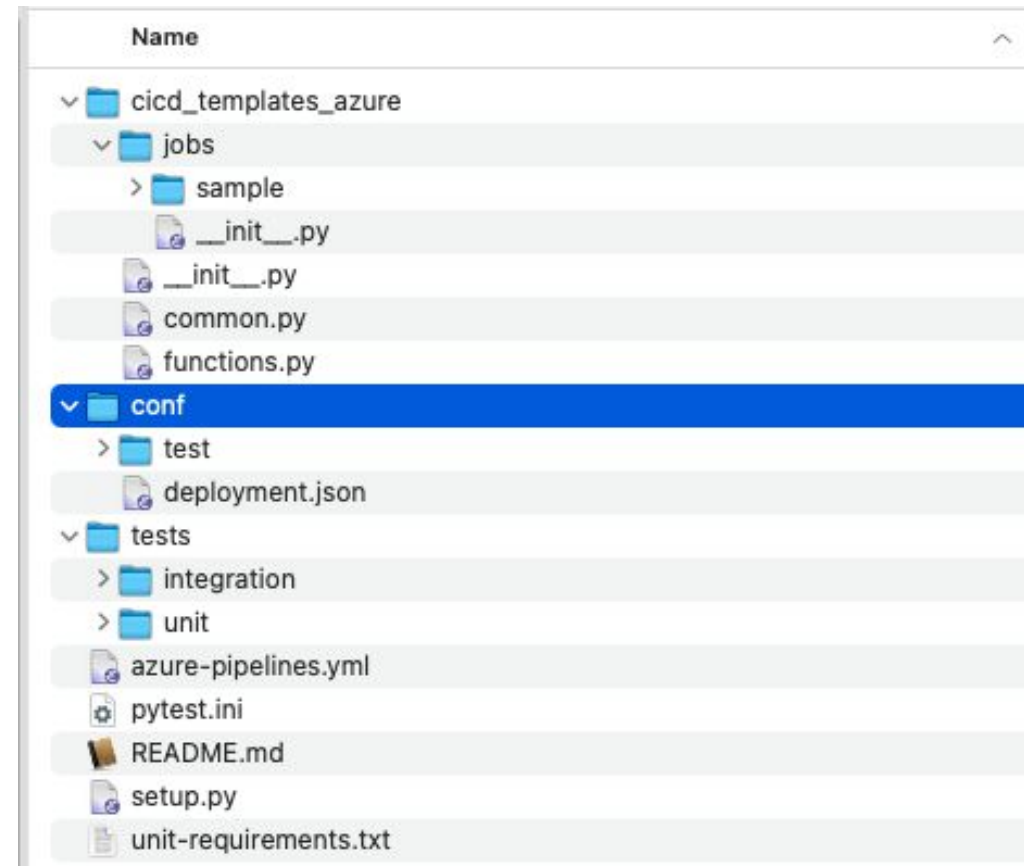
- Workflow:
  - Create separate notebook with functions that you want to test
  - Create separate notebook for tests
  - Import defined functions from notebook **%run**
  - Depending on the language, defined test functions or test classes
  - Execute tests
- This workflow may require special approach when using test frameworks – they often rely on the annotations and auto-discovery, so may not work with notebooks:
  - Solution – create test suites explicitly & execute them



# CI/CD for Databricks using IDE

# Code organization

- Organize code into testable chunks:
  - Library functions
  - Jobs – doing actual data processing
- Unit tests for library functions
  - Use specialized testing frameworks
- Integration tests for jobs
- Use dbx init to generate a skeleton of Python project



# Development flow

- Make changes using the IDE of choice
- Run unit tests locally (see [examples](#)), debug errors
- Commit the code into a repository
- CI/CD process pickups changes & performs:
  - Run unit tests
  - Run integration tests on Databricks
  - Publish test results
  - Promote changes to other environments if necessary
- Use dbx & other tooling to implement CI/CD pipelines

# CI/CD integrations for Databricks

# Existing integrations


- Azure DevOps
  - Use dbx init to generate pre-configured project template
  - [Continuous integration and delivery on Azure Databricks using Azure DevOps](#)
  - [Implement CI/CD with Azure DevOps](#) (MS Learning)
  - Additional DevOps integrations via [DevOps for Azure Databricks by Microsoft DevLabs](#) or [3rd party package by Data Thirst](#)
- Jenkins
  - [Continuous integration and delivery on Azure Databricks using Jenkins](#)
- Github Actions
  - Use dbx init to generate pre-configured project template

# DBX

# dbx: extended support for IDE development

**dbx** is an extension of the Databricks CLI that makes it **simple** to:

- Build, run, and test your code on Databricks from your local IDE
- Quickly iterate on your project while retaining the efficiency of using an IDE for modularity and testing
- Nimblely manage multiple execution environments and deployment configurations
- Get access to the full feature set in Databricks Runtime via holistic, batch execution (e.g., Photon, Unity Catalog, Feature Store, AutoML)



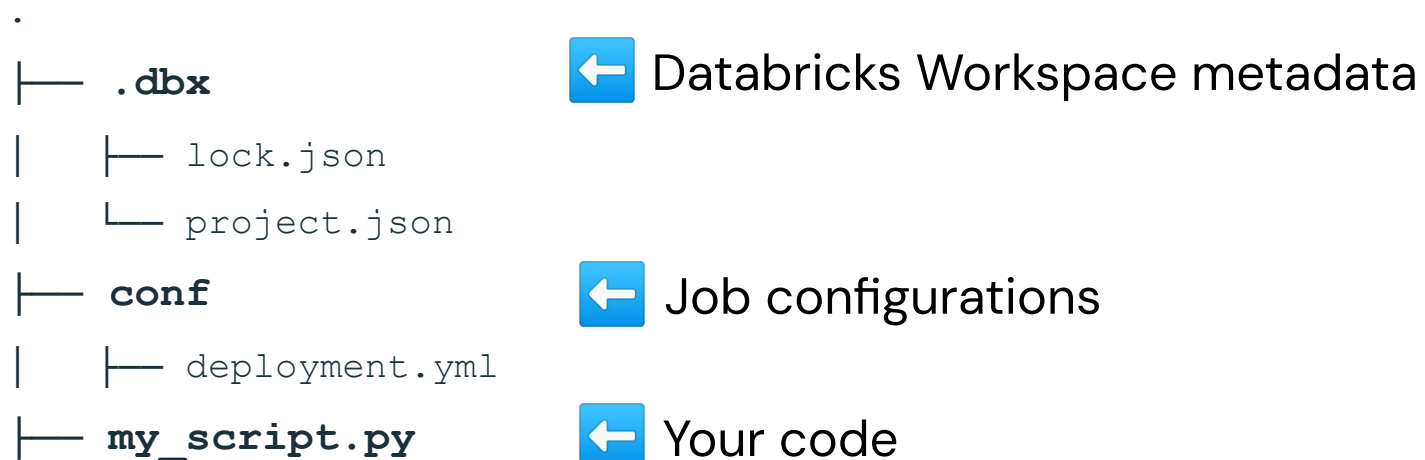
- Build, run and test:
  - › **dbx configure** Configure your project to run against Databricks
  - › **dbx deploy** Deploy a Databricks Workflow as scripts or a versioned build artifact
  - › **dbx launch** Trigger the Workflow locally and check the results on Databricks
- Quickly iterate
  - › **dbx execute** Submit code directly on All-Purpose Compute
  - › **dbx sync** Replicate local changes to Databricks Repos
- Bootstrap your project
  - › **dbx init** Initialize a template for a Python package with tests, CI/CD, and more



- Supports deploying single or multi-task workflows
  - Define tasks and execution environments in flexible YAML or JSON
    - Databricks workspace(s) to run in
    - Compute resources and library dependencies
- Ship code to different environments from the terminal
  - Develop against dev:
    - › **dbx execute --environment=dev model-train**
  - Test in staging:
    - › **dbx deploy --environment=staging model-tests**
    - › **dbx launch --environment=staging model-tests**
  - Deploy to production:
    - › **dbx deploy --environment=prod model-train**
    - › **dbx launch --environment=prod model-train**

**dbx**

# Minimal Project



```
# run on All-Purpose Compute
> dbx execute --cluster-name='Shared Autoscaling' --no-rebuild --no-package my-job
```

# Databricks CLI

# Databricks CLI

- Installation:
  - `pip install databricks-cli`
- Configuration
  - Requires personal access token (PAT)
  - Execute: `databricks configure` <- it will generate config file on the disk
  - Or use environment variables to dynamically change configuration (handy for CI/CD pipelines, etc.)
- It's possible to have multiple profiles in the configuration file & select them via command-line
  - `databricks --profile <profile-name> command ...`

# Working with files on DBFS

- databricks fs <subcommand>
  - ls, mkdirs, cp, mv, rm, cat
- We can use it to publish artifacts – jars, wheels, eggs, ...
  - `databricks fs cp --overwrite target/scala-2.11/some.jar 'dbfs:/FileStore/jars/'`

# Working with clusters & jobs

- databricks clusters <subcommand>
  - create, delete, edit, list, get, resize, start, restart, ...
- databricks jobs <subcommand>
  - list, get, create, delete, reset, run-now
- databricks runs <subcommand>
  - list, get, submit, get-output
- databricks libraries <subcommand>
  - install, uninstall, list, ...
- When in doubt, try databricks clusters --help

# Library management

# Artifacts

Libraries & other artifacts built by build server are pushed:

- to DBFS – dedicated location, versioned.
  - pros: simple to implement
  - cons: harder to maintain, when multiple workspaces are used
- to ADLS, S3 or GCS (depends on the DBR version, also need to have cloud storage authentication configured correctly)
- to Artifact store – Azure DevOps, Artifactory, ...
  - pros: easy access from multiple workspaces, integrated with CI/CD systems
  - cons: need external system, may not always work well with authenticated services, ...



# Installing from private repositories

- Python, R:
  - Change global settings using [cluster or global init scripts](#)
- Private Maven repositories:
  - Before DBR 11.0 – no, because libraries were resolved in the control plane. Need to put to DBFS/ADLS, or use init script to install libraries
  - DBR 11+ – yes, libraries are now resolved locally on the cluster (still may need to use init scripts to update config files)

```
#!/bin/bash
```

```
cat << 'EOF' > /etc/pip.conf  
[global]  
timeout = 60  
index-url =  
https://<host>/<path>/simple  
trusted-host = <host>  
EOF
```

# Recommendations for building Scala/Java artifacts

- Generate uberjar (or fat jar) only with your code and internal dependencies (like, that are only available in the private Maven repositories)
- Don't include into uberjar following dependencies as will conflict with Databricks Runtime versions (mark them as provided):
  - Core Spark dependencies (spark-core, spark-sql, ...) and spark-sql-kafka
  - Delta Lake
- Open source dependencies it's better not include into uberjar, but attach to clusters/declare in the jobs definitions (mark them as provided):
  - This decrease the size of the fat jar
  - It's easier to change them in case of new versions released, especially to fix vulnerabilities or other critical issues

# Attaching artifacts to clusters/jobs

- Prefer to use IaC (Infrastructure-as-Code) – Databricks Terraform provider:
  - pros: versioned, easy to rollback not working changes
  - pros: changes are made via pull requests, tested by build pipelines, published by release pipelines to multiple environments
  - pros: could be completely automated on release of artifacts
  - cons: need to have/learn one more tool
- Alternatives – scripts using Databricks CLI or REST API
  - cons: harder to maintain

# Jobs Scheduling

# Jobs Scheduling & orchestration

- Databricks Workflows: Built-in job scheduling ([doc](#)):
  - Periodic scheduling of the jobs (cron-like right now)
  - Execute notebook / jar / Python script / Spark-submit / DLT / DBSQL / DBT
- Apache Airflow
  - [Airflow provider for Databricks](#)
  - Execute notebook / jar / Python script / DLT pipeline as a job
  - Execute SQL queries against Databricks cluster / SQL Endpoint
- Azure Data Factory
  - Execute notebook / jar / Python script
  - Triggers: time-based, ADF Event-based

# Databricks Workflows

- Scheduled or executed one time
- Permanent (created via UI/REST API/Terraform) or ephemeral (submitted via REST API / Azure Data Factory / Airflow)
- Can run on ephemeral (cheaper) or existing clusters (more expensive)
- Multiple tasks in a job
- Jobs cluster reuse for permanent jobs!
- We can pass parameters to a job (use `dbutils.widgets.get()` to get them in notebook. In “normal” programs – they are just command-line parameters)

# Task types in Databricks Workflows

We can execute:

- Databricks notebooks
- Jar files
- Python code – source file & wheels
- Arbitrary code supported by spark-submit (for example, R code)
- DBSQL Dashboards/Queries/Alerts
- Delta Live Tables pipelines
- DBT

Type \*

Notebook ▼

Notebook ✓

Python script

Python wheel

SQL **New**

Delta Live Tables pipeline

dbt

JAR

Spark Submit

Task name \* ?

Test\_task

Type \*

Notebook

Source \* ?

Workspace

Path \* ?

/Users/alexey.ott@databricks.com/Test

Cluster \* ?

Shared\_job\_cluster 126 GB · 36 Cores · DBR 10.4 LTS · Spark 3.2.1 · Scala 2.12

Parameters ?

+ Add

Advanced options

Add dependent libraries

Edit notifications

Edit retry policy

Edit timeout



Job details

Job ID 130431168113409

Creator alexey.ott@databricks.com

Run as alexey.ott@databricks.com

Tags + Tag



Git

Not configured

Add Git settings



Schedule

At 09:00 AM (UTC+00:00 — UTC)

Edit schedule Pause Delete



Compute

Shared\_job\_cluster

Driver: Standard\_DS3\_v2 · Workers: Standard\_DS3\_v2 · 8 workers · 10.4 LTS (includes Apache Spark 3.2.1, Scala 2.12)

Configure Swap



Notifications

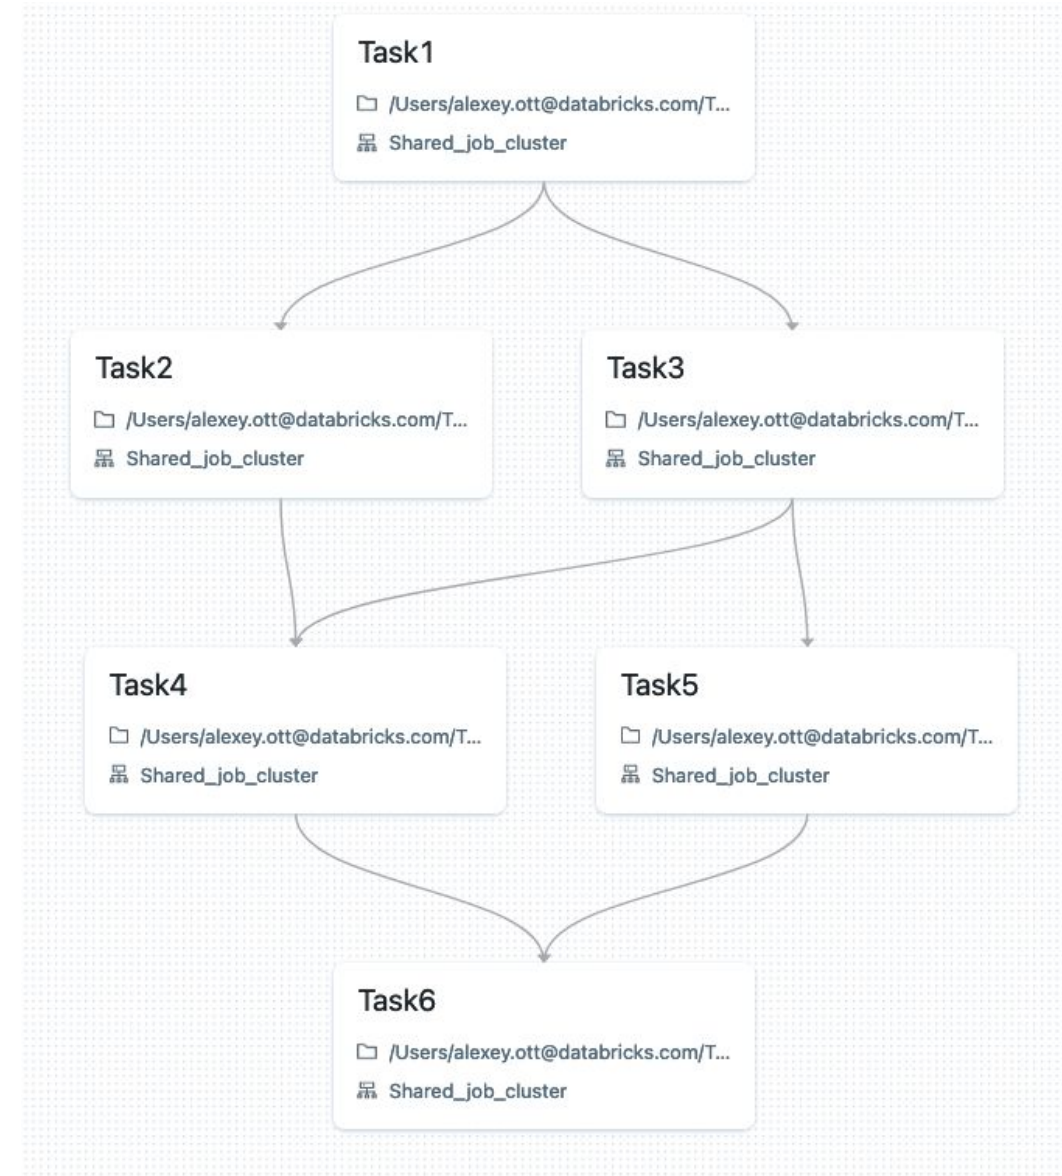
No notifications

Edit notifications



# Multiple tasks in Databricks Workflows

- It's possible to specify multiple tasks in a single job
- Linear, or non-linear dependencies
- Job cluster(s) reuse – faster start times for next task(s)
- It's possible to pass values between tasks
- It's possible to rerun failed tasks (repair)
- Fully supported for automation via REST API / Terraform / Databricks CLI



# Databricks Job Scheduling

- Create a job through the UI

Schedule ?

Every  at  :  (UTC+00:00) ...

- Create a job using the API (using cron
- Advanced job options
  - Alerts
  - Max concurrent runs
  - Timeout
  - Retries
  - Permissions

Schedule ?

(UTC+00:00) ...

☒ Show Cron Syntax

# Summary


- We have covered the basic concepts of Continuous Integration (CI) & Continuous Development (CD) which helps in building reusable patterns of delivering code into production.
- Covered the basic elements of Data Engineering using Databricks and how notebooks and/or IDEs can be used for efficient coding.
- Use services such as Azure DevOps, Github Actions, etc to manage CI/CD and promote code into different environments.
- Databricks tooling such as DBX, Databricks-CLI to manage communication between Databricks APIs and other services such as Azure DevOps, Github Actions.
- Databricks Orchestration using native Workflows other external services such as Airflow, Azure Data Factory, etc.

# Labs

- As part of the lab today, we will be working with retail dataset
- Dataset consists of sales, stores, customer and product info
- The lab will focus on using Notebooks and perform unit testing on individual batch transformations. Further, we will run some integration tests to ensure nothing breaks upon running black box tests from bronze to silver to gold layers.
- Post the unit and integration tests, we shall create an artefact and deploy into production.

# Setup

- Integrating github account with Databricks Repos
  - Github --> Settings --> Developer Settings --> Personal access tokens
    - Generate new token (classic) --> Enter password --> Enter note
    - Select repo and workflow scope --> Submit --> Copy token
  - Login to Databricks --> User Settings --> Git Integration
    - Git provider (Github) --> Git provider username or email (your git credentials) --> Token



## User Settings

Password Access tokens **Git integration** Notebook settings Email preferences Language settings

### With co-versioned repo

Databricks Repos allows you to clone a remote Git repo, which you can specify when you add a repo. [Learn more](#)

### With individual notebooks

Although we recommended using co-versioned repo for Git integration, Databricks supports notebook version control integration

### Set your Git provider and credentials

You can also set your Git provider credentials via API. [Learn More](#)

Select your Git provider from the dropdown.

Git provider

GitHub

Git provider username or email ⓘ

shivam.panicker@gmail.com

To generate a GitHub personal access token, follow the [GitHub documentation](#). The token must have the "repo" and "workflow" (i

Token


Token with repo read/write permissions

Save

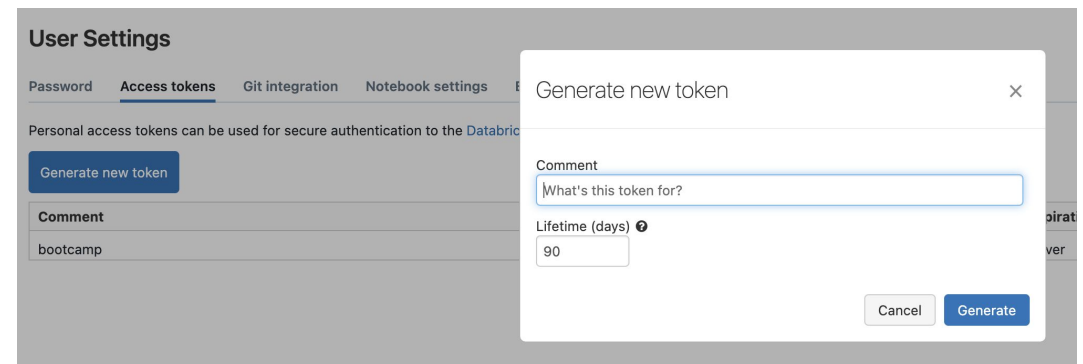

# Setup

- Fork the repo– [https://github.com/shivampanicker/cicd\\_with\\_databricks.git](https://github.com/shivampanicker/cicd_with_databricks.git)

- Create 2 branches
  - develop – on: main
  - feature/<username>: on develop



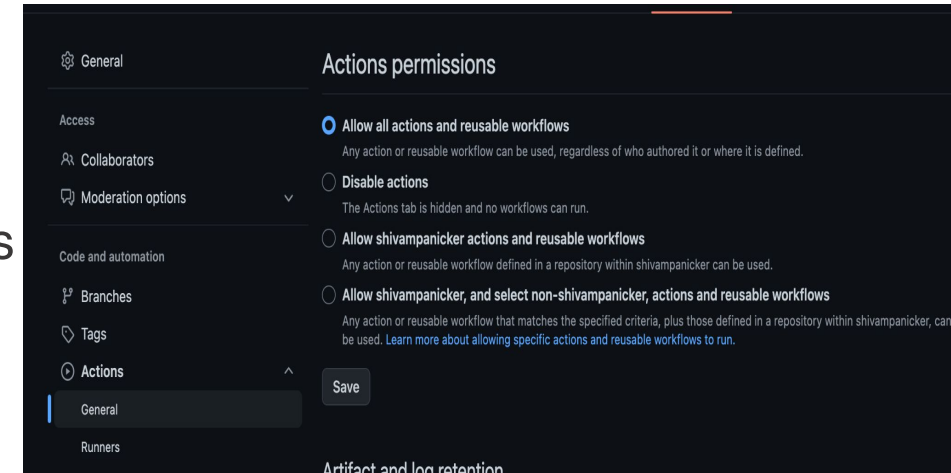
- Generate a personal access token and from Databricks Workspace and copy it in your notepad.



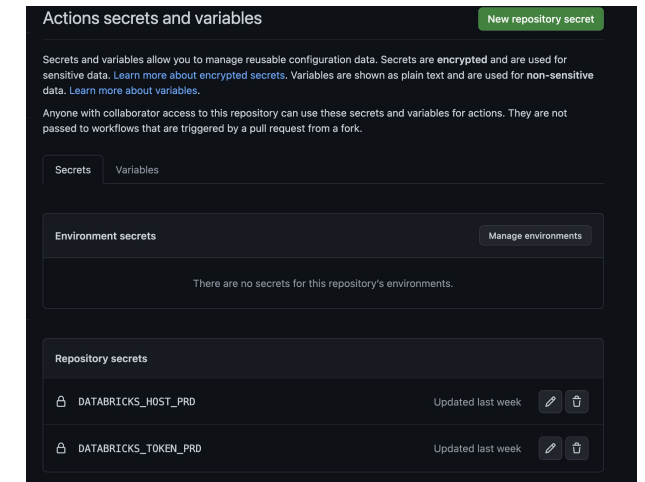


# Setup

- Configure github actions (CI/CD)
  - Go to the forked repository.
  - Settings --> Actions
  - General --> Allow all Actions and reusable workflows



- Configure secrets
  - Go to Settings --> Secrets and Variables
  - Actions- Add new repository secret.
    - DATABRICKS\_HOST\_PRD https://<databricks\_workspace\_name>.com/
    - DATABRICKS\_TOKEN\_PRD <personal access token generated in DB>



# Action time!

- Go to Databricks Repos and clone your repository.
- Checkout branch- *feature/<username>*
- There are few code level changes required before one raises pull request.
  - Navigate to this file: *src/main/tests/bronze/test\_load\_data\_into\_bronze*
  - Set `expected_num_files = 2`
  - Check the Silver layer unit tests in the *src/main/tests/silver/* folder and add an assertion. Feel free to write any test case.
- Review the `integration_suite` test which uses Files in Repos feature and fill in the missing elements.
  - `src/main/python/gold/gold_layer_etl.py`
  - `src/main/tests/integration_suite/test_integration_gold_layer_etl`
- Review the dbx deployment file
  - Open `cicd_with_databricks/deployment/deploy-job.yaml` and update `notebook_path` variable to your Databricks Repos location.
- Review files under `.github/workflow/` to understand the CI/CD plan.

# Action time!

- Go to Databricks Repos and clone your repository.
- Checkout branch- `feature/<username>`.
- There are few code level changes required before one raises pull request:
  - Navigate to this file: `src/main/tests/bronze/test_load_data_into_bronze`
  - Set `expected_num_files = 2`
  - Check the Silver layer unit tests in the `src/main/tests/silver/` folder and add an assertion. Feel free to write any test case.

# Action time!

- Review the `integration_suite` test which uses Files in Repos feature and fill in the missing elements:
  - `src/main/python/gold/gold_layer_etl.py`
  - `src/main/tests/integration_suite/test_integration_gold_layer_etl`
- Review the dbx deployment file
  - Open `deployment/deploy-job.yaml`
  - Update `notebook_path` variable to your Databricks Repos location.
- Review files under `.github/workflow/` to understand the CI/CD plan.

# Action time!

- Create a pull request\* and view the CI/CD unit testing job that spins up in github → actions.
- Once it succeeds, merge the pull request into develop branch\* and view the CI/CD integration testing job that spins up
- Once integration tests are completed on develop branch, raise a PR from develop branch into main\*. View the CI/CD job that spins up, runs unit & integration tests.
- Once it succeeds, merge the pull request into main\* and view the CI/CD job that creates Databricks workflow jobs and launches them.

**\*Make sure that you are requesting to merge to the forked repository**

Thank you

# Appendix

- [https://en.wikipedia.org/wiki/Continuous\\_integration](https://en.wikipedia.org/wiki/Continuous_integration)
- [https://en.wikipedia.org/wiki/Continuous\\_delivery](https://en.wikipedia.org/wiki/Continuous_delivery)
- DevOps- <Definition>