



aws INNOVATE

DATA EDITION

AUGUST 30, 2023

Accelerate insights from real-time streaming data at scale with databases, analytics and ML

Arun Balaji
Principal Prototyping Engineer
AWS India

Gopalakrishnan Subramanian
Principal Database Specialist Solutions Architect
AWS India

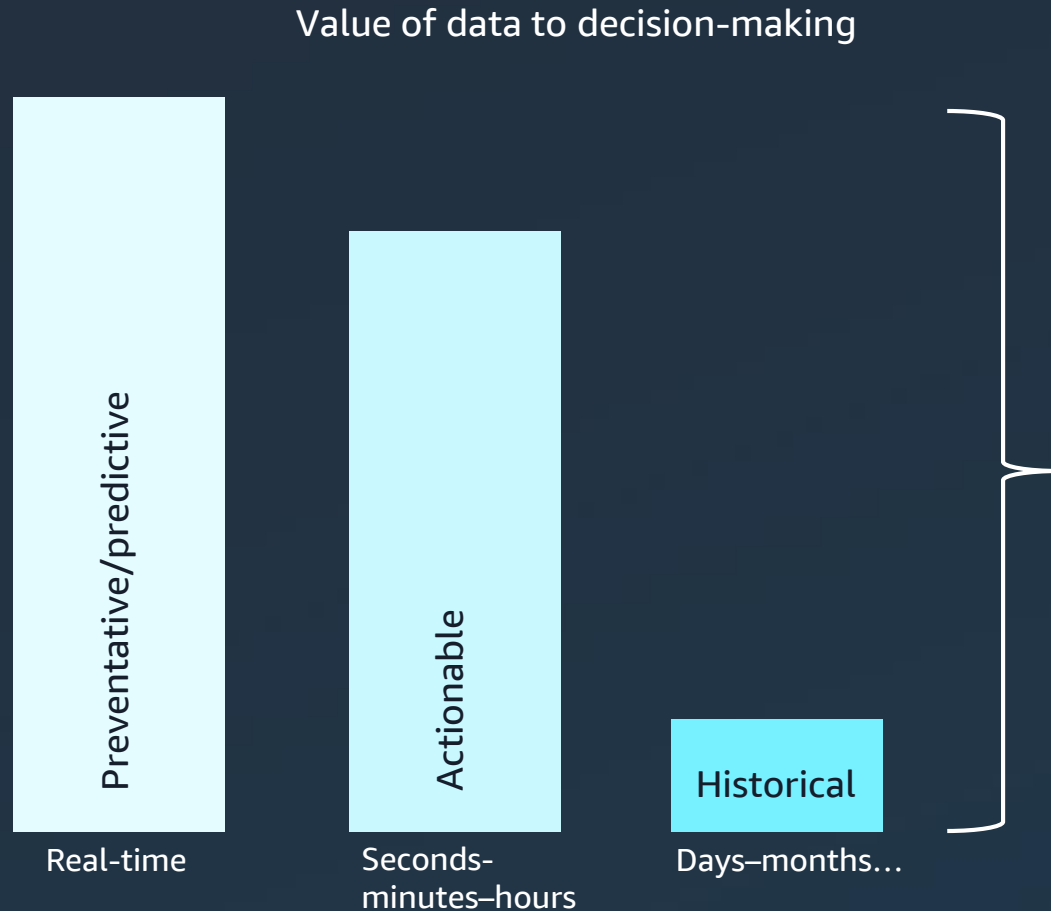
Agenda

- Why invest in real-time data analytics
- Building real-time analytics on AWS
- Use cases and architectures
- Demo: Near-real-time fraud detection using Amazon Redshift Streaming Ingestion with Amazon Kinesis Data Streams and Amazon Redshift ML
- How to get started

Why invest in real-time data analytics ?

The value of data diminishes over time

Data has a short shelf life of actionability. Real-time data lets you act **as fast as the market dictates**.

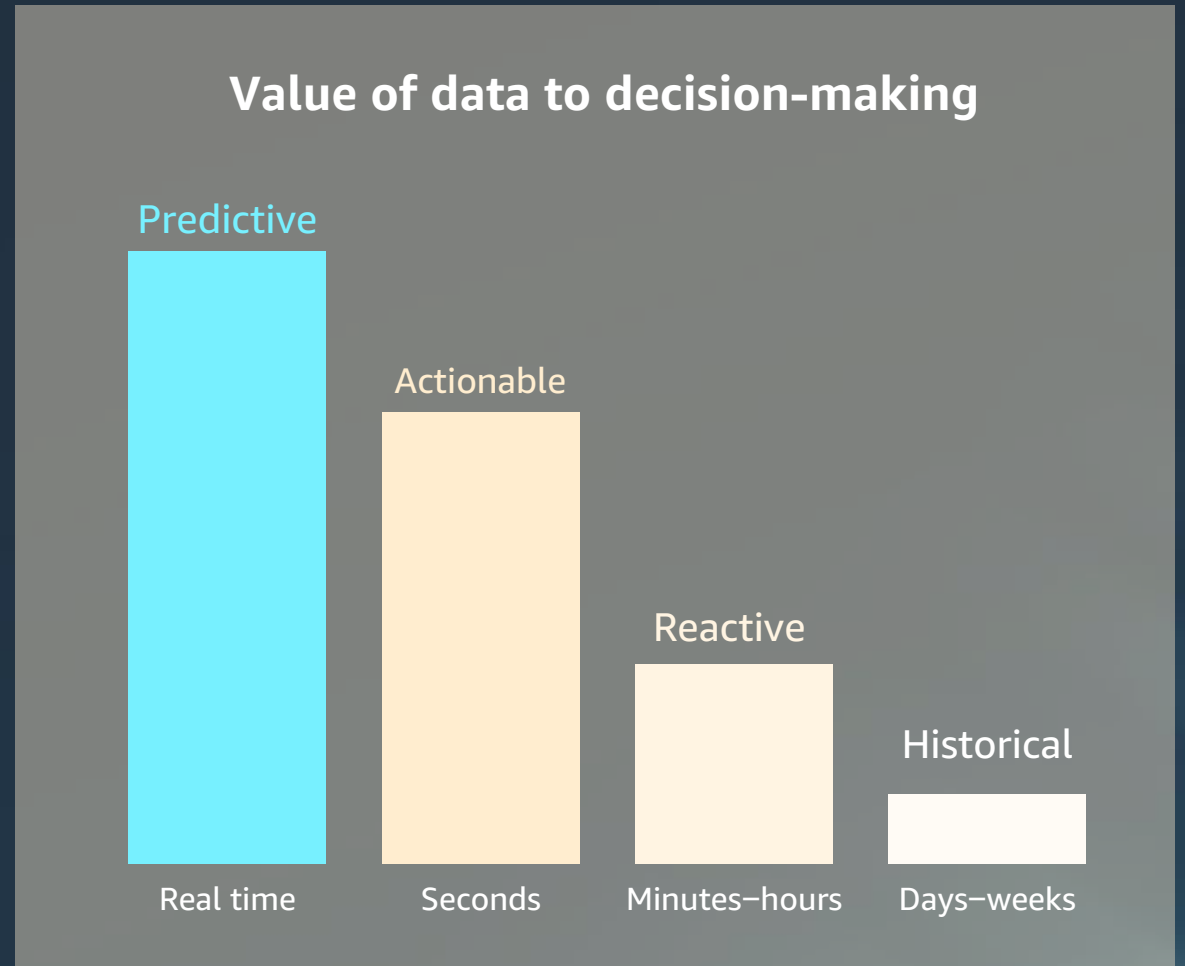


Majority of organizations surveyed believe untimely data inhibited business opportunities¹

The value of the data for insights reduces drastically **within a day**²

Failing to **act in real-time** can translate to real losses

Data have a short shelf life of actionability





Personalization

Chatbots, live agents,
or virtual assistants

Loyalty programs

Automated customer service

Real-time notifications
and feedback

Real-time insights

Real-time monitoring

More accurate ML models

Predictive analytics

Customer 360

Anomaly detection and
fraud prevention

IoT and supply chain
management

Predictive maintenance

Dynamic pricing

Security and threat
management

New use cases

New business models

New products

New services

New experiences

**Enhance customer
experience**




**Better and faster
decision-making**

Improve operations

**Innovate and stay
ahead of competitors**

Streaming analytics anatomy

Ingest, process, and analyze high volumes of high-velocity data from a variety of sources in **real time**



1. Source

Devices and applications that produce real-time data at high velocity

2. Stream ingestion

Data from tens of thousands of data sources can be collected and ingested in real time

3. Stream storage

Data is stored in the order it was received for a set duration of time and can be replayed indefinitely during this time

4. Stream processing

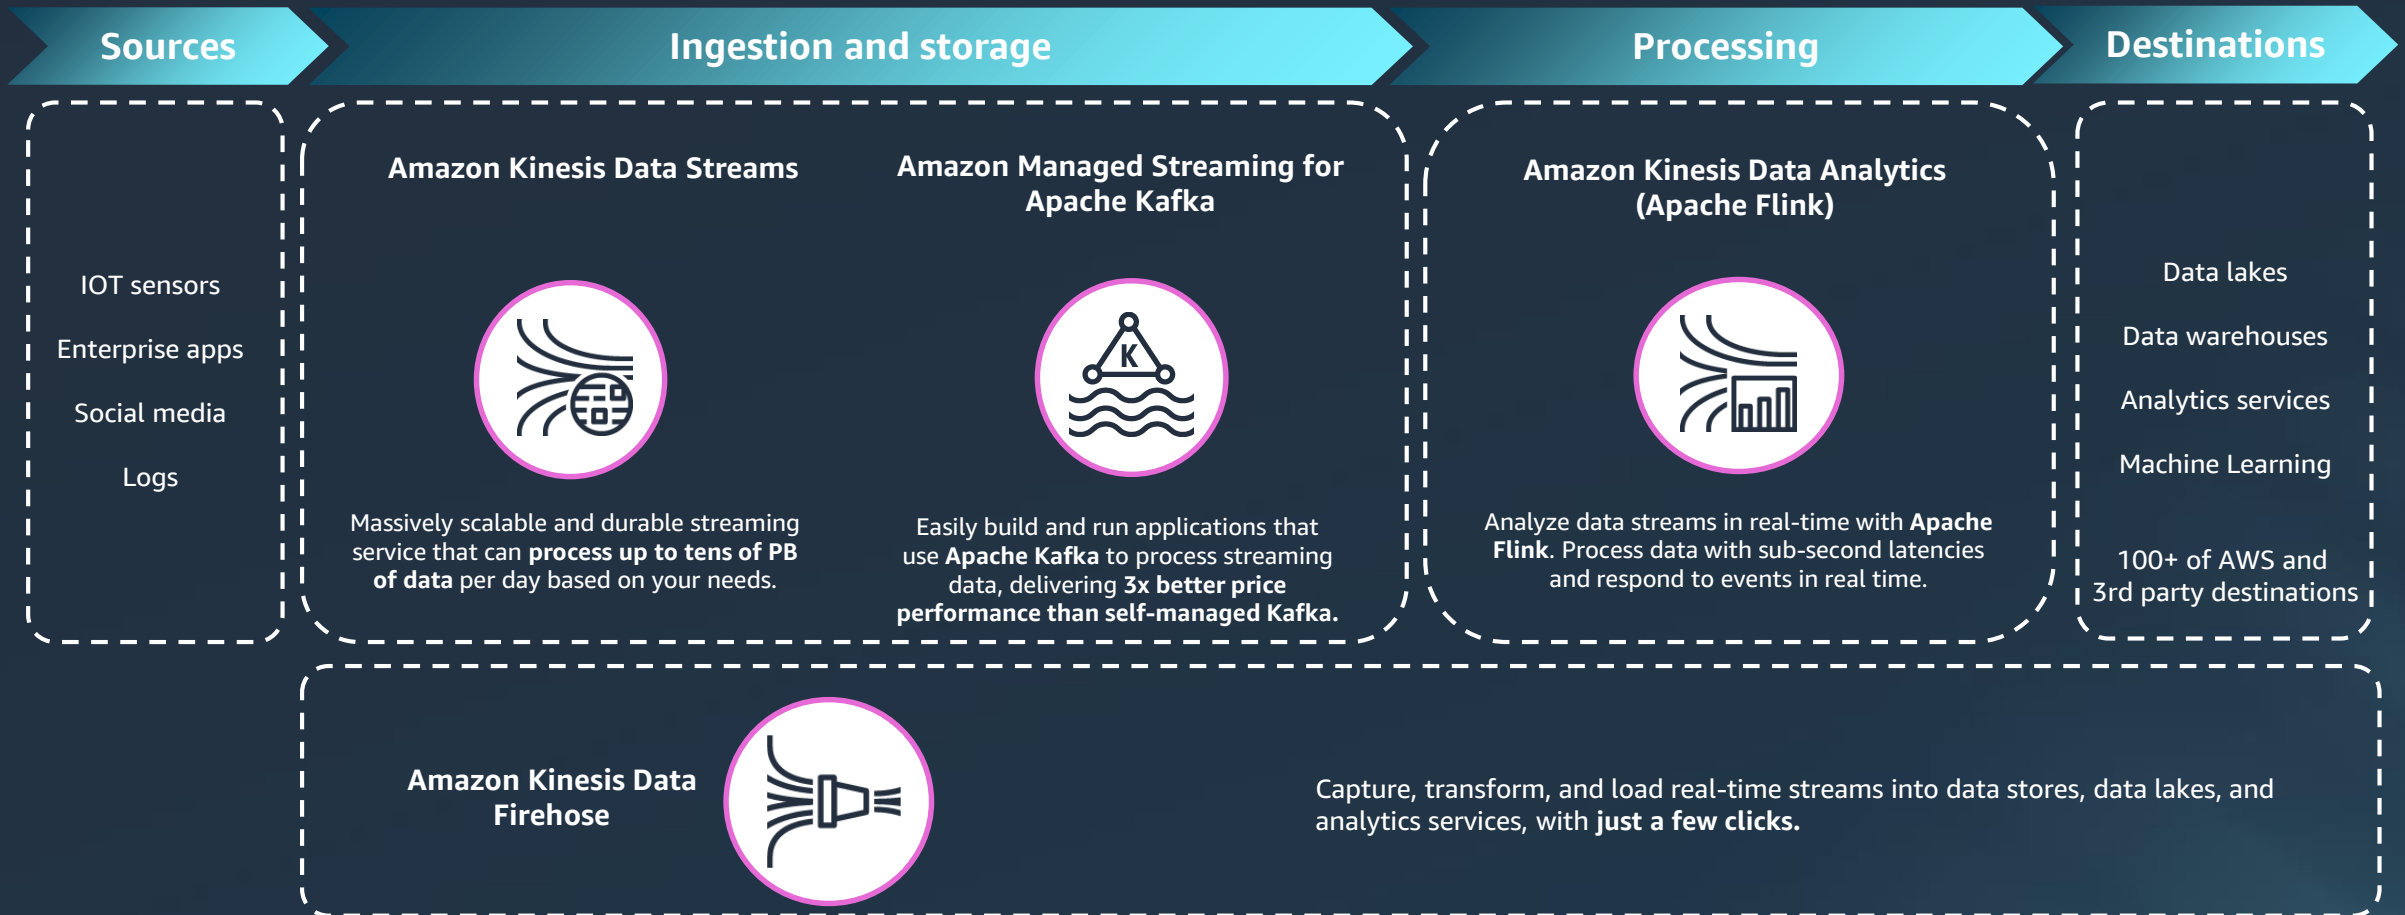
Records are read in the order they are produced, enabling real-time analytics or streaming ETL

5. Destination

Data lake/warehouse (most common)

Database, search

Building end-to-end streaming data pipelines is easier with AWS



Realize greater cost savings by running streaming data on AWS

Reduced Infrastructure Costs

- ✓ **Pay only for what you use**, scale up or down without having to commit to long-term investment or paying for unused resources
- ✓ **No upfront investment** in hardware and software

Reduced Operational Costs

- ✓ **Reduce development and maintenance costs** by handing off infrastructure management to AWS
- ✓ **Free-up IT** so your team can focus on driving innovation

Greater Efficiency and Price Performance

- ✓ Modernize your technology to realize greater savings
- ✓ Amazon MSK delivers **3x better price performance than self-managed Kafka**

Customer examples

Accelerate revenue growth with real-time analytics

Poshmark

A leading social marketplace for new and second hand style fashion goods with over **80M+ customers and 200M+ available listings**

Goal

Drive **top-line revenue growth** through **enhancing user experience and personalization** and **fraud prevention**

Challenges

The initial approach of using batch processing for personalization and security did not meet expectations for customer experience improvement

Solution

Designed a **real-time personalization and fraud prevention solution** using real-time event capture and real-time data enrichment leveraging AWS streaming data services, machine learning capabilities, and storage solutions

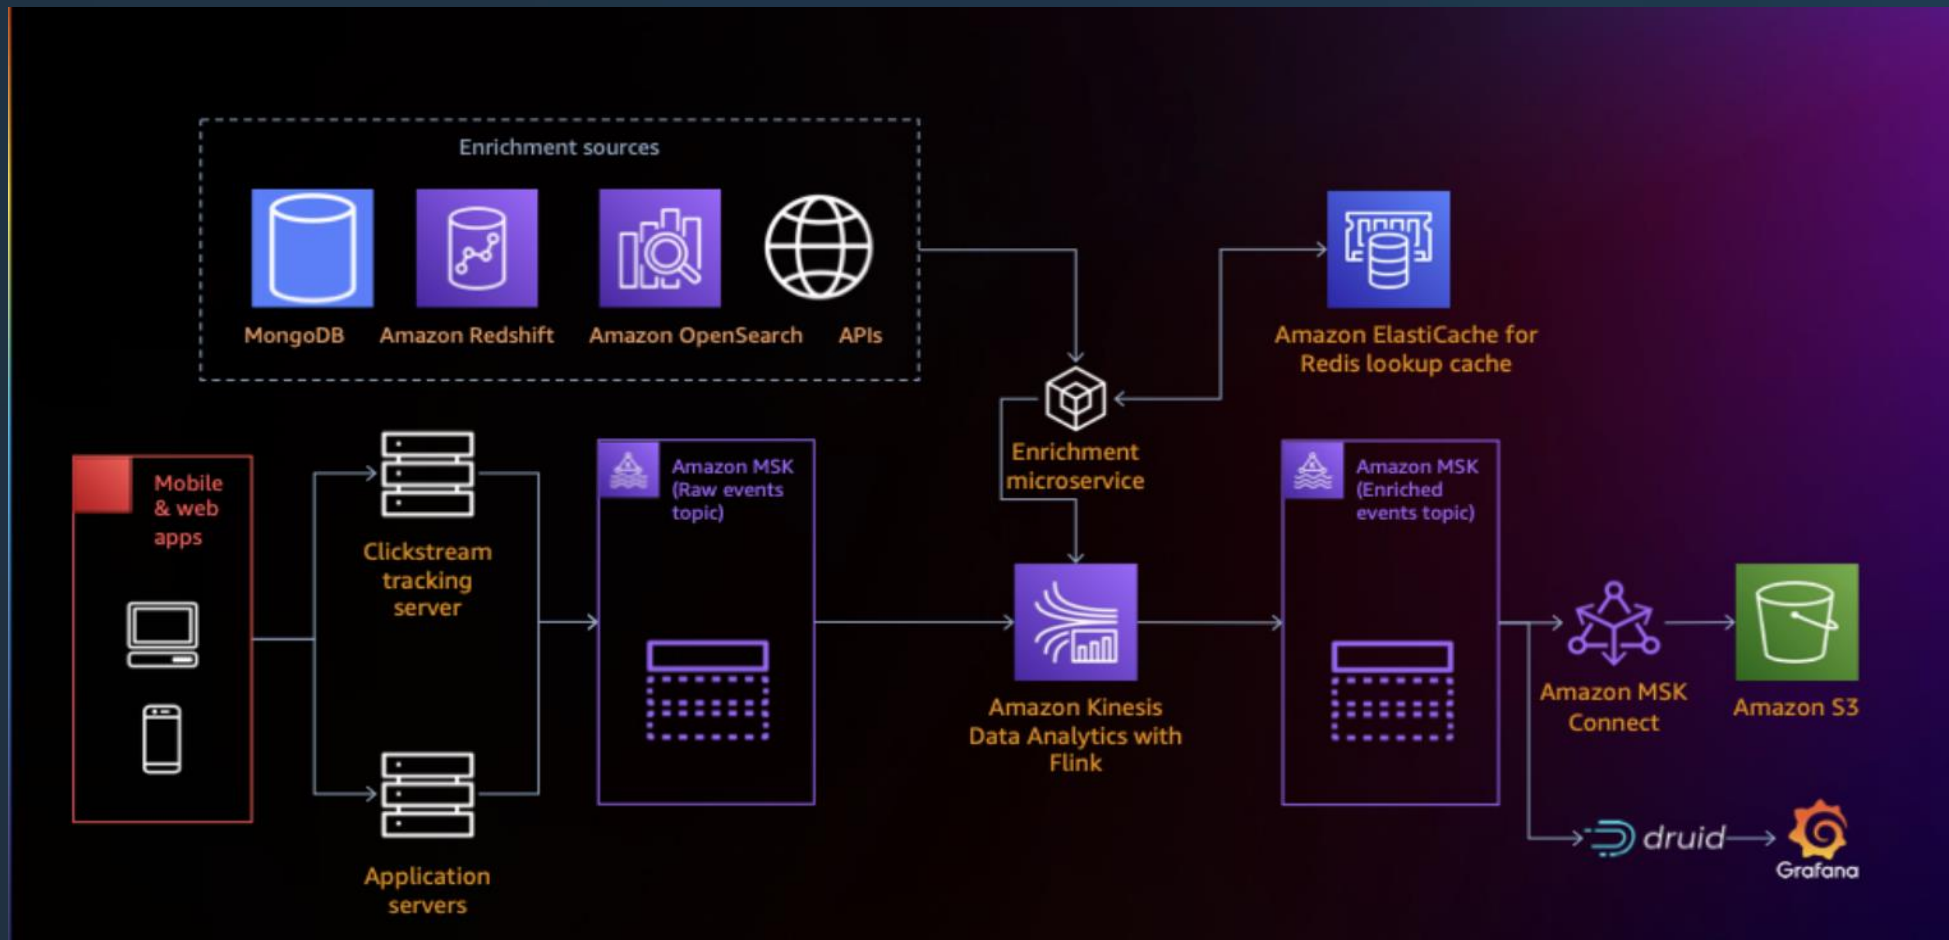
Results

- ✓ Drove top-line revenue growth, **increasing customer click through rate by 8% (vs batch processing)**
- ✓ Improved alerting of Account Take Over(ATO)/Fraud from 45% with batch based system to **80% with real-time monitoring**

[Learn more](#)



Real-time analytics solution architecture

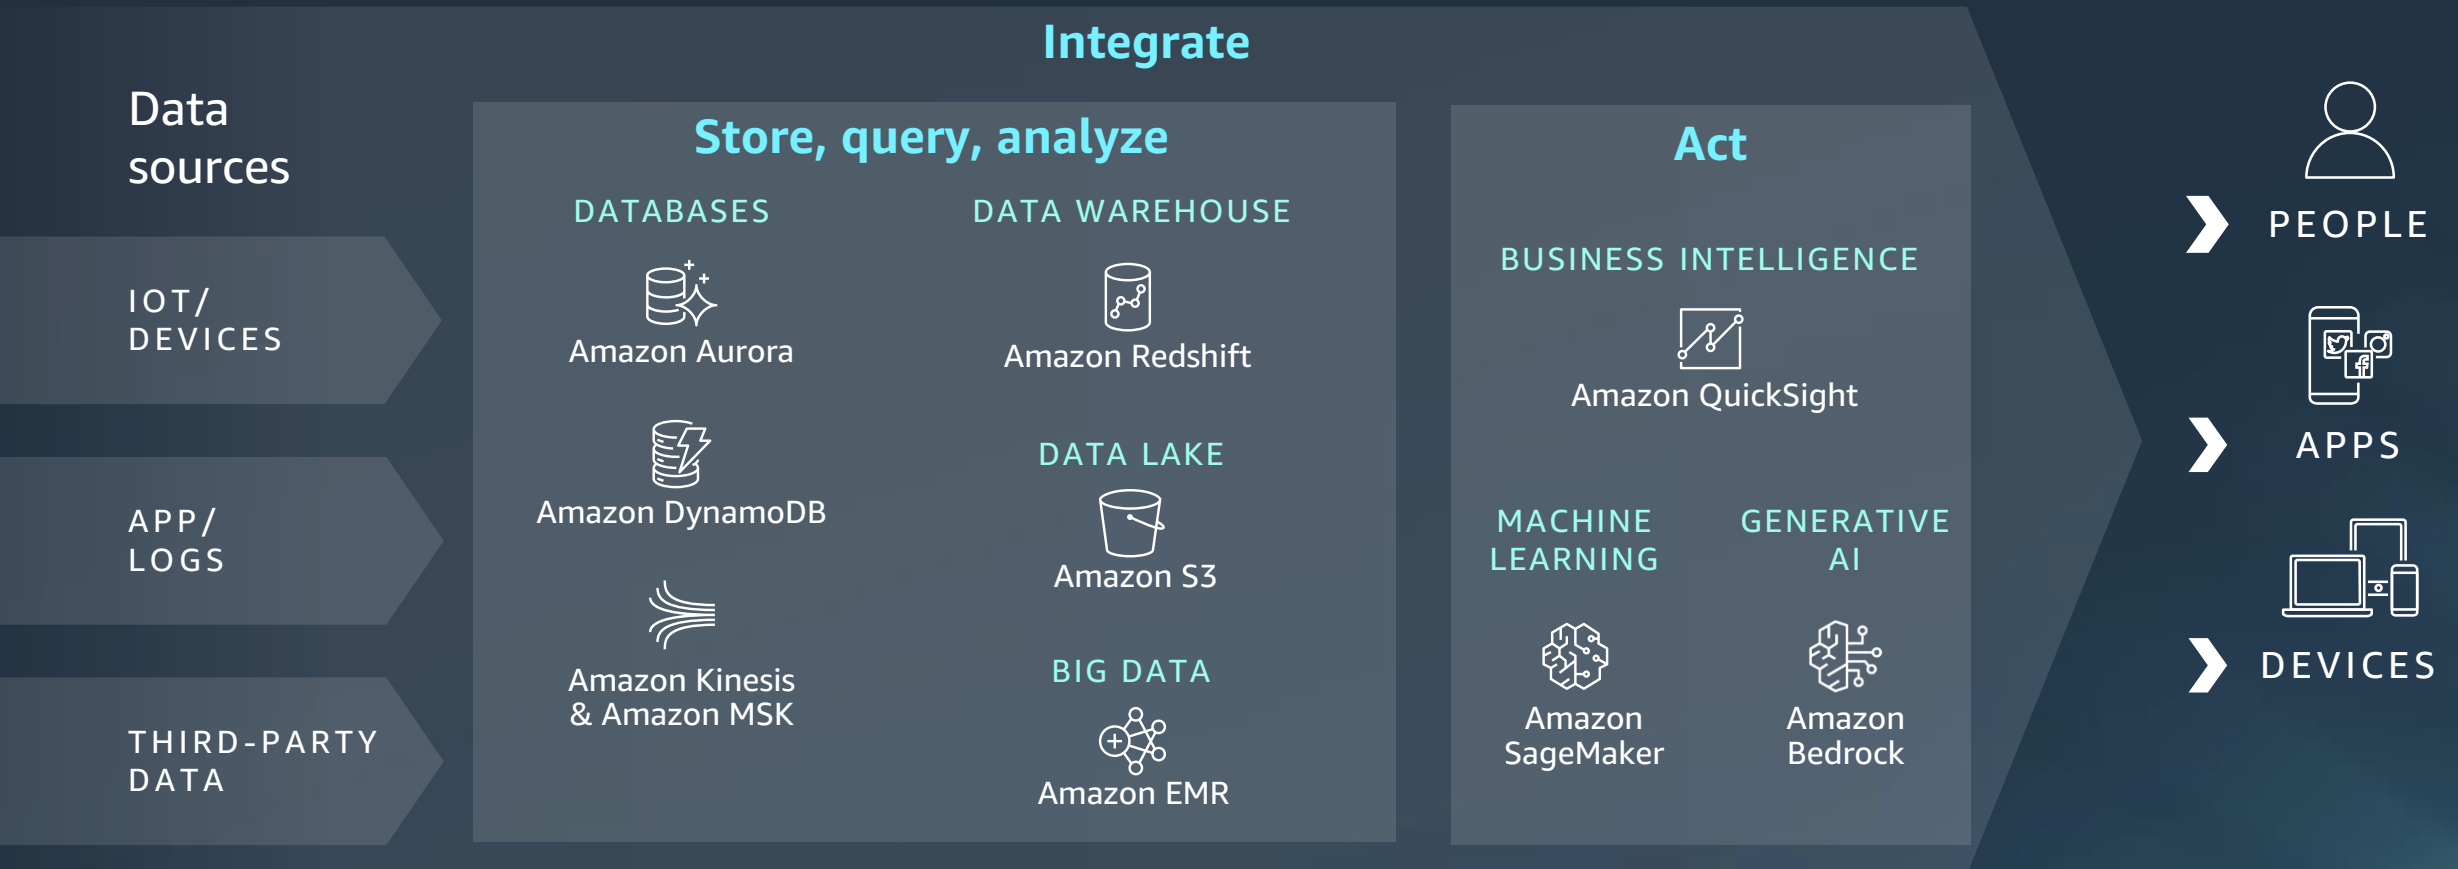


<https://aws.amazon.com/blogs/big-data/accelerating-revenue-growth-with-real-time-analytics-poshmarks-journey>

Streaming analytics with AI/ML

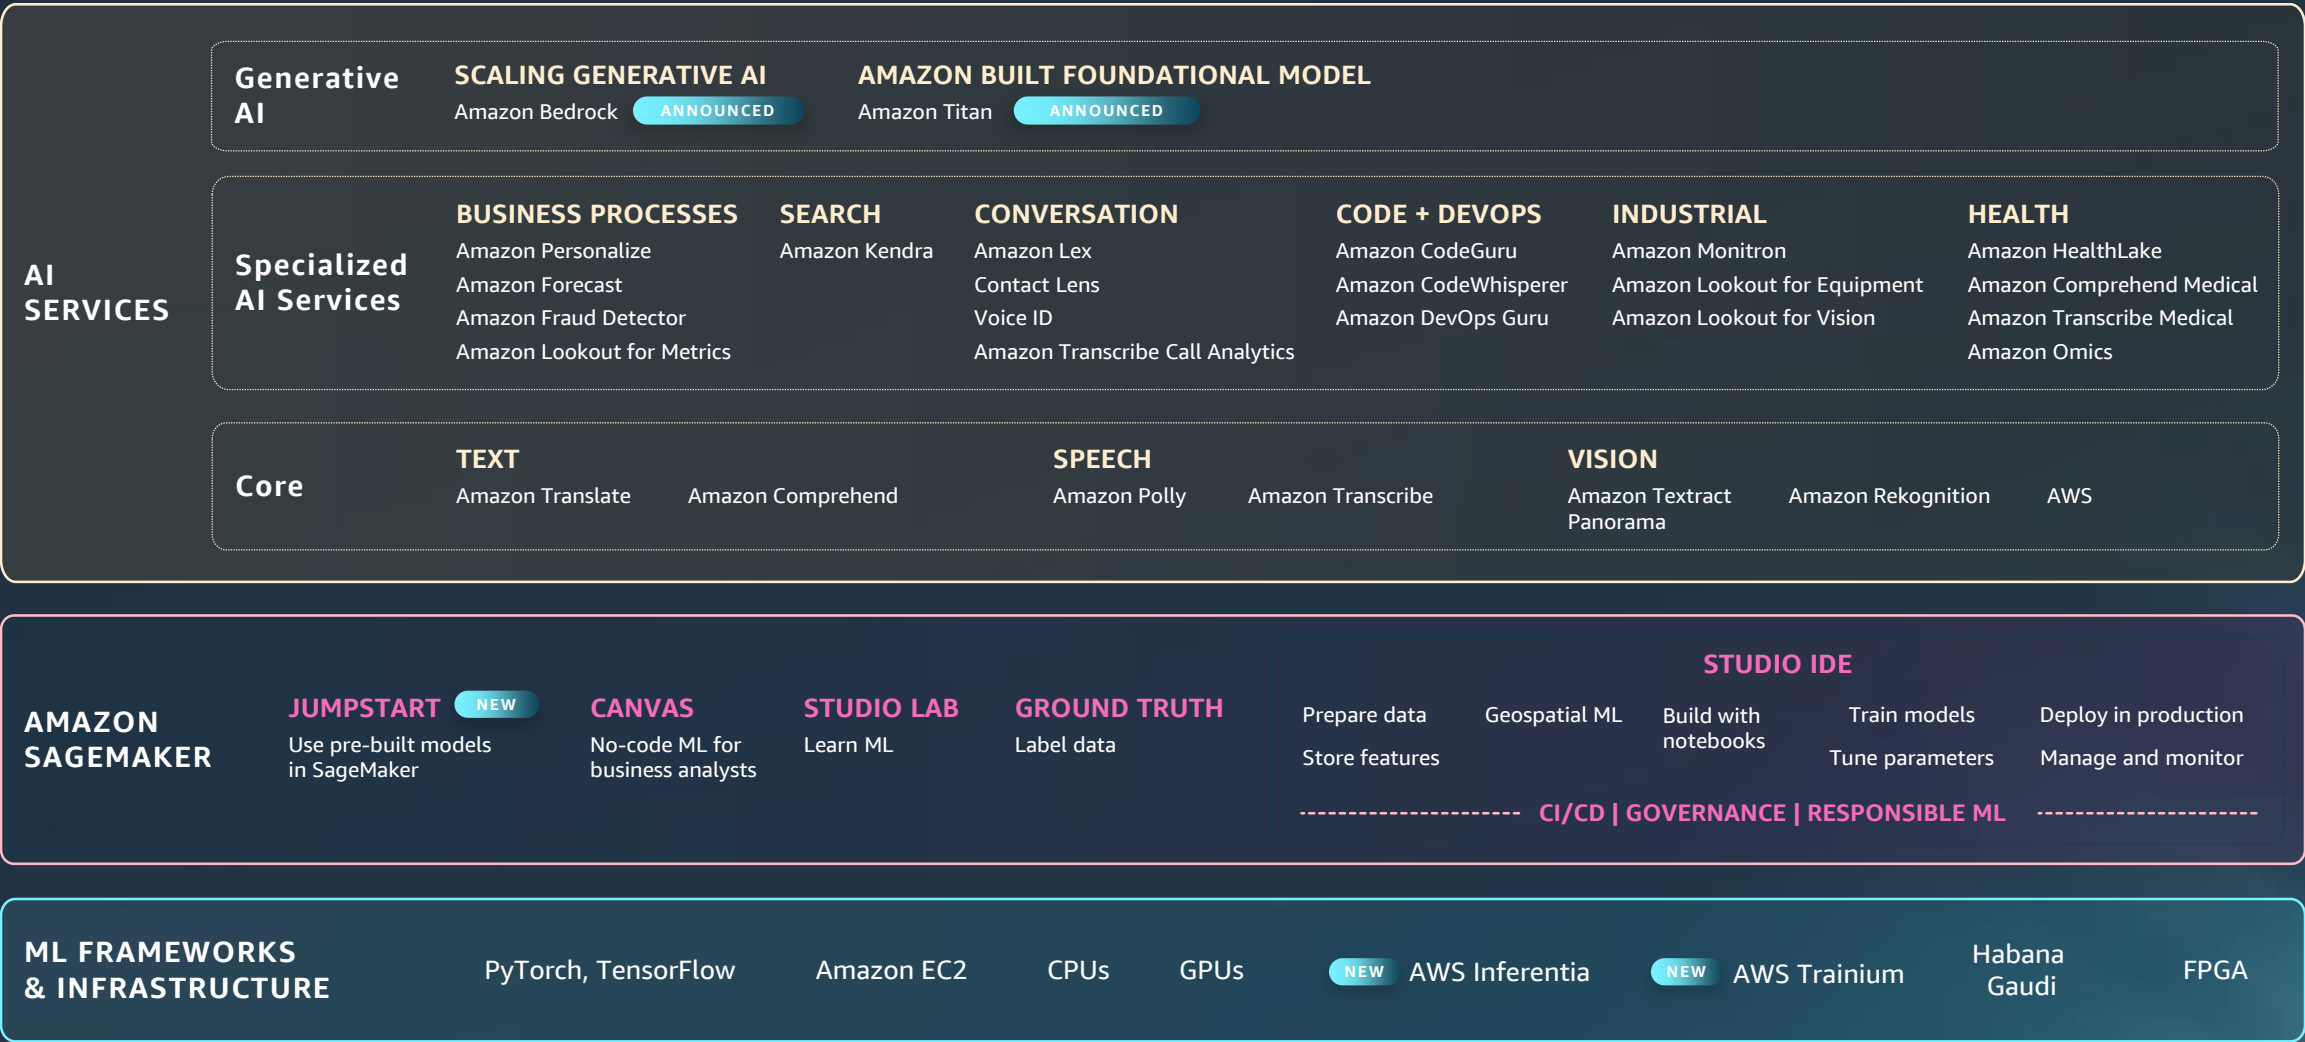


End-to-end data foundation



Catalog and govern | AWS Lake Formation and Amazon DataZone

The AWS AI/ML Stack



AWS brings ML closer to data



Databases

+



Data warehouses
+ data lakes

+



Business
intelligence tools

AMAZON
AURORA ML



AMAZON
NEPTUNE ML



AMAZON
REDSHIFT ML



AMAZON
ATHENA ML

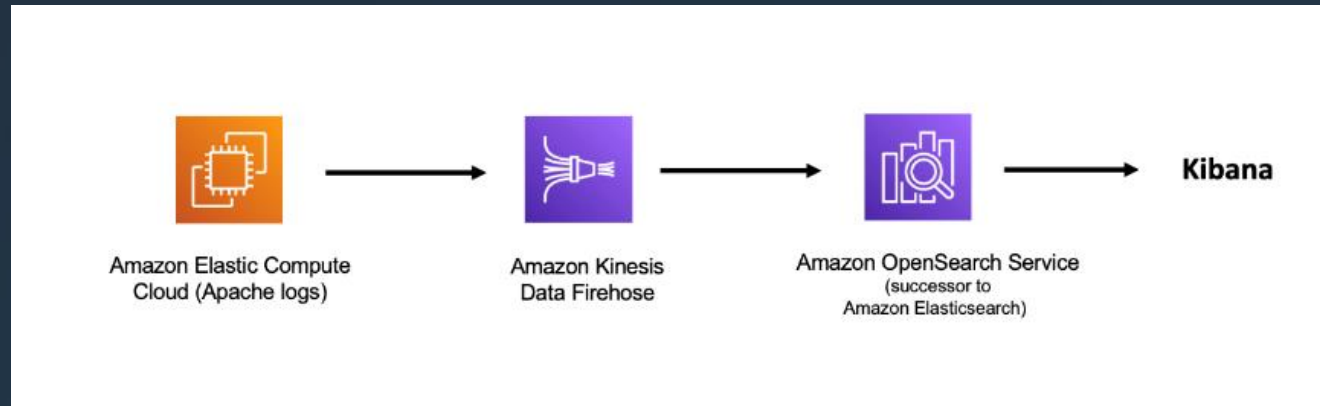
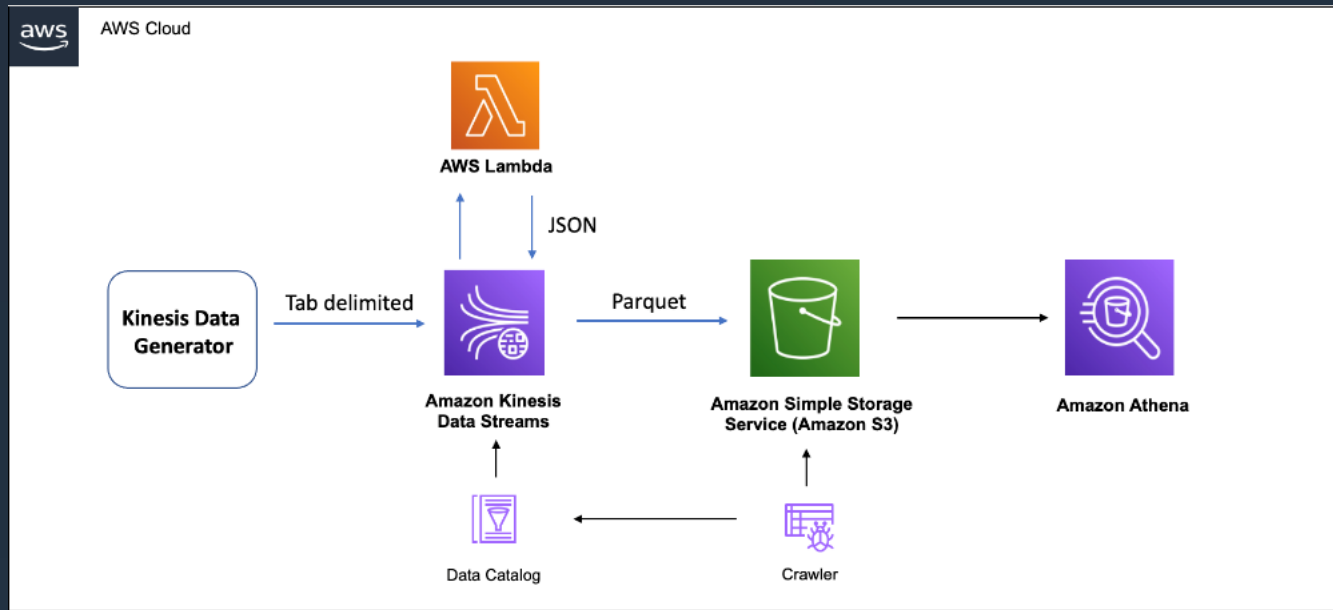


AMAZON
QUICKSIGHT ML

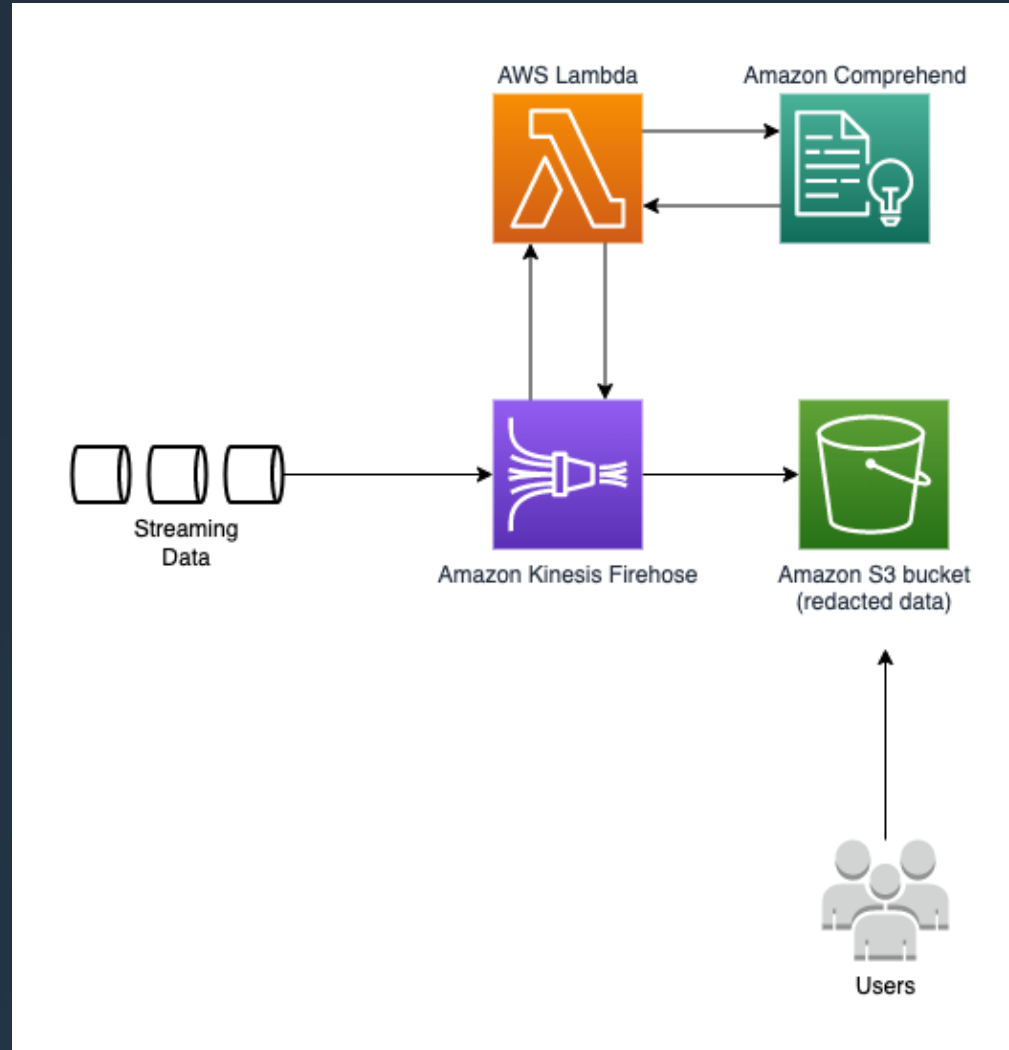


Bringing it together: Use cases and architectures

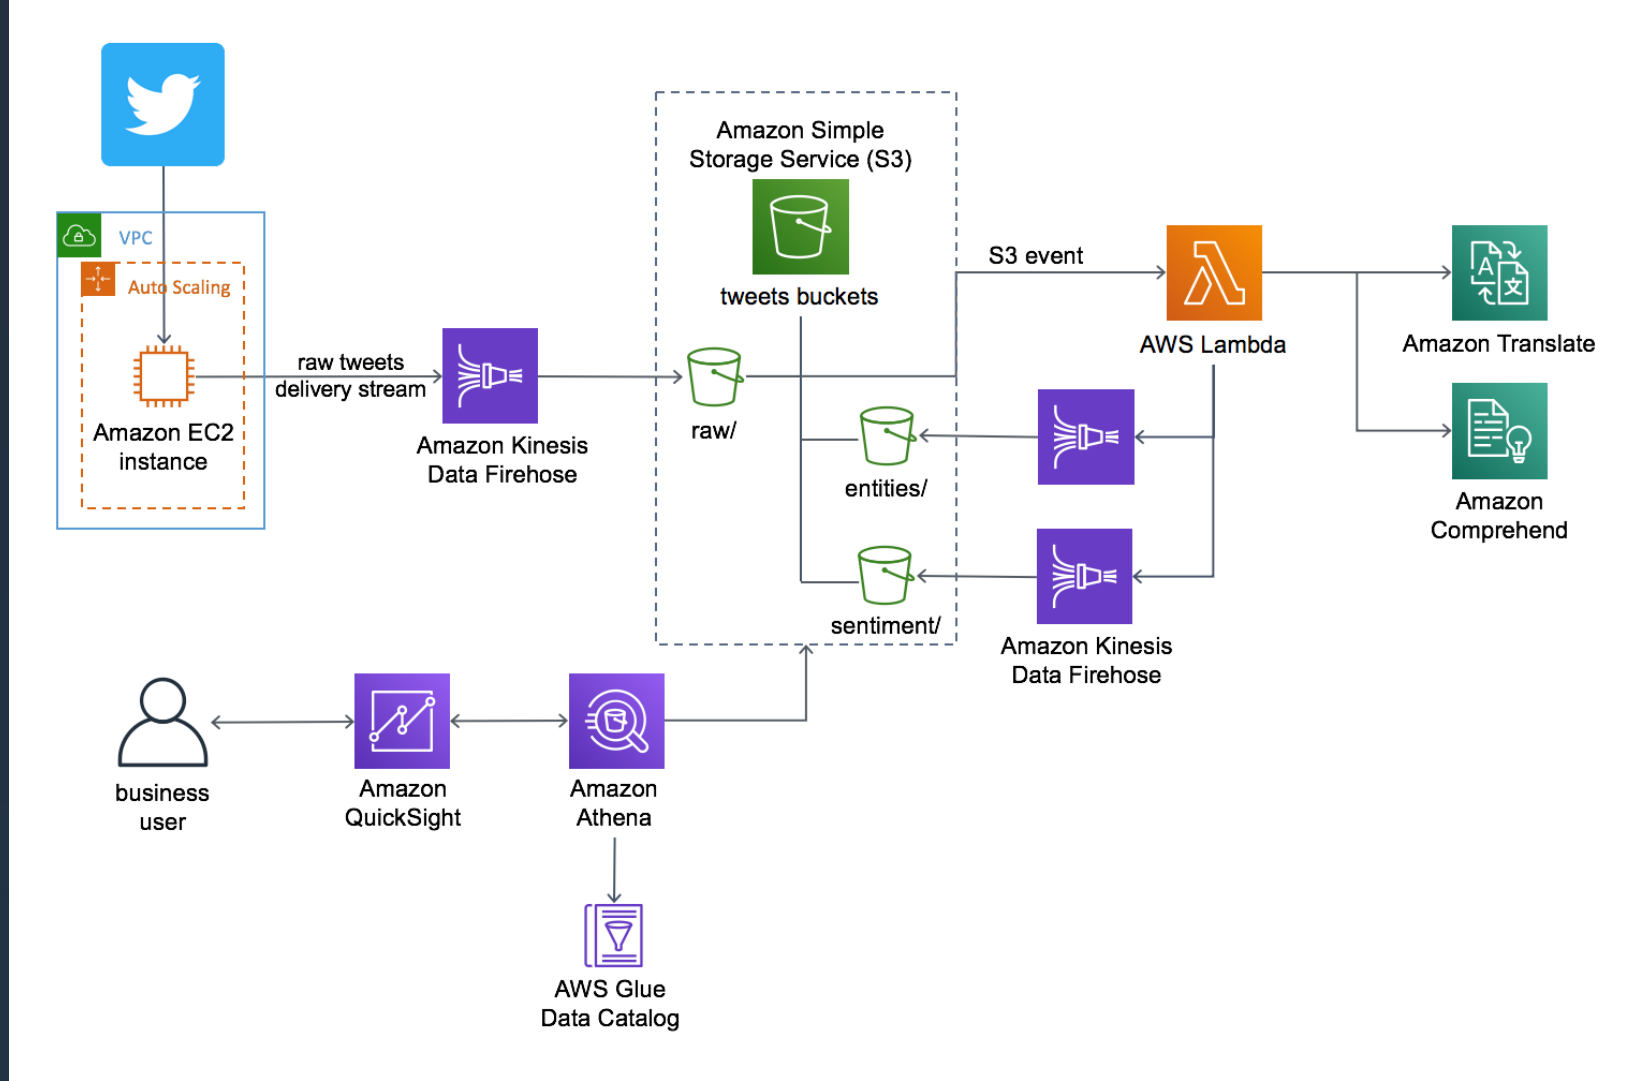
Use case 1: Simple use case of real time log analytics with Amazon Kinesis Firehose



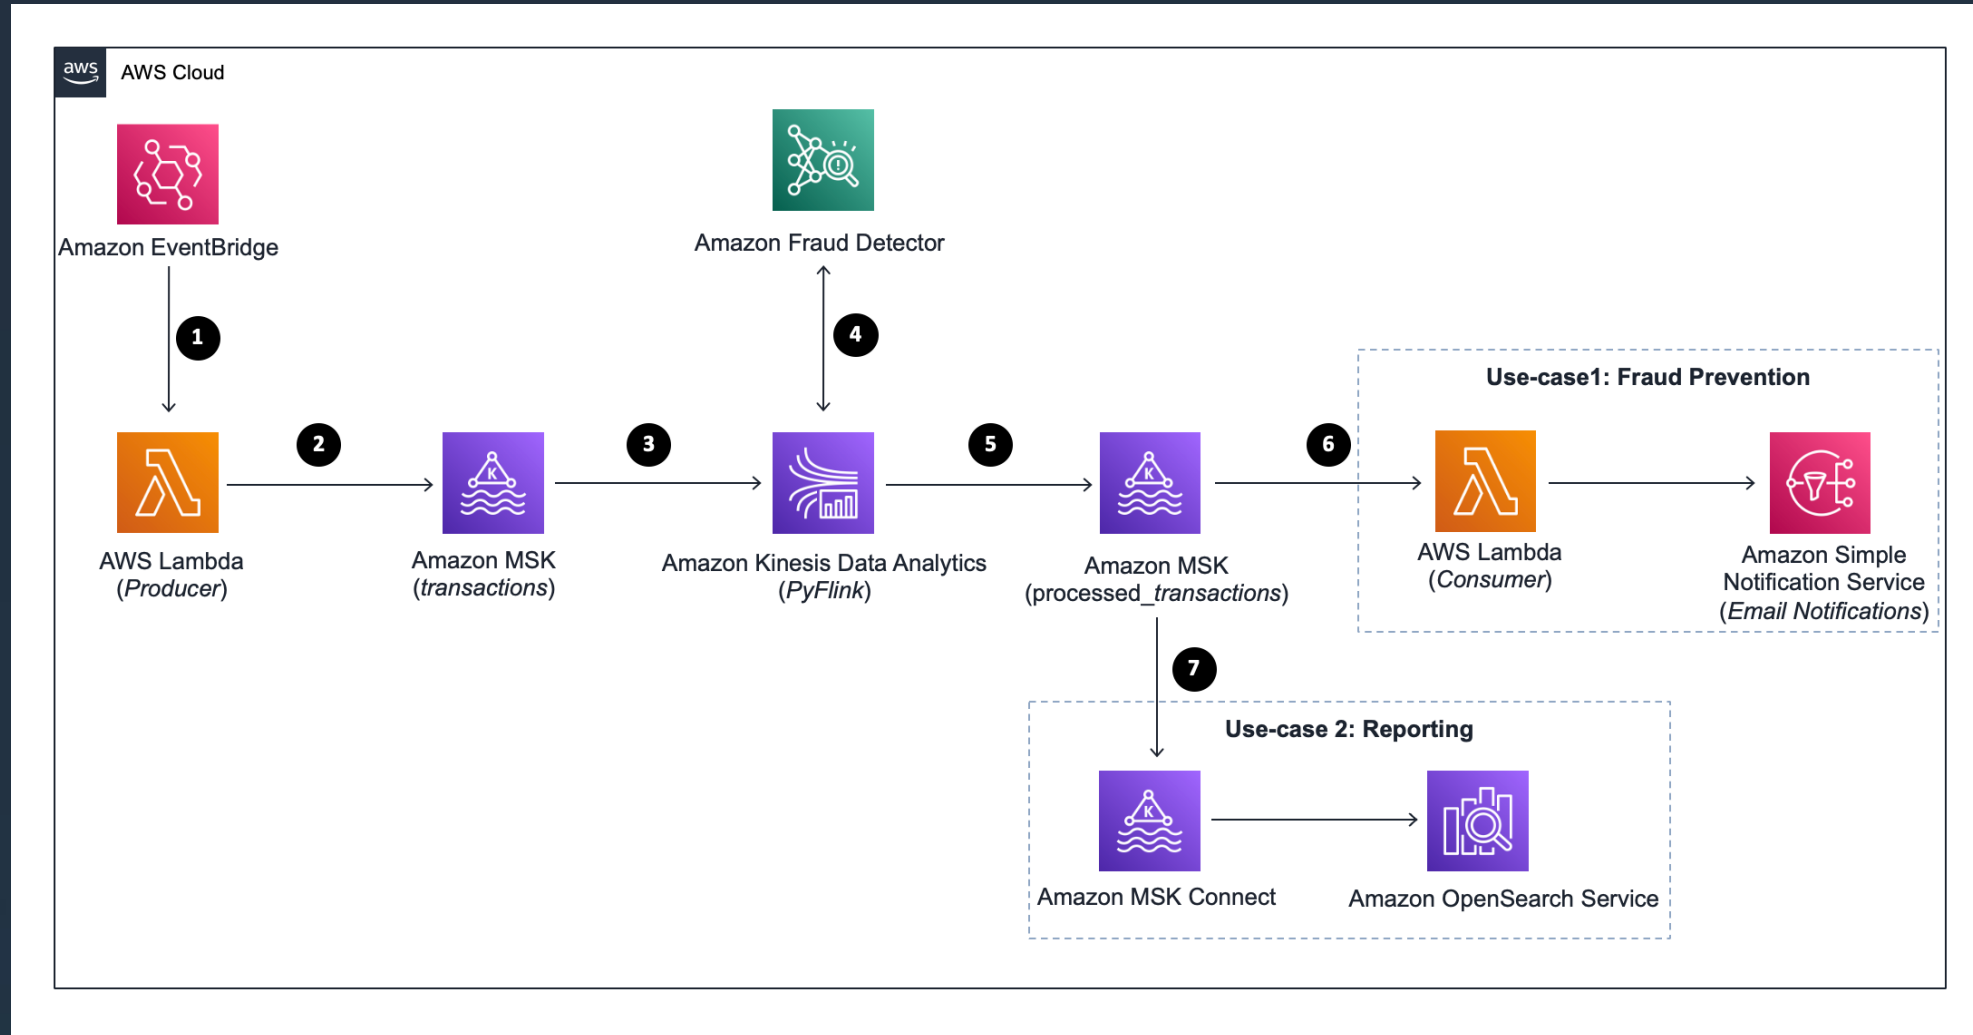
Use case 2: Redact sensitive data from streaming data in near-real time using Amazon Comprehend and Amazon Kinesis Data Firehose



Use case 3: AI-driven social media dashboard



Use case 4: Build and visualize a real-time fraud prevention system using Amazon Fraud Detector



Score user activity in real-time with Amazon Fraud Detector

Identify potential fraudulent activity with a fully managed service that uses ml models and fraud detection expertise from AWS and amazon.com



User performs an online activity



App or website captures details about the activity, session, and user

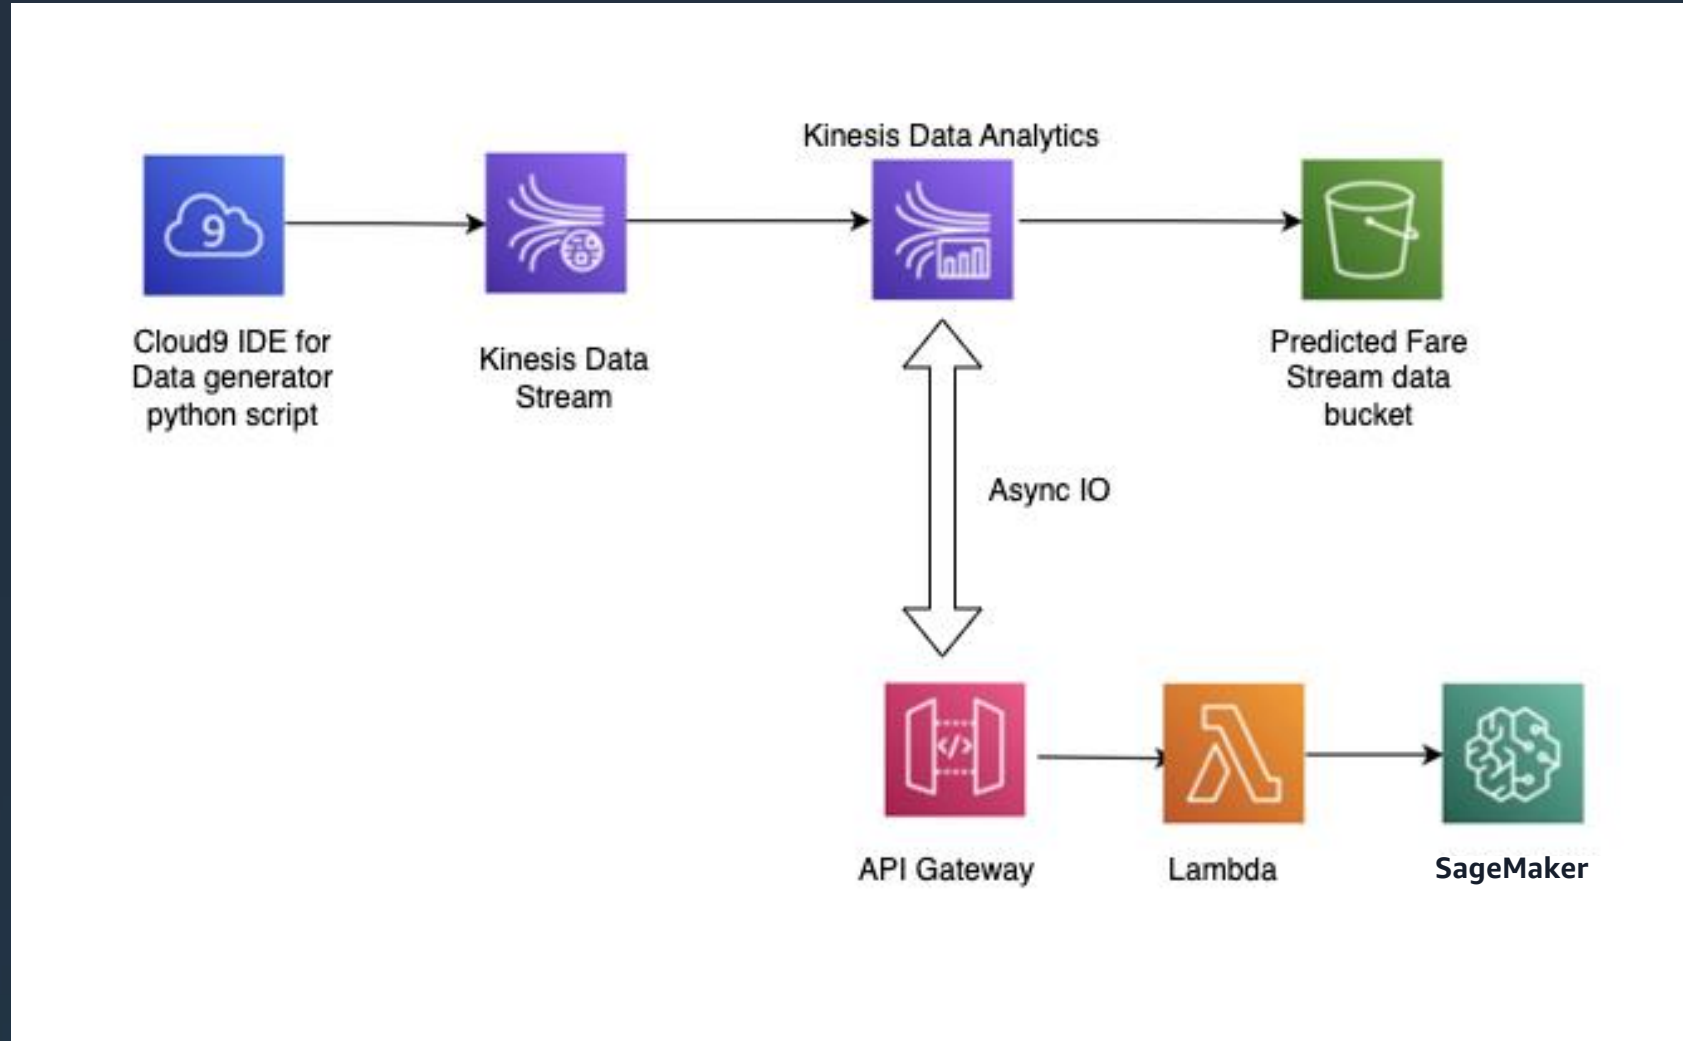


Details are sent to Amazon Fraud Detector activity fraud model API



Outcomes are sent back in real-time to allow automation of decisions

Use case 5: Real-time analytics and ML inference – Predict fares



Challenges with building time-series solutions

RELATIONAL DATABASES



Inefficient at processing time-series data




Data management issues with rigid schema



Limited integrations for ML, analytics, and data collection




EXISTING TIME-SERIES SOLUTIONS



Difficult to scale for large volumes of data



Minimal data lifecycle management



Real-time and historical data are decoupled



Addressing time-series data challenges with Amazon Timestream

Fast, scalable, and serverless time-series database

Serverless and easy to use



No servers to manage or instances to provision; software patches, indexes, and database optimizations are handled automatically

Performance at scale



Capable of ingesting trillions of events daily; the adaptive SQL query engine provides rapid point-in-time queries with its in-memory store, and fast analytical queries through its magnetic store

Purpose built for time-series data



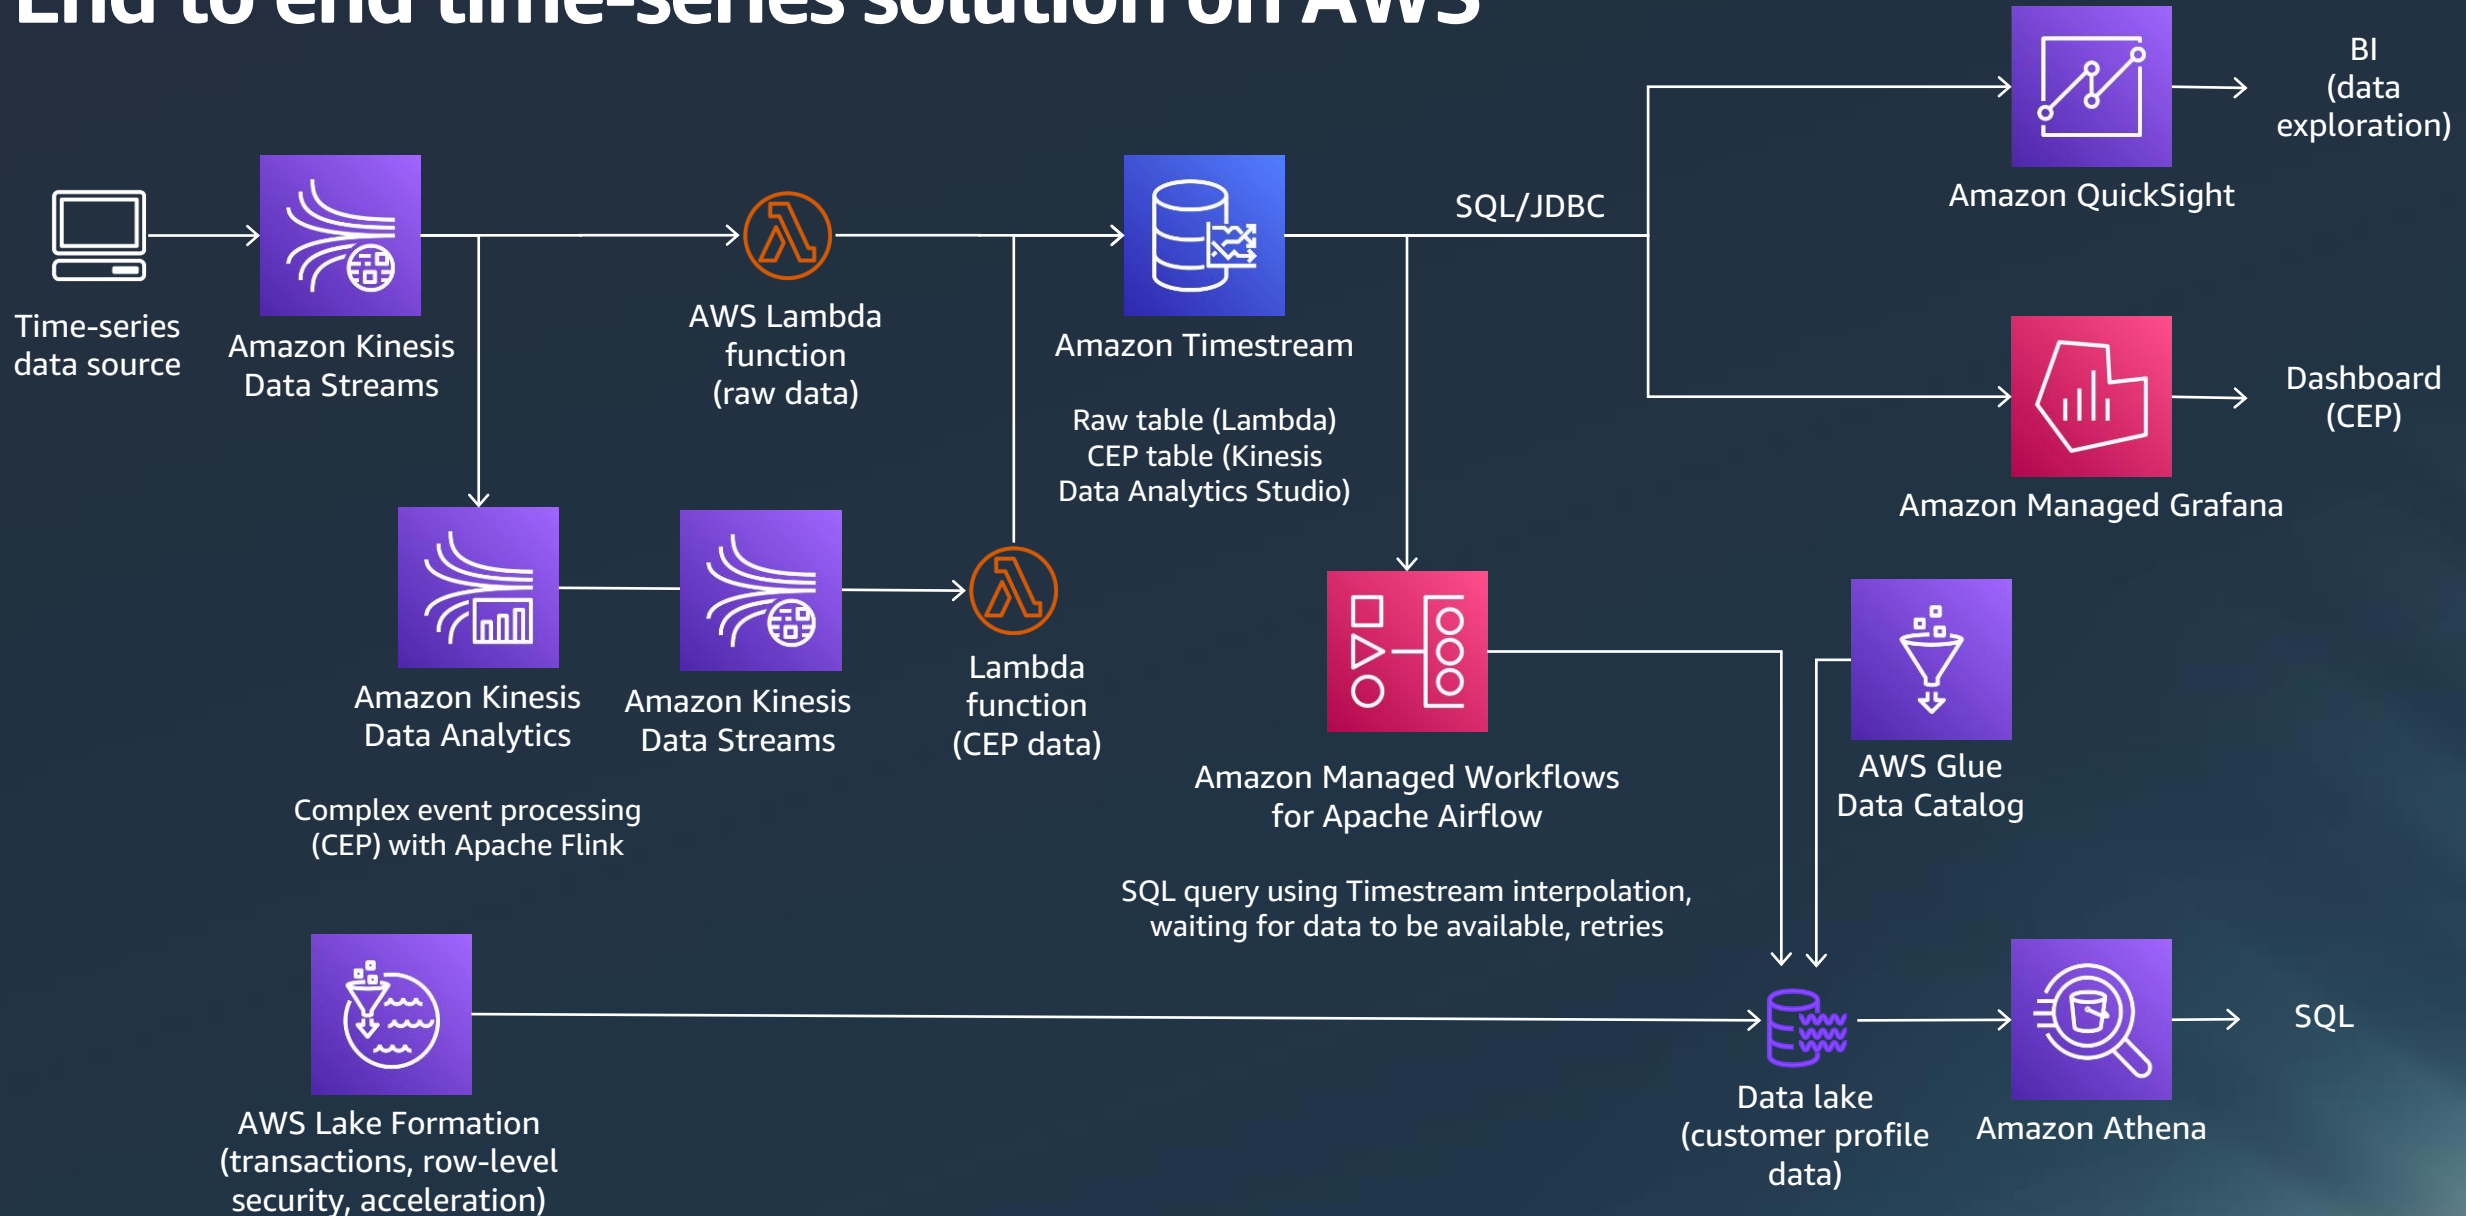
Built-in analytics using standard SQL with added interpolation and smoothing functions to identify trends, patterns, and anomalies

Secure from the ground up

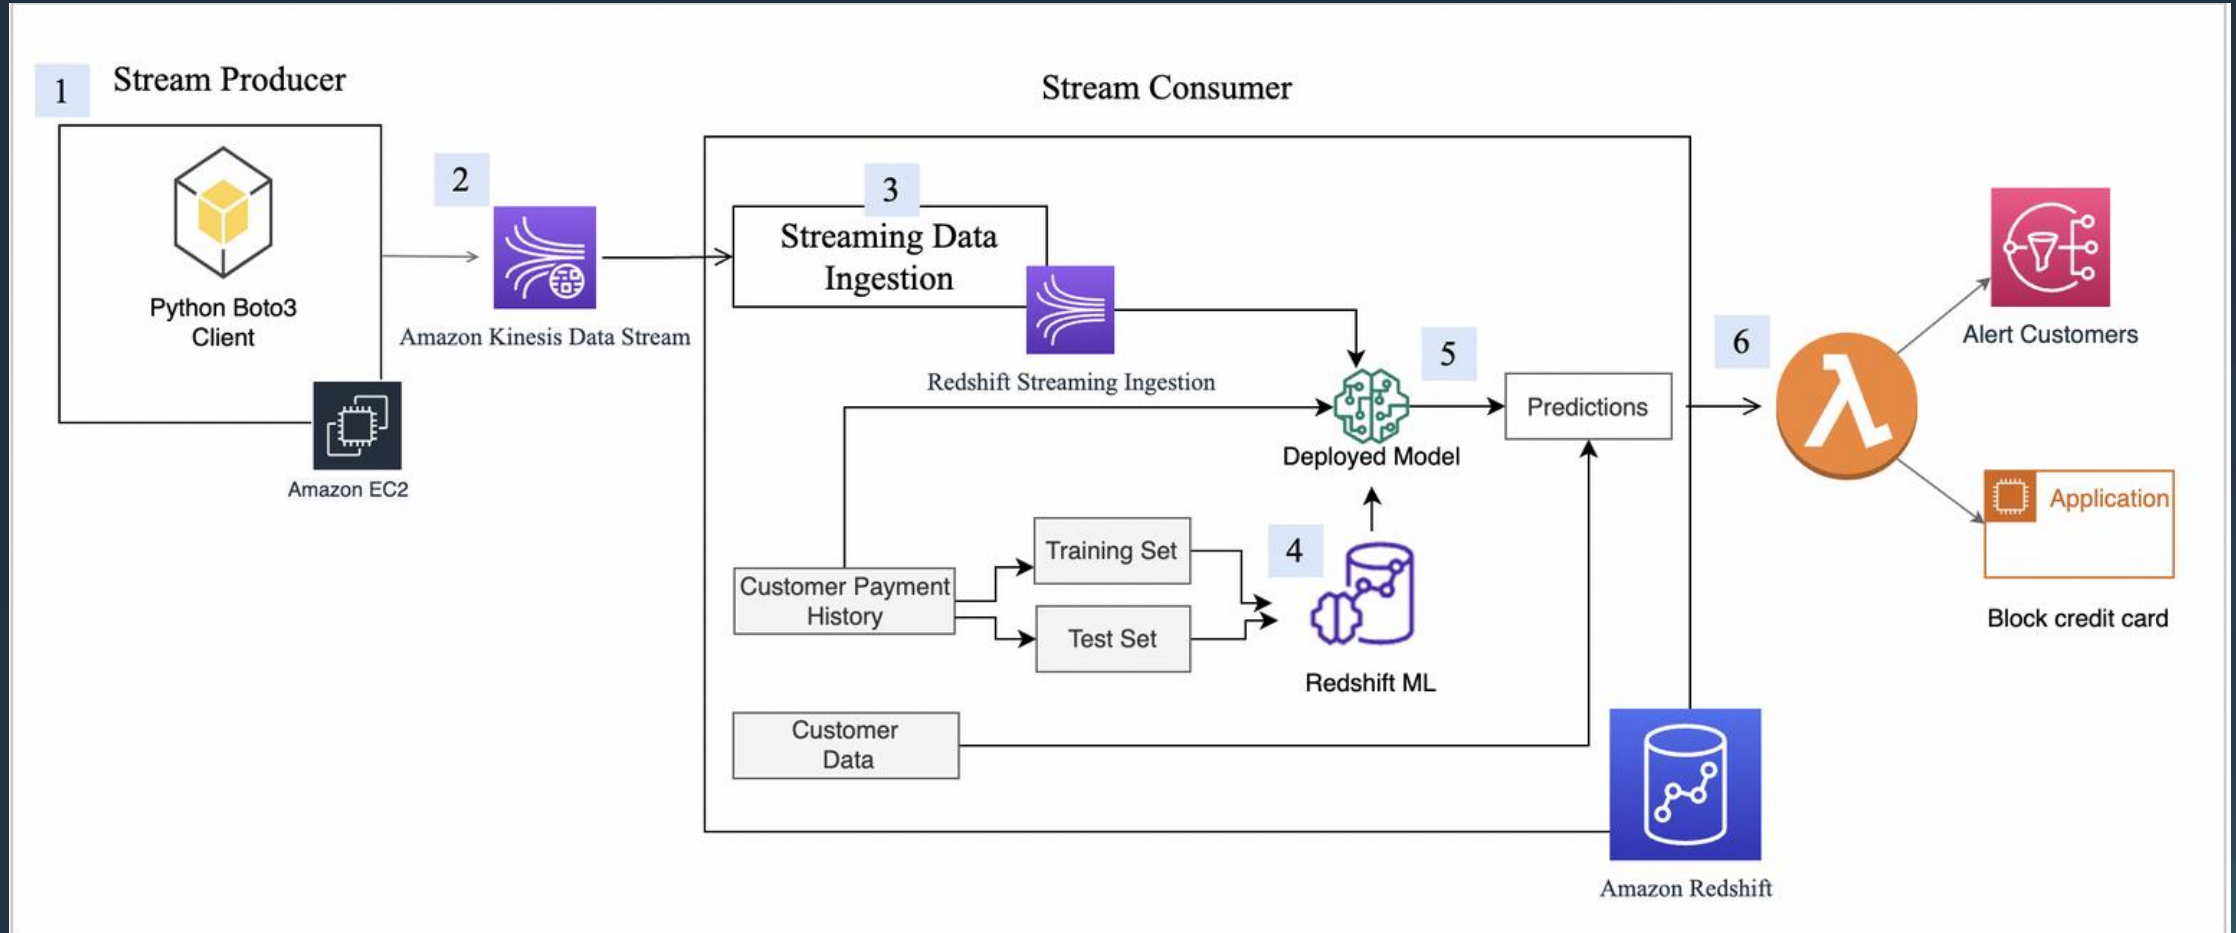


All data is encrypted in flight and at rest using AWS Key Management Service (AWS KMS) with customer-managed keys (CMKs)

End to end time-series solution on AWS



Demo : Near-real-time fraud detection using Amazon Redshift Streaming Ingestion with Amazon Kinesis Data Streams and Amazon Redshift ML



Demo



Recap

- Why is real time analytics important
- What is streaming ingestion with AWS
- How to integrate AWS AI/ML with streaming data
- Where to apply: Use cases and architectures
- Demo: How to build near real time fraud detection

Getting started

Try it yourself

- Near real time fraud detection - [here](#)
- Streaming ingestion to feature store - [here](#)
- Real time ML inference workshop - [here](#)
- Customer story - [here](#)
- AWS solution example - [here](#)

Visit the AWS Data Resource Hub

A modern data strategy can help you manage, act on, and react to your data so you can make better decisions, respond faster, and uncover new opportunities. Dive deeper with these resources today.

- Maximizing business value with data: 6 essential data-driven use cases for leaders
- The ultimate guide to developing an end-to-end data strategy
- 5 key elements of a data-driven culture
- A builder's guide to AWS analytics services
- ... and more!



<https://tinyurl.com/data-hub-aws>

[Visit resource hub](#)



AWS Training and Certification

Access 600+ free digital courses with AWS Skill Builder

Focus on the cloud skills and services that are most relevant to you across 30+ AWS solutions, including digital self-paced learning plans and ramp-up guides

- Build your future in the AWS Cloud at your own pace
- Advance your skills and knowledge with learning plans
- Validate your cloud expertise with AWS Certification



LEARN YOUR WAY [SKILLBUILDER.AWS](https://skillbuilder.aws)



Thank you for attending AWS Innovate – Data Edition

We hope you found it interesting! A kind reminder to **complete the survey**.
Let us know what you thought of today's event and how we can improve the event experience for you in the future.



aws-apj-marketing@amazon.com



twitter.com/AWSCloud



facebook.com/AmazonWebServices



youtube.com/user/AmazonWebServices



slideshare.net/AmazonWebServices



twitch.tv/aws

Thank you!

