



# Simplify your Streaming: Delta Live Tables



---

**Jerrold Law**

**Specialist Solutions Architect**



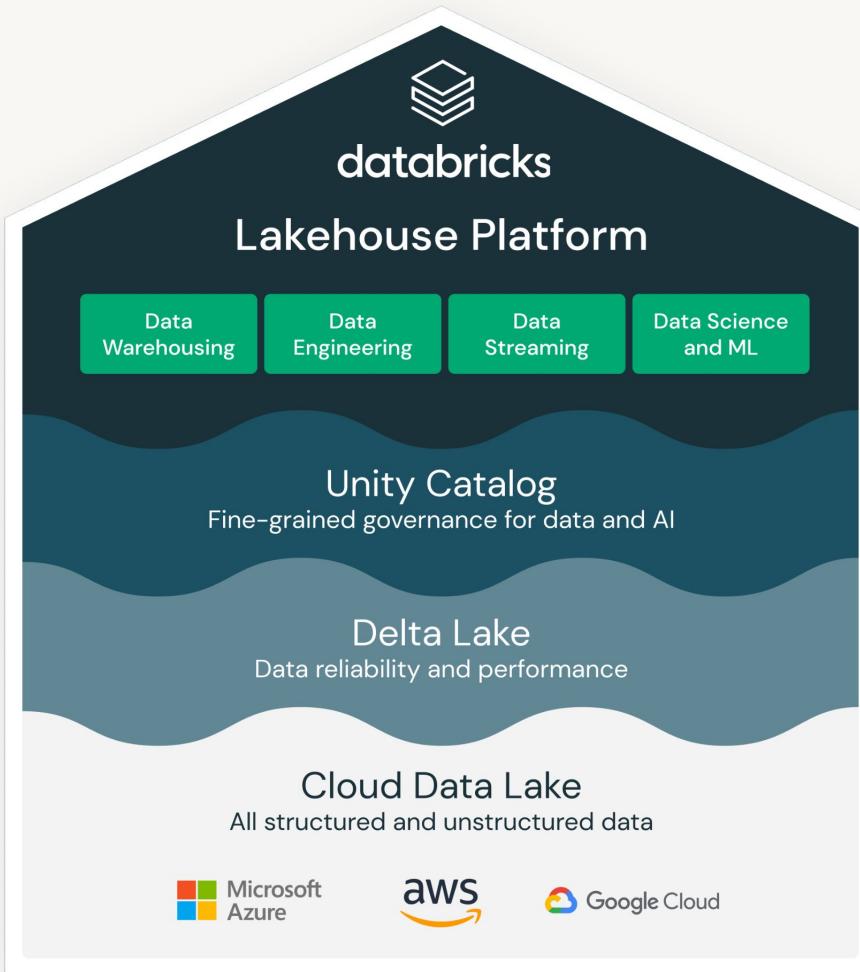
# Housekeeping

- Your connection will be muted
- We will share recording with all attendees after the session
- Submit questions in the Q&A panel
- If we do not answer your question during the event, we will follow-up with you to get you the information you need!



# The best data warehouse is a lakehouse

Why is the lakehouse your next data warehouse?



- 1 Data Engineering, Analytics, Streaming, and AI in one place
- 2 World-class performance with data lake economics
- 3 One source of truth for all your data





# Good data is the foundation of a Lakehouse

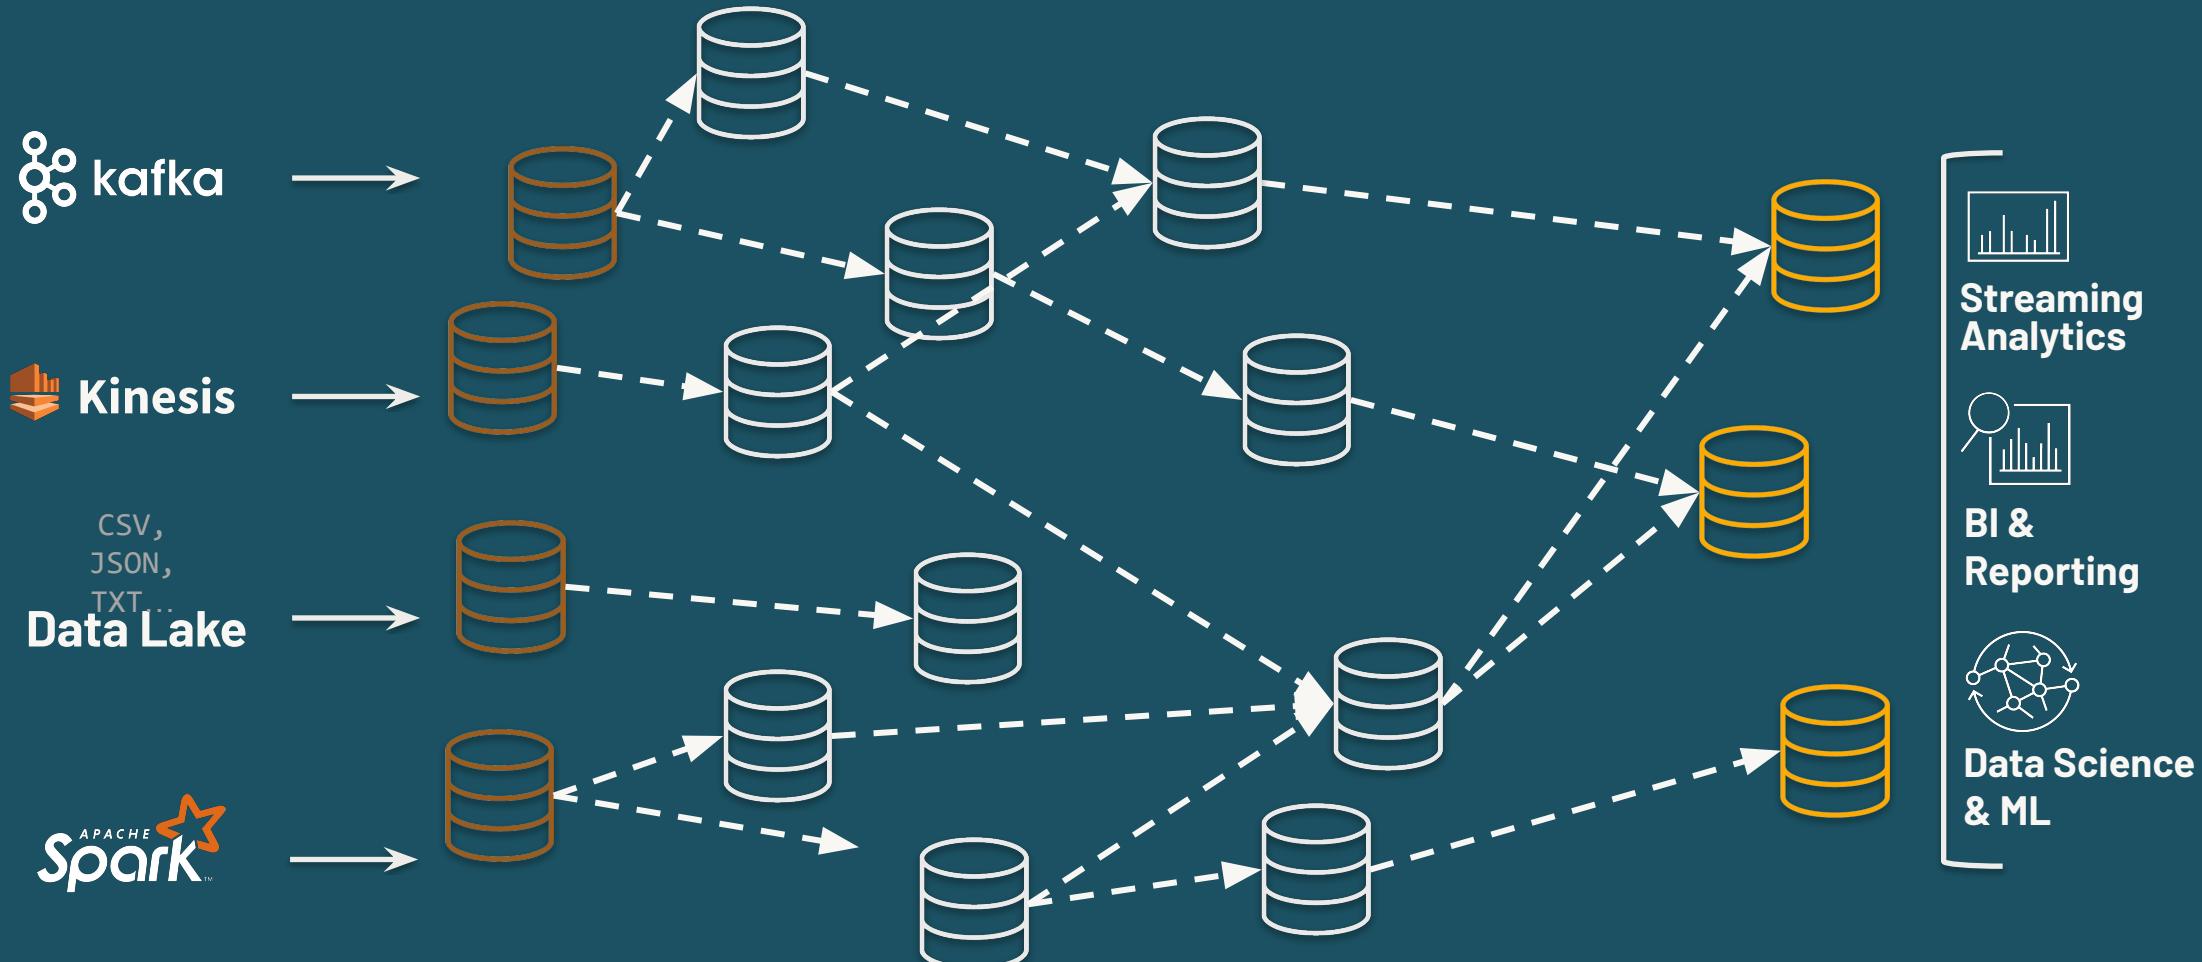
All data professionals need clean, fresh and reliable data.





# But the reality is not so simple

Maintaining data quality and reliability at scale is often **complex and brittle**



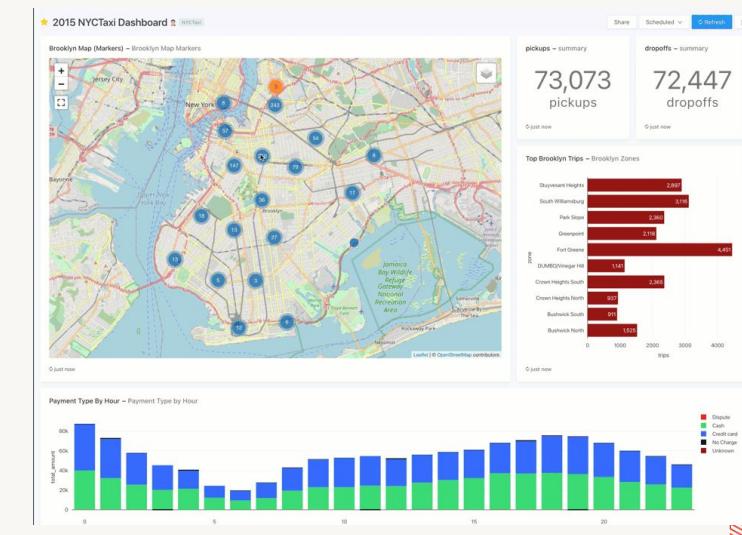
# Life as a data professional...

From: **The CEO** <ali@databricks.com>  
**Subject:** Need an analysis ASAP!  
To: Michael Armbrust  
<michael@databricks.com>

Hey Michael, I need a quick analysis of our net customer retention and how it has changed over the past few quarters. Raw data can be found at  
s3://our-data-bucket/raw\_data/...

```
1 %fs ls /data/sensors
```

	path
1	dbfs:/data/sensors/_SUCCESS
2	dbfs:/data/sensors/_committed_3908896360792309052
3	dbfs:/data/sensors/_started_3908896360792309052
4	dbfs:/data/sensors/part-00000-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
5	dbfs:/data/sensors/part-00001-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
6	dbfs:/data/sensors/part-00002-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
7	dbfs:/data/sensors/part-00003-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
8	dbfs:/data/sensors/part-00004-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3



# Going from query to production

The tedious work required to turn SQL queries into reliable ETL Pipelines

From: **The CEO** <ali@databricks.com>

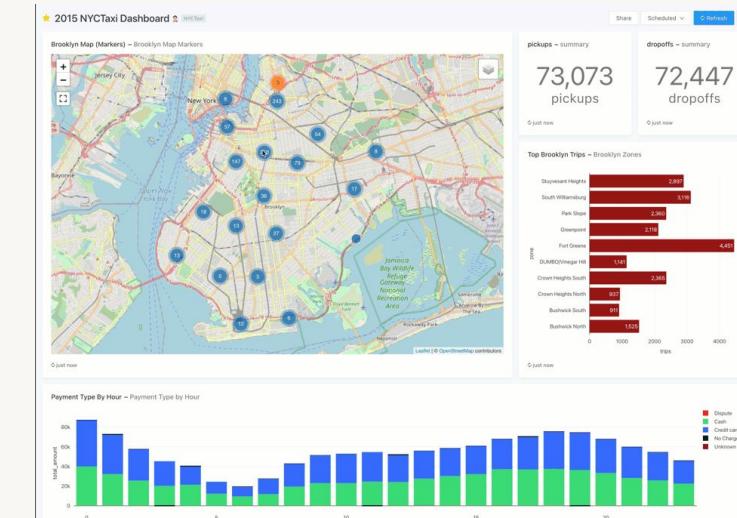
**Subject:** Need an analysis ASAP!

To: Michael Armbrust <michael@databricks.com>

every minute

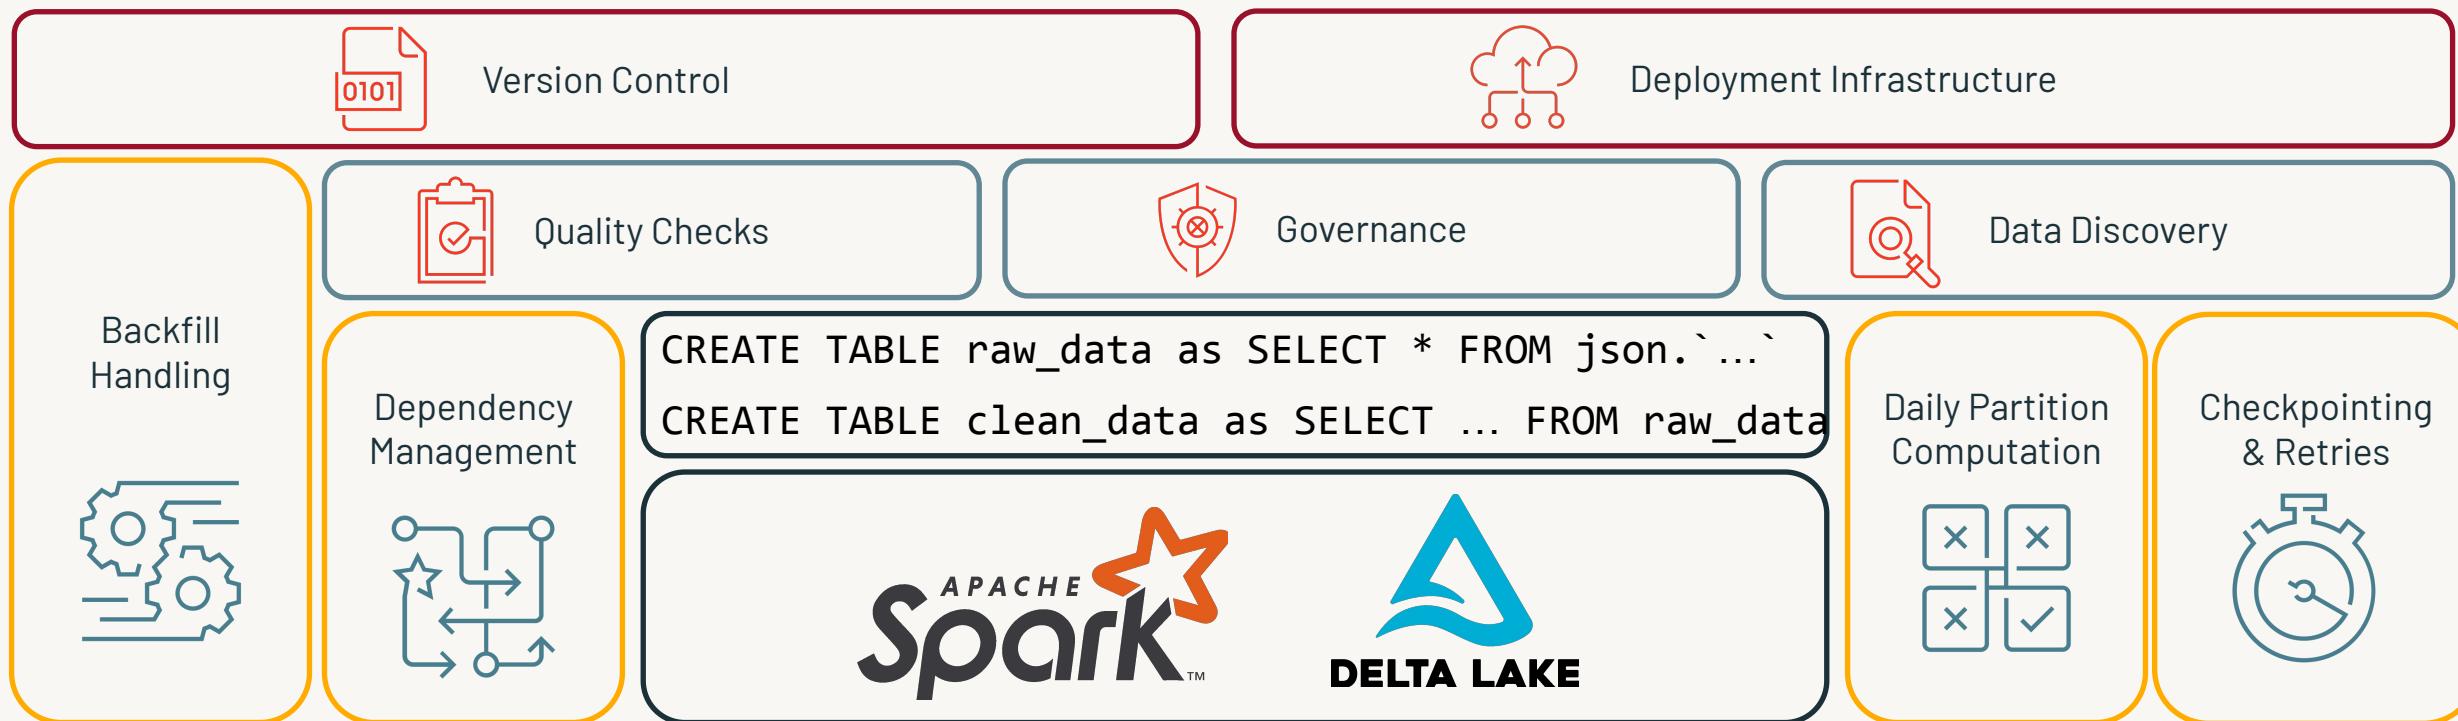
Great report! Can you update it everyday?

```
CREATE TABLE raw_data as SELECT * FROM  
json` `CREATE TABLE clean_data as SELECT ... FROM raw_data
```



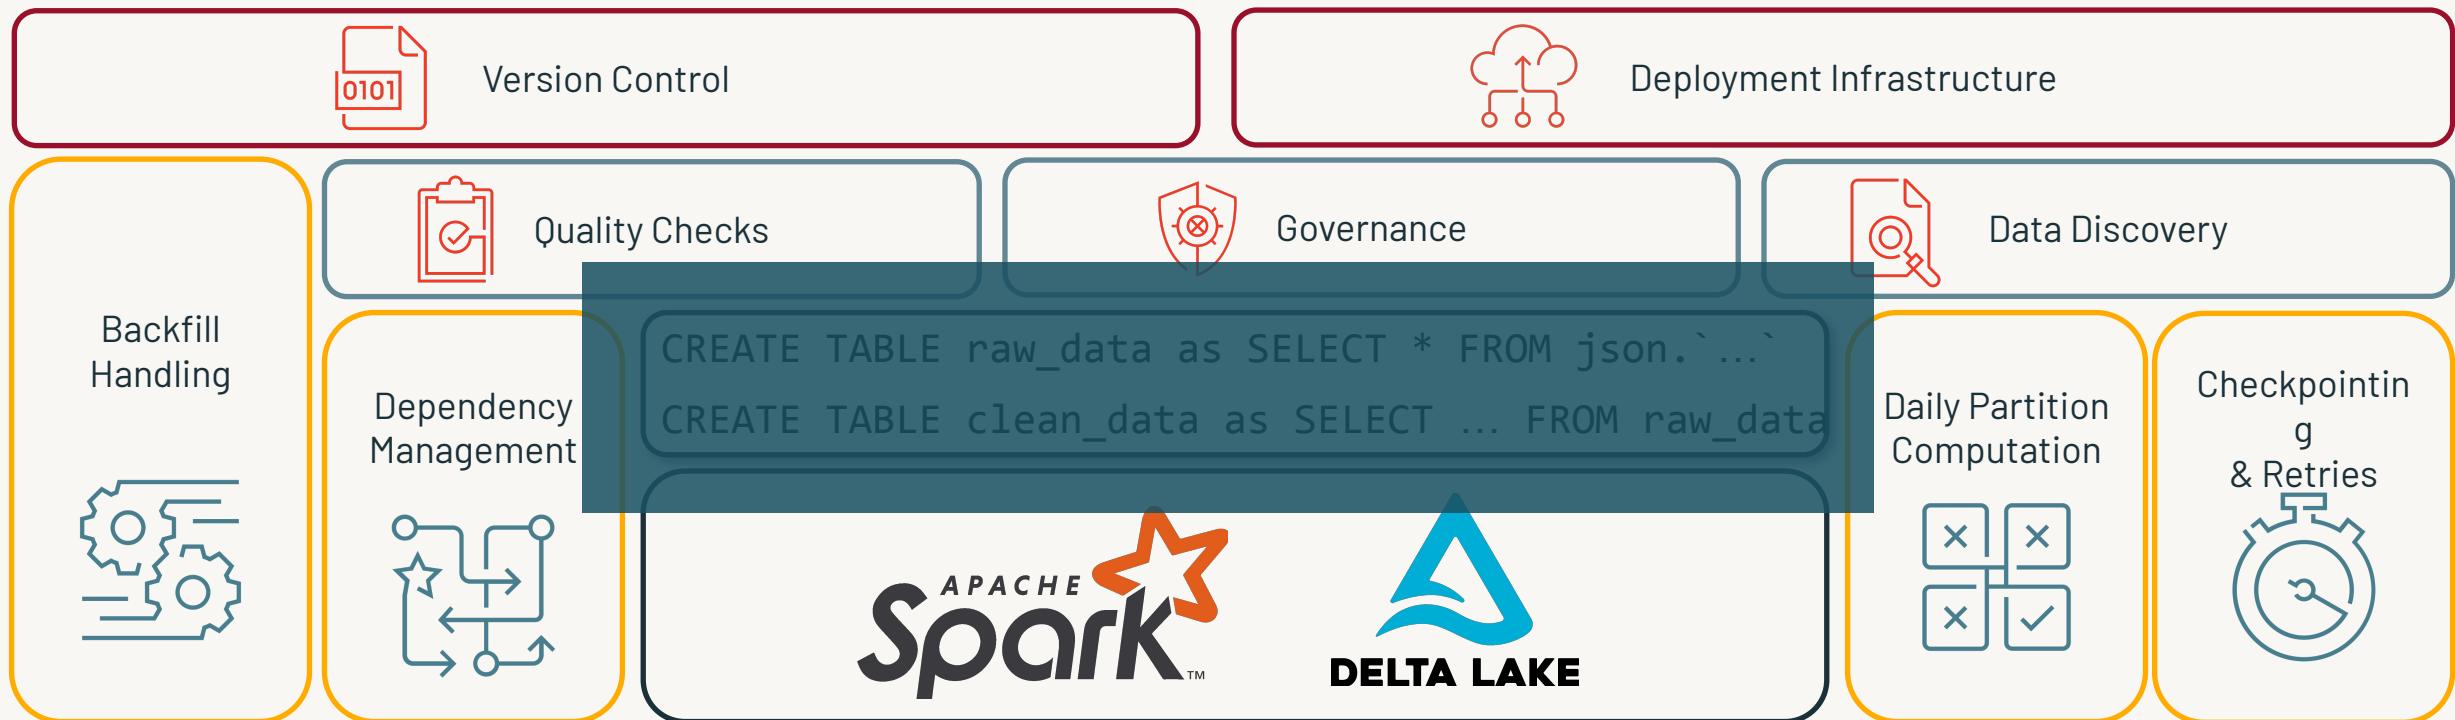
# The **slog** from query to production

The **tedious work** required to turn SQL queries into **reliable ETL Pipelines**



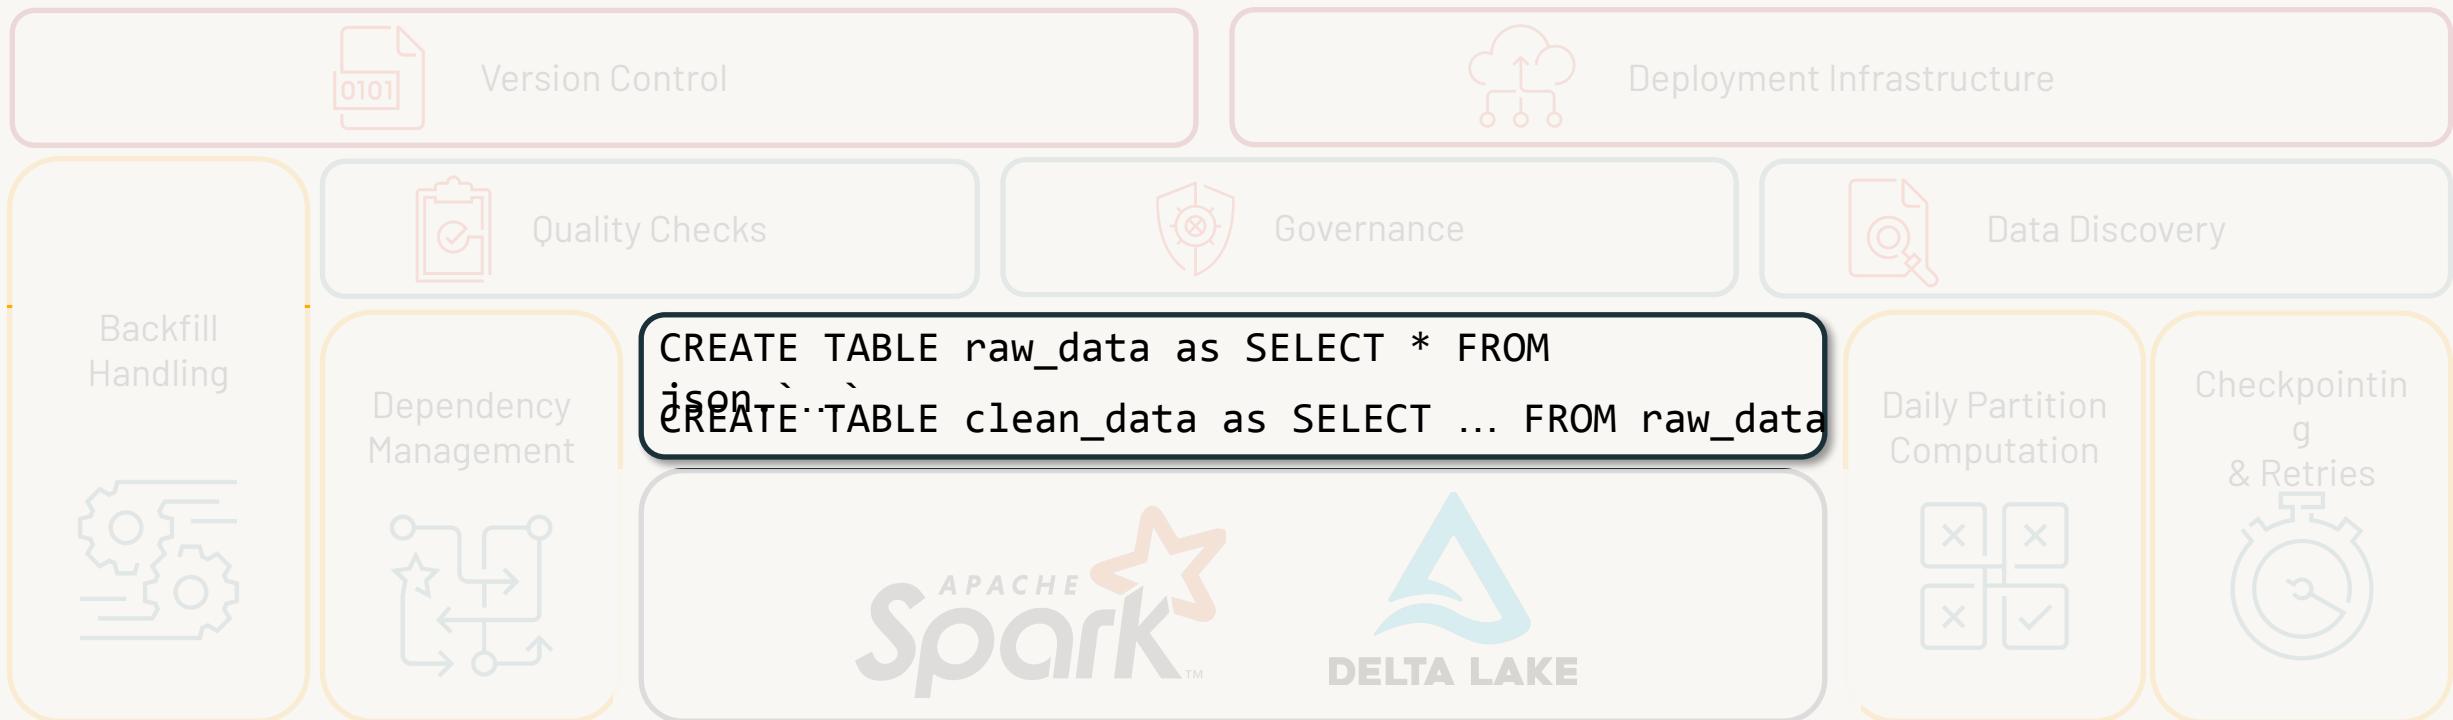
# Operational complexity dominates

Time is spent on **tooling** instead of on **transforming**



# Where you should focus your time

Getting value from data



# Introducing Delta Live Tables

From query to **production pipeline** just by adding **LIVE**.

```
CREATE LIVE TABLE raw_data as SELECT * FROM json.`...`  
CREATE LIVE TABLE clean_data as SELECT ... FROM LIVE.raw_data
```



Repos



Unity Catalog\*



Databricks Workflows

## Delta **Live** Tables

Full Refresh



Dependency Management



Expectations



Incremental Computation\*



Checkpointing & Retries



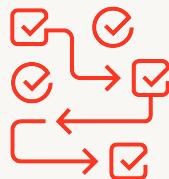
# What is Delta Live Tables?

Modern software engineering for ETL processing

Delta Live Tables (DLT) is the first ETL framework that uses a simple, declarative approach to building reliable data pipelines. DLT automatically manages your infrastructure at scale so data analysts and engineers can spend less time on tooling and focus on getting value from data.



**Accelerate ETL  
Development**



**Automatically manage  
your infrastructure**



**Have confidence  
in your data**



**Simplify batch and  
streaming**

<https://databricks.com/product/delta-live-tables>



# What is a LIVE TABLE?

# What is a Live Table?

Live Tables are materialized views for the lakehouse.

A **live table** is:

- Defined by a SQL query
- Created and kept up-to-date by a pipeline

```
LIVE  
CREATE OR REPLACE TABLE report  
AS SELECT sum(profit)  
FROM prod.sales
```

Live tables provides tools to:

- Manage dependencies
- Control quality
- Automate operations
- Simplify collaboration
- Save costs
- Reduce latency

# What is a Streaming Live Table?

Based on Spark™ Structured Streaming

A **streaming live table** is “stateful”:

- Ensures exactly-once processing of input rows
- Inputs are only read once

- **Streaming Live tables** compute results over append-only streams such as Kafka, Kinesis, or Auto Loader (files on cloud storage)
- Streaming live tables allow you to **reduce costs and latency** by avoiding reprocessing of old data.

```
CREATE STREAMING LIVE TABLE report
AS SELECT sum(profit)
FROM cloud_files(prod.sales)
```

# How do I use DLT?

# Creating Your First Live Table Pipeline

SQL to DLT in three easy steps...

## Write create live table

- Table definitions are written (**but not run**) in notebooks
- Databricks Repos allow you to **version control** your table definitions.

```
1 CREATE LIVE TABLE daily_stats  
2 AS SELECT sum(rev) - sum(costs) AS profits  
3 FROM prod_data.transactions  
4 GROUP BY day
```

## Create a pipeline

- A Pipeline picks **one or more notebooks** of table definitions, as well as any **configuration** required.



Delta Live Tables

## Click start

- DLT will **create or update** all the tables in the pipelines.



# Development vs Production

Fast iteration or enterprise grade reliability

## Development Mode

- Reuses a **long-running cluster** running for **fast iteration**.
- **No retries** on errors enabling **faster debugging**.

## Production Mode

- **Cuts costs** by **turning off clusters** as soon as they are done (within 5 minutes)
- **Escalating retries**, including cluster restarts, **ensure reliability** in the face of transient issues.

In the Pipelines UI:



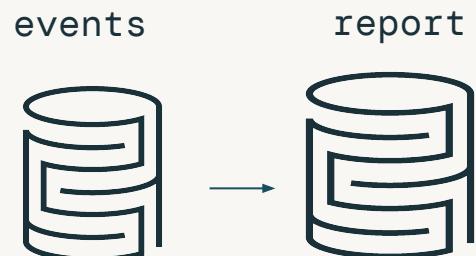
# What if I have many tables?

# Declare LIVE Dependencies

Using the **LIVE virtual schema**.

```
CREATE LIVE TABLE events  
AS SELECT ... FROM prod.raw_data
```

```
CREATE LIVE TABLE report  
AS SELECT ... FROM LIVE.events
```

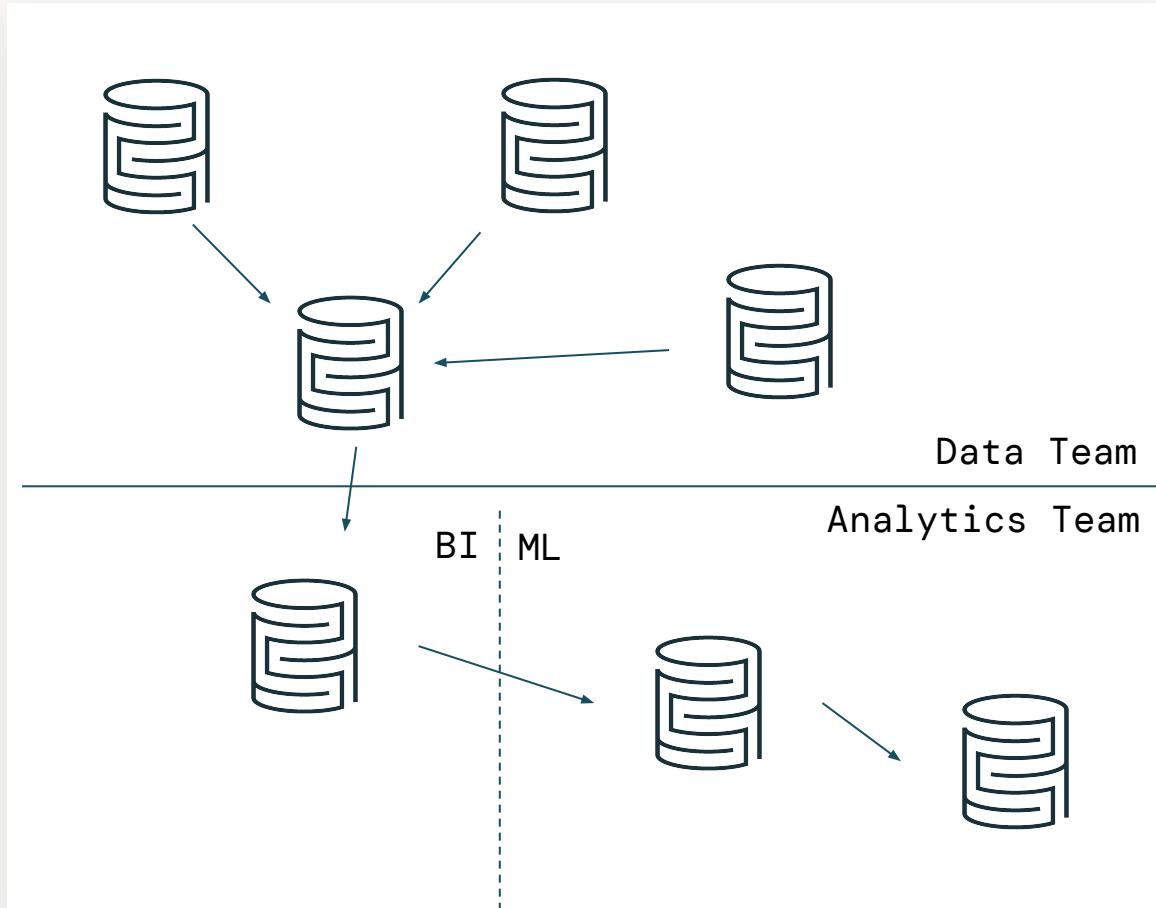


- Dependencies owned by **other producers** are just read from the **catalog or spark data source** as normal.
- **LIVE dependencies**, from the **same pipeline**, are read from the **LIVE schema**.
- DLT **detects LIVE dependencies** and executes all operations in **correct order**.
- DLT handles **parallelism** and captures the **lineage** of the data.

# Choosing pipeline boundaries

Break up pipelines at **natural external divisions**.

- Larger pipelines can utilize cluster resources more efficiently **reducing TCO**.
- Use **more than one pipeline**:
  - At **team boundaries**
  - At helpful **application-specific** boundaries.
- Tables do not necessarily need to have inter-dependencies
- Advantages of co-locating Tables in a single pipeline:
  - Reduces overall Pipeline Count
  - Reduce Orchestration Complexity
  - Provides observability & lineage



# Use autoscaling to reduce costs

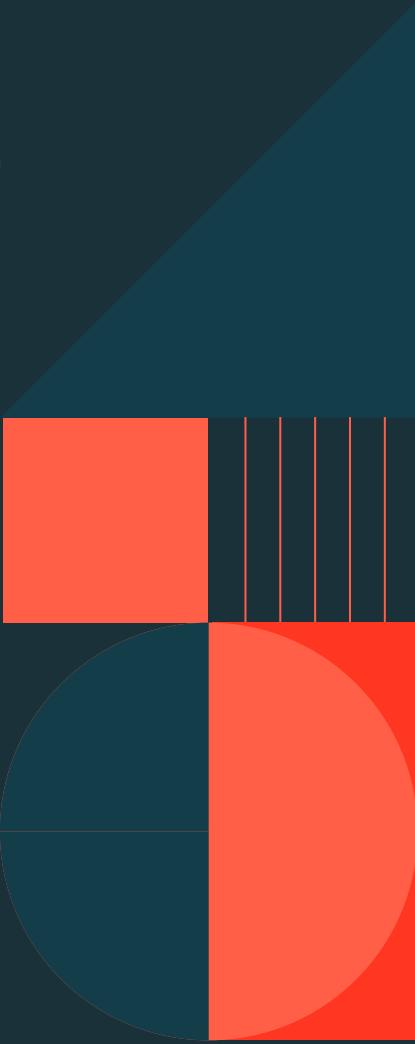
Autoscaling lets you set a budget and leave the tuning to us

How to think about setting the various configurations in **production**:

- Min Nodes
  - Leave default
- Max Nodes
  - Maximum you would pay for a **timely answer**
- Instance Type
  - unset – leaves more **freedom** for DLT to pick the right instance type\*

This strategy leverages **cloud elasticity**:

- Autoscaling will only **add more resources** if it thinks they will **speed things up**.
- **Resources are freed** as soon as they are no longer useful. **Clusters are shut down** as soon as all updates are complete.
- A larger maximum size makes you **robust to changes** in input volumes to avoid missing SLAs.

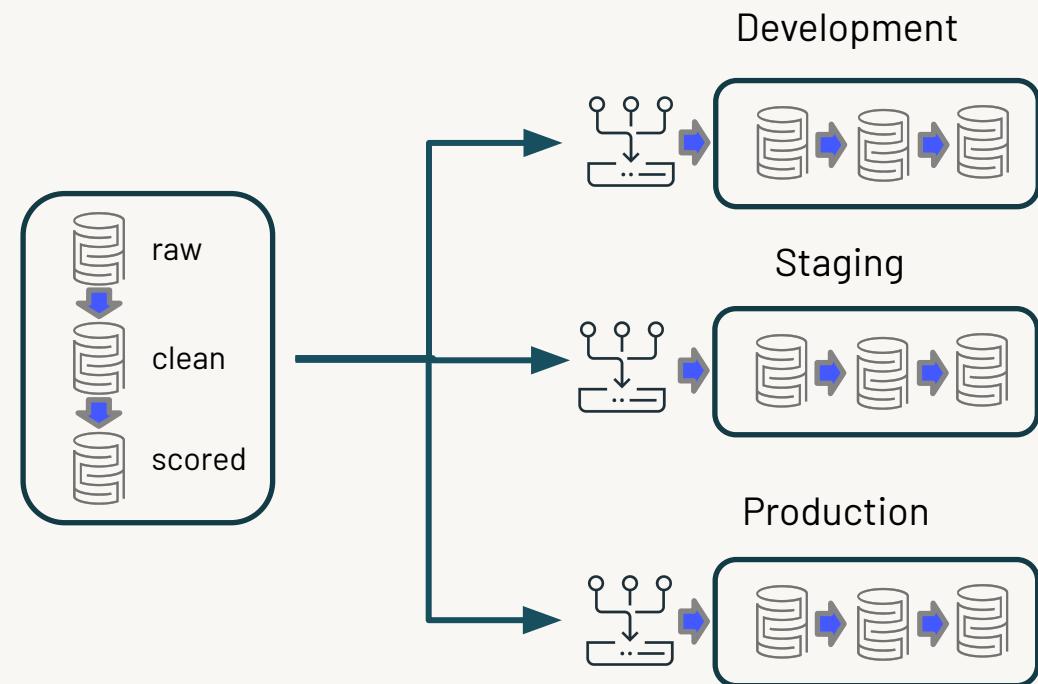


# How can I develop safely with my coworkers?

# Continuous Integration and Deployment (CI/CD) requires testing

DLT allows the same version of your code to read and write to different, isolated environments

- DLT accomplishes this with **targets** and the **LIVE** schema.
- A **target** is a pipeline-level parameter that defines a schema to publish table(s) to
- The **LIVE** schema is a **reserved keyword** that locates your dependencies in the **target**.

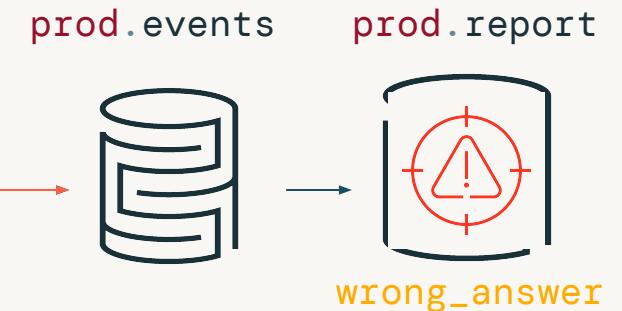
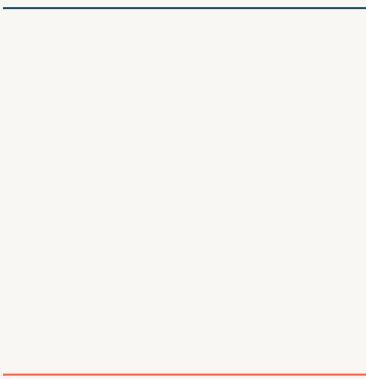


# Pitfall: hard-code sources & destinations

Problem: Hard coding the source & destination makes it impossible to test changes outside of production, breaking CI/CD

## Production Job

```
CREATE OR REPLACE TABLE prod.report  
AS SELECT right_answer FROM prod.events
```



## Struggling Developer

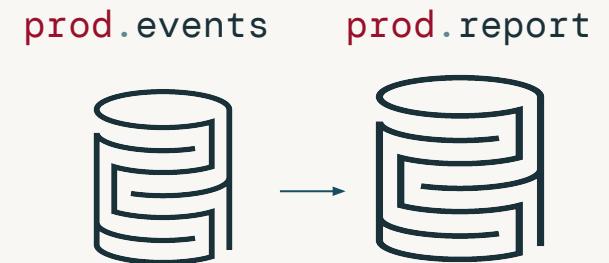
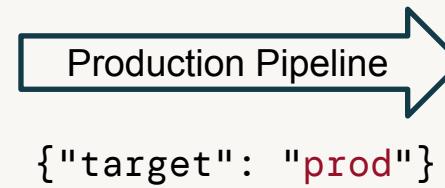
```
CREATE OR REPLACE TABLE prod.report  
AS SELECT wrong_answer FROM prod.events
```

# Best Practice: Keep environments isolated

DLT automatically locates data using the pipeline's "target"

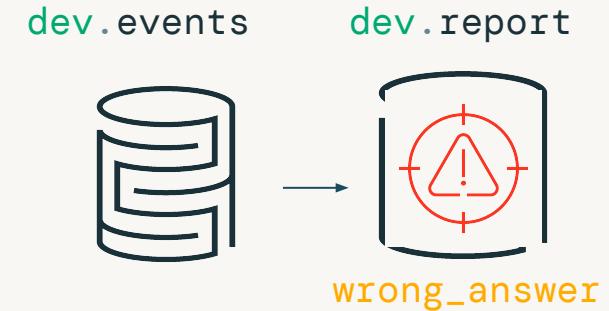
## Production Job

```
CREATE LIVE TABLE report  
AS SELECT right_answer FROM LIVE.events
```



## Struggling Developer

```
CREATE LIVE TABLE report  
AS SELECT wrong_answer FROM LIVE.events
```

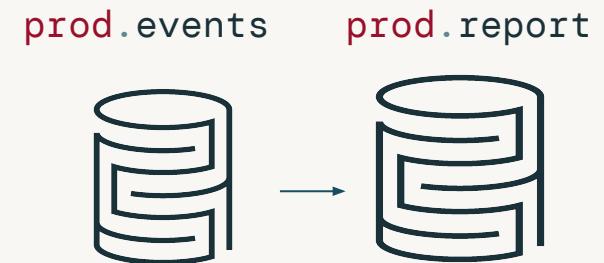
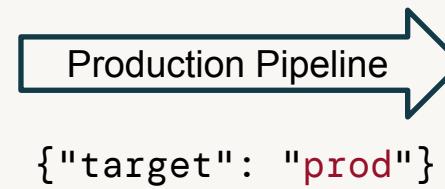


# Best Practice: Use version control

One code base that is promoted to production **after validation**

## Production Job

```
CREATE LIVE TABLE report  
AS SELECT right_answer FROM LIVE.events
```



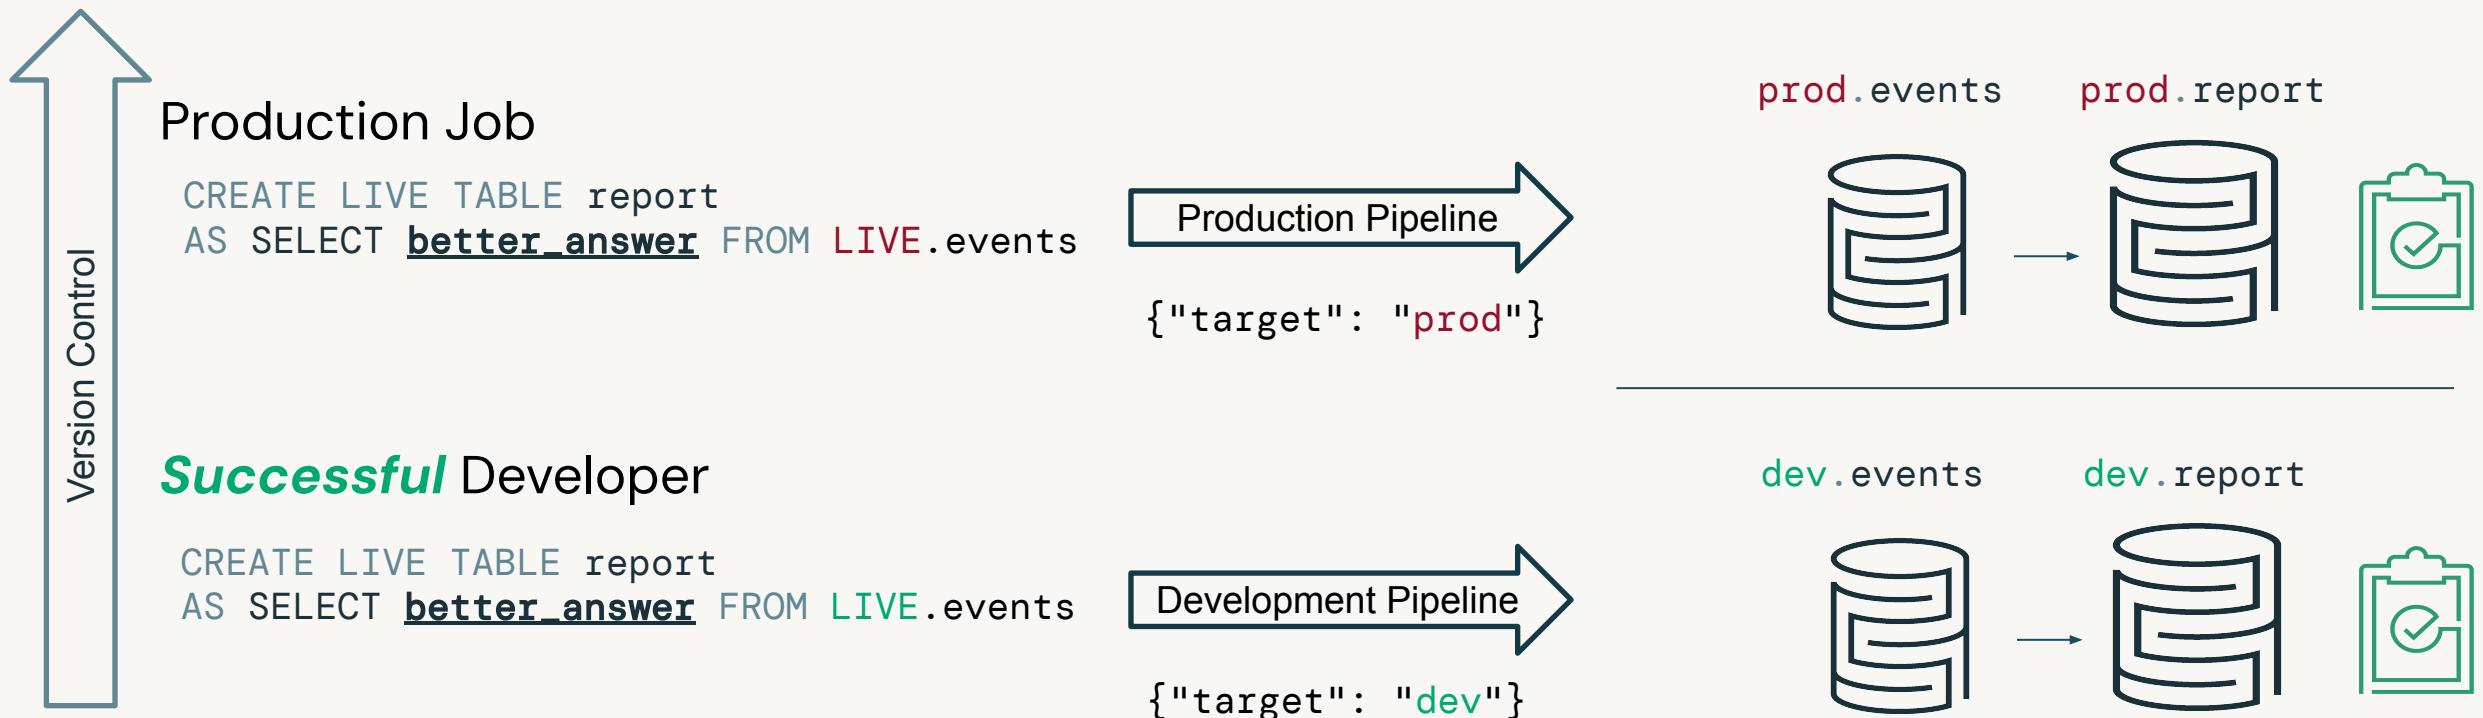
## **Successful** Developer

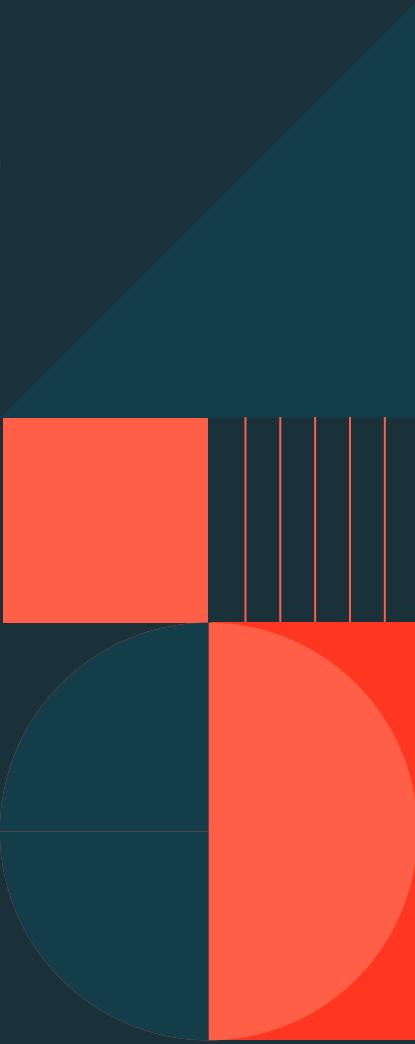
```
CREATE LIVE TABLE report  
AS SELECT better_answer FROM LIVE.events
```



# Best Practice: Use version control

One code base that is promoted to production **after validation**





# How do I know my results are correct?

# Ensure correctness with Expectations

Expectations are tests that ensure data quality in production

```
CONSTRAINT valid_timestamp  
EXPECT (timestamp > '2012-01-01')  
ON VIOLATION DROP
```

```
@dlt.expect_or_drop(  
    "valid_timestamp",  
    col("timestamp") > '2012-01-01')
```

Expectations are true/false expressions that are used to validate each row during processing.

DLT offers flexible policies on how to handle records that violate expectations:

- Track number of bad records
- Drop bad records
- Abort processing for a single bad record

# Expectations using the power of SQL

Use SQL aggregates and joins to perform complex validations

-- Make sure a primary key is always unique.

```
CREATE LIVE TABLE report_pk_tests(  
    CONSTRAINT unique_pk EXPECT (num_entries = 1)  
)  
AS SELECT pk, count(*) as num_entries  
FROM LIVE.report  
GROUP BY pk
```

# Expectations using the power of SQL

Use SQL aggregates and **joins** to perform complex validations

- Compare records between two tables,
- or validate foreign key constraints.

```
CREATE LIVE TABLE report_compare_tests(  
    CONSTRAINT no_missing EXPECT (r.key IS NOT NULL)  
)  
  
AS SELECT * FROM LIVE.validation_copy v  
LEFT OUTER JOIN LIVE.report r ON v.key = r.key
```

# What if I need more than SQL?

# Using Python

## Write advanced DataFrame code and UDFs

```
import dlt

@dlt.table
def report():
    df = spark.table("LIVE.events")
    return df.select(...)
```

---

```
spark.udf.register(
    "complex_function",
    lambda x: ...)
```

- Add `@dlt.table` to any function that returns a DataFrame.
- Mix and match SQL/Python notebooks in a single pipeline.
- But a single notebook must be all Python or all SQL (cannot mix/match at notebook level)
- DataFrames can be constructed using PySpark, Koalas, or SQL strings.

# Installing libraries with pip

pip is a package installer for python

- Access a wide range of python libraries using the %pip magic command.
- Python libraries installed in one notebook are available to all notebooks in the pipeline.
- Order of notebooks & dependencies does not matter.
- Compile external python files as wheels and %pip install:  

```
%pip install  
/path/to/my_package.whl
```

# Metaprogramming in python

Create dynamic pipelines by programmatically creating tables

- Automate ingestion of many similar inputs
- Integrate with external metadata system like schema registries or other catalogs

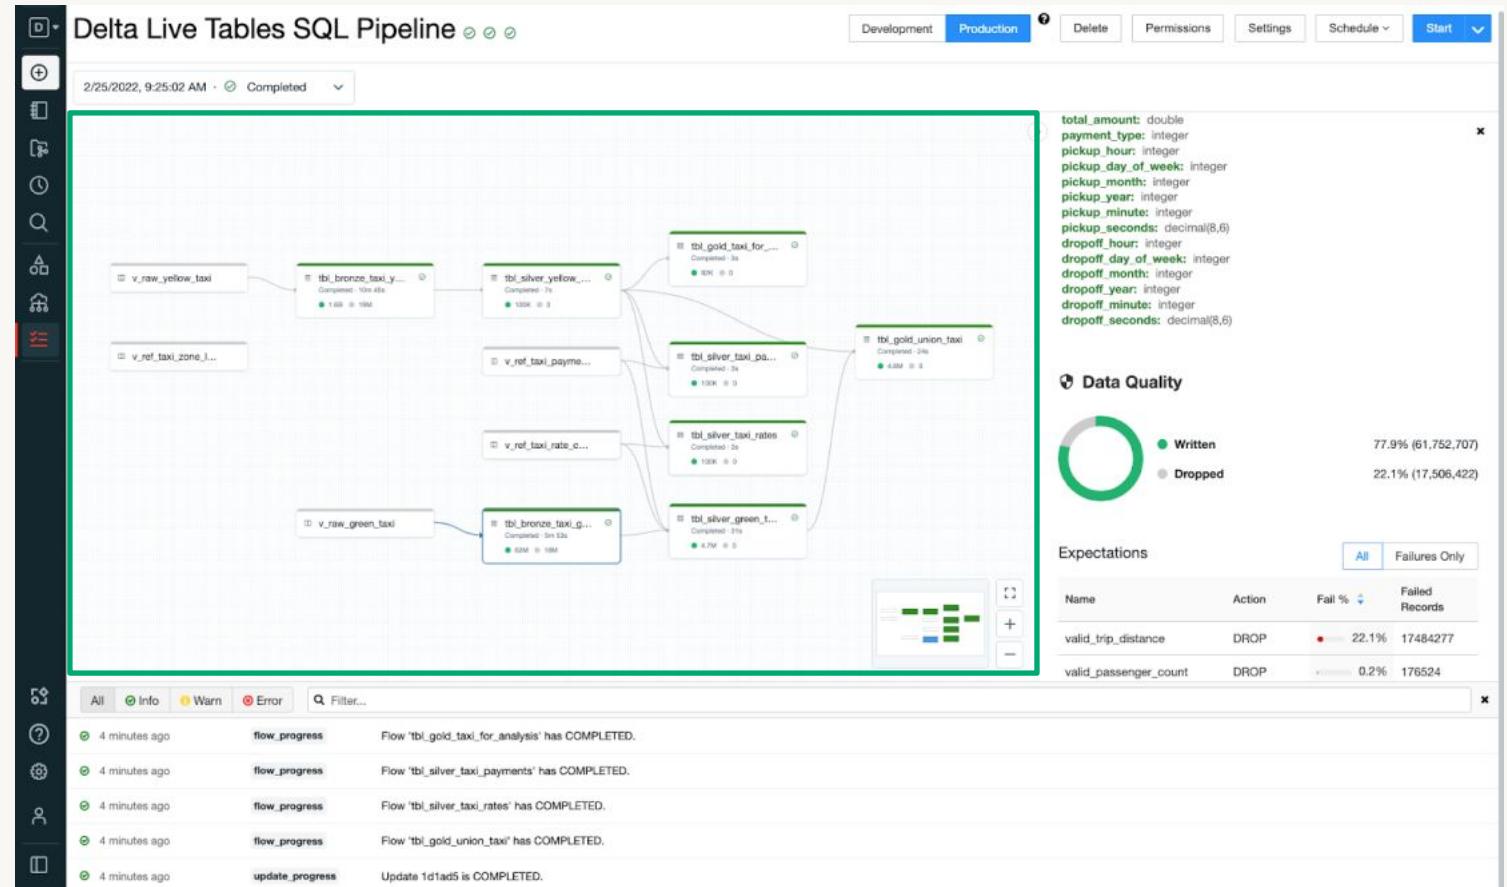
```
for t in tables:  
    @dlt.table(name=t)  
def report():  
    df = spark.table(t)  
    return df.select(...)
```

# What about operations?

# Pipelines UI

A one stop shop for ETL debugging and operations

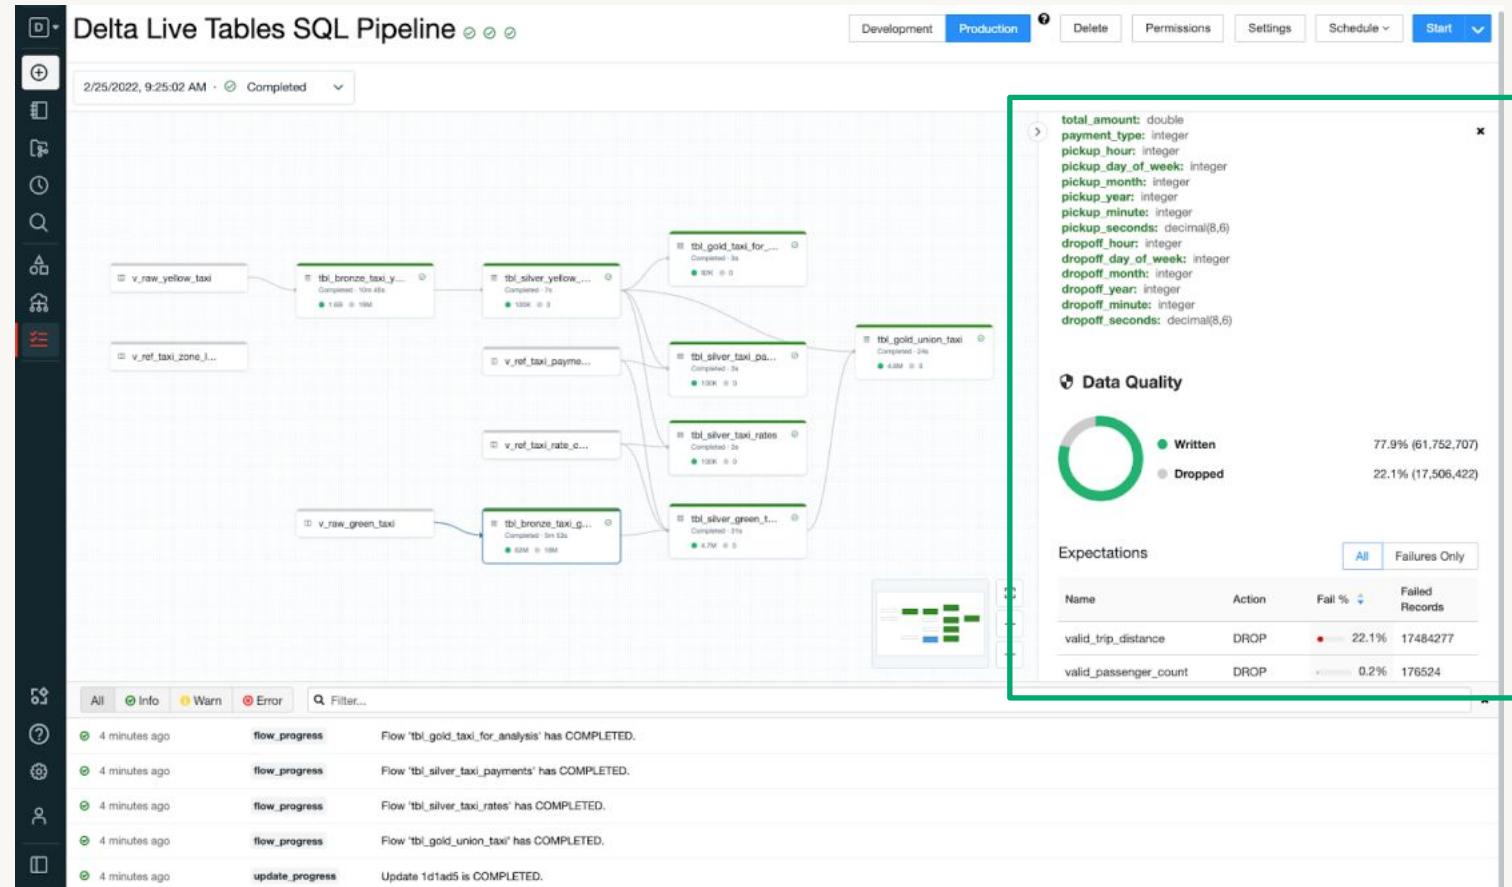
- Visualize data flows between tables



# Pipelines UI

A one stop shop for ETL debugging and operations

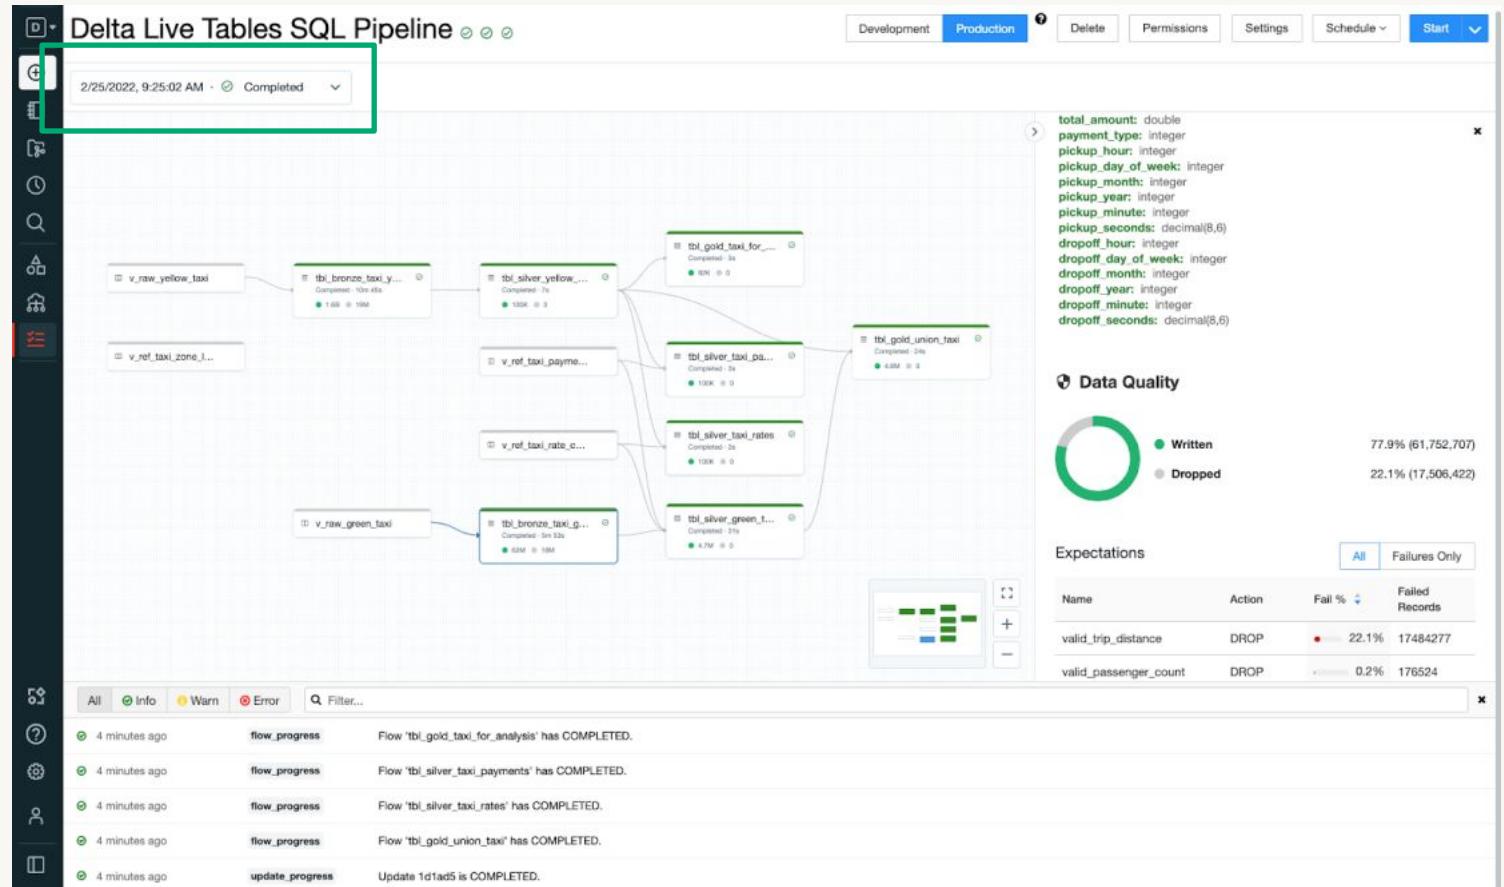
- Visualize data flows between tables
- Discover metadata and quality of each table



# Pipelines UI

A one stop shop for ETL debugging and operations

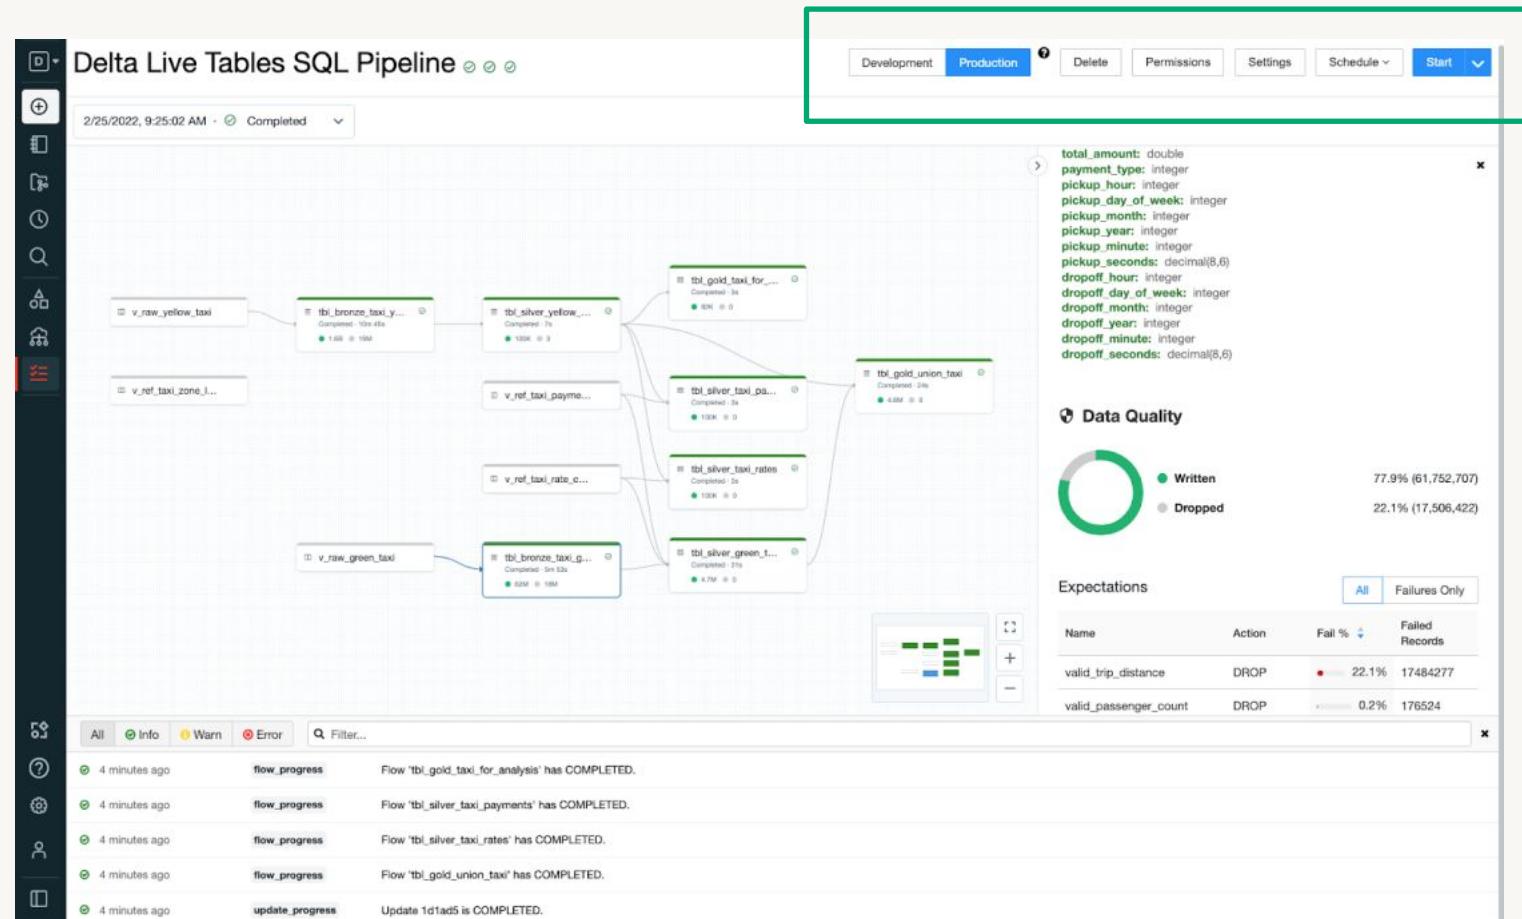
- Visualize data flows between tables
- Discover metadata and quality of each table
- Access to historical updates



# Pipelines UI

A one stop shop for ETL debugging and operations

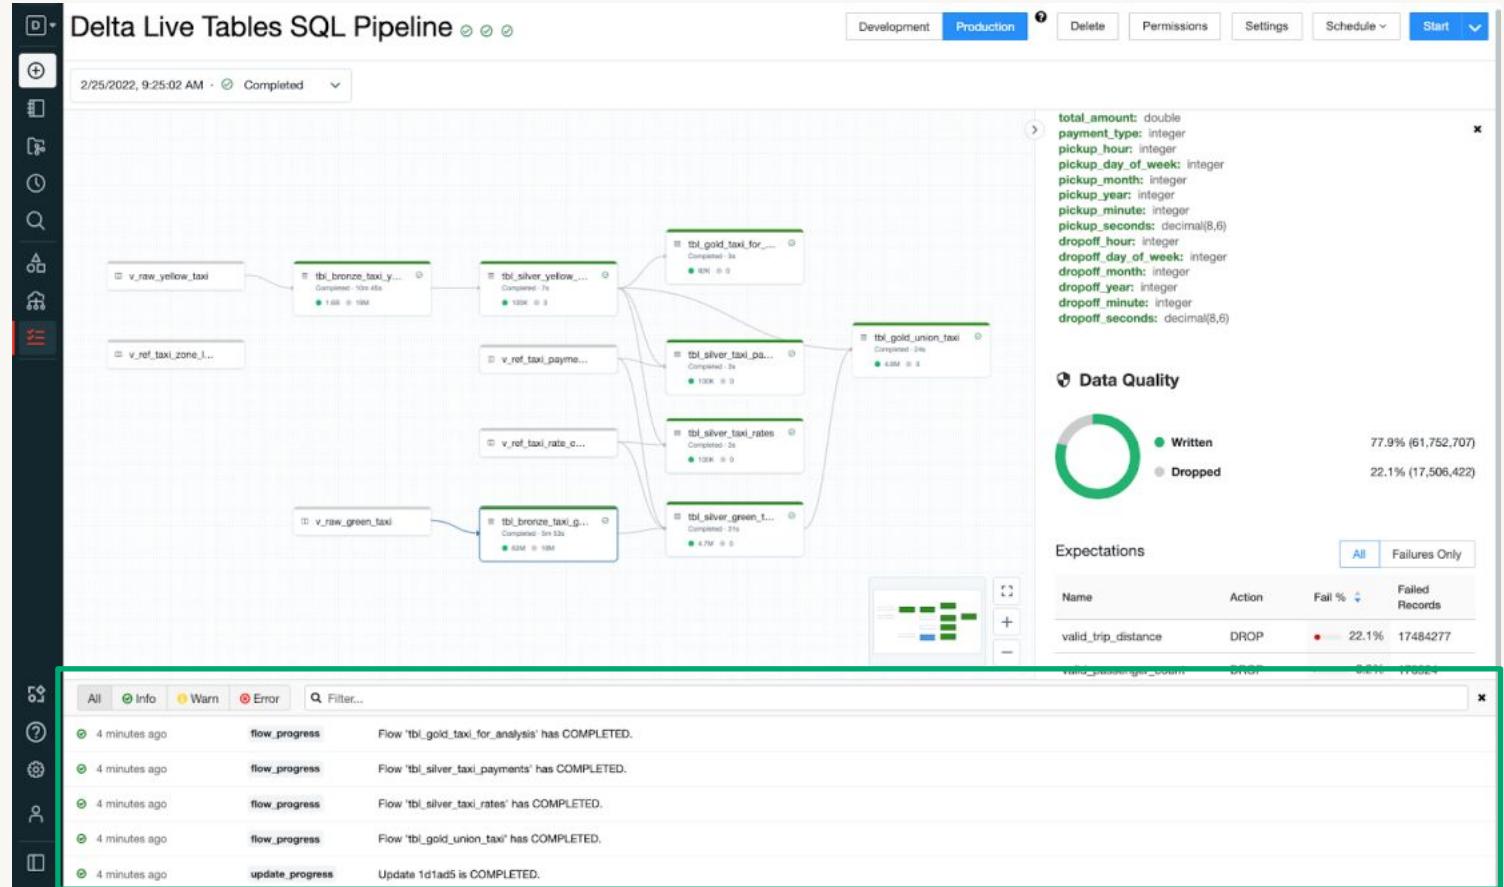
- Visualize data flows between tables
- Discover metadata and quality of each table
- Access to historical updates
- Control operations



# Pipelines UI

A one stop shop for ETL debugging and operations

- Visualize data flows between tables
- Discover metadata and quality of each table
- Access to historical updates
- Control operations
- Dive deep into events



# The Event Log

The event log **automatically records all pipelines operations.**

## Operational Statistics

Time and current status, for all operations

Pipeline and cluster configurations

Row counts

## Provenance

Table schemas, definitions, and declared properties

Table-level lineage

Query plans used to update tables

## Data Quality

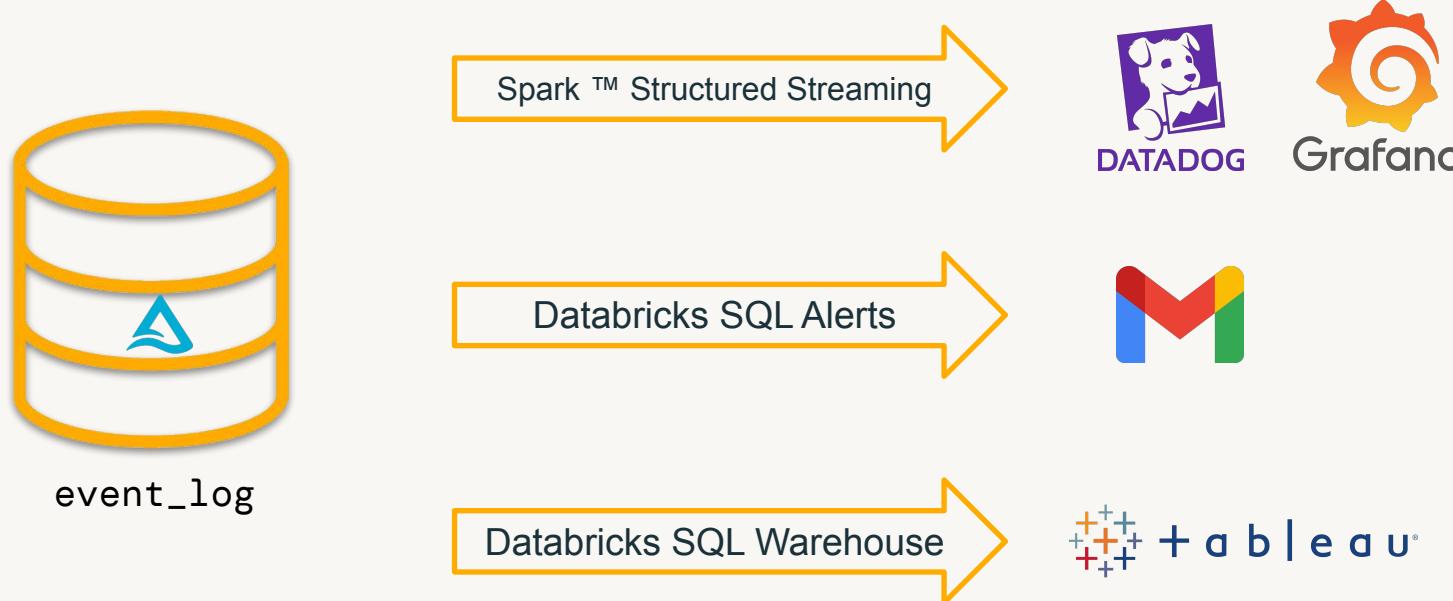
Expectation pass / failure / drop statistics

Input/Output rows that caused expectation failures

# Best Practice: Integrate using the event log

Use the information in the event log with your **existing operational tools**.

The event log is just a delta table created for each pipeline.



# How does DLT handle failures?

# DLT Automates Failure Recovery

Transient issues are handled by built-in retry logic

- DLT uses **escalating retries** to balance speed with reliability
  - Retry the individual transaction
  - Next, restart the cluster in case its in a bad state
  - If DBR was upgraded since the last success, automatically try the old version to detect regressions.
- Transactionally
  - Operations on tables are atomic.
  - Updates to different tables in a pipeline are not atomic. Best effort to update as many as possible.



# When should I use streaming?

# What is Spark™ Structured Streaming?

Basis for **Streaming Live Tables**. Runs **queries** on continually arriving data.

**Computation Model:** Input is an ever-growing **append-only table**

- Files uploaded to cloud storage
- Message busses like kafka, kinesis, or eventhub
- Delta tables with delta.appendOnly=true
- Transaction logs of other databases

Rather than **wait until all data has arrived**, structured streaming can produce **results on demand**.

- **Lower latency** by processing less data each update
- **Lower costs** by avoiding redundant work

# Using Spark™ Structured Streaming for ingestion

## Easily ingest files from cloud storage as they are uploaded

```
CREATE STREAMING LIVE TABLE raw_data  
AS SELECT *  
FROM cloud_files("/data", "json")
```

This example creates a table with all the json data stored in "/data":

- `cloud_files` keeps track of which files have been read to **avoid duplication and wasted work**
- Supports both listing and notifications for **arbitrary scale**
- Configurable **schema inference** and **schema evolution**

# Using Spark™ Structured Streaming for ingestion

## Easily ingest records from message buses

```
import dlt

@dlt.table
def kafka_data():
    return spark.readStream \
        .option("format", "kafka") \
        .option("subscribe", "events") \
        .load()
```

This example creates a table with all the records published to the Kafka topic “event”.

- The Kafka source + DLT automatically track which partitions / offsets have already been read.
- Any structured streaming source included in DBR can be used with DLT
- Message buses provide the lowest latency for ingesting data

# Using the SQL STREAM() function

## Stream data from any Delta table

```
CREATE STREAMING LIVE TABLE mystream  
AS SELECT *  
FROM STREAM(my_table)
```

Pitfall: `my_table` must be an append-only source.

e.g. it may not:

- be the target of `APPLY CHANGES INTO`
- define an aggregate function
- be a table on which you've executed DML to delete/update a row (see GDPR section)

- `STREAM(my_table)` reads a stream of new records, instead of a snapshot
- Streaming tables must be an append-only table
- Any append-only delta table can be read as a stream (i.e. from the live schema, from the catalog, or just from a path).

# Use Delta for infinite retention

Delta provides cheap, elastic and governable storage for transient sources



Use a **short retention** period to **avoid compliance risks** and reduce costs

CREATE STREAMING  
LIVE TABLE AS ...

```
TBLPROPERTIES (  
    pipelines.reset.allowed=false  
)
```

bronze



CREATE STREAMING  
LIVE TABLE AS ...

Avoid complex transformations that could have bugs or drop important data

Retain **history**  
Easy to perform **GDPR** and other compliance tasks

Setting **pipelines.reset.allowed=false** ensures that downstream computation can be **full-refreshed** without losing data

# Streaming does not always mean expensive

Delta live tables lets you choose how often to update the results.

## Triggered: Manually

**Costs:** lowest

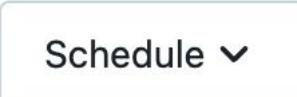
**Latency:** highest



## Triggered: On a schedule using Databricks Jobs

**Costs:** depends on frequency

**Latency:** 10 minutes to months



Every Day at 22 : 14 (UTC-07:00) Pacific Ti...

## Continually

**Costs:** highest

**Latency:** minutes to seconds  
(for some workloads)

Pipeline Mode

Triggered  Continuous

# When should I chain streaming?

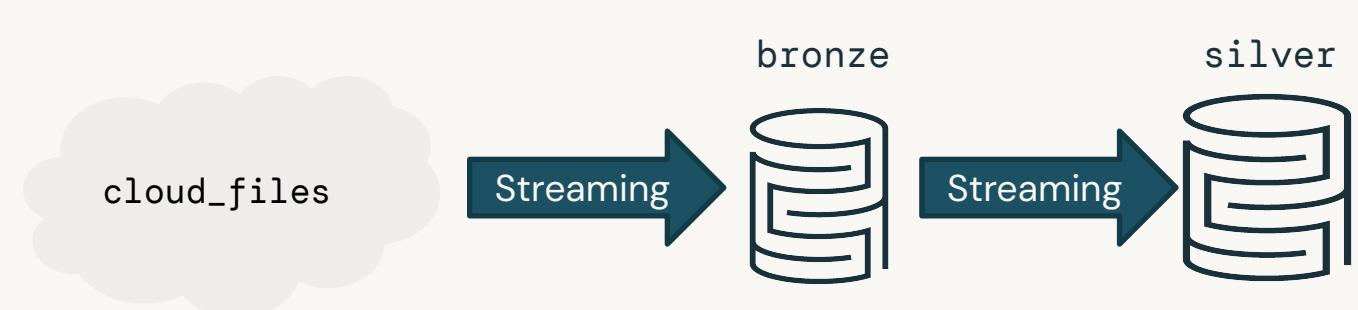
# Multi-hop streaming for cost and latency

Append-only live tables can be both a source and a sink

```
CREATE STREAMING LIVE TABLE bronze  
AS SELECT * FROM cloud_files("/data/", "json")  
  
CREATE STREAMING LIVE TABLE silver  
AS SELECT  
    CAST(ts AS TIMESTAMP) AS timestamp  
    ...  
FROM STREAM(LIVE.bronze)
```

Chain multiple streaming jobs for workloads with:

- Very large data
- Very low latency targets



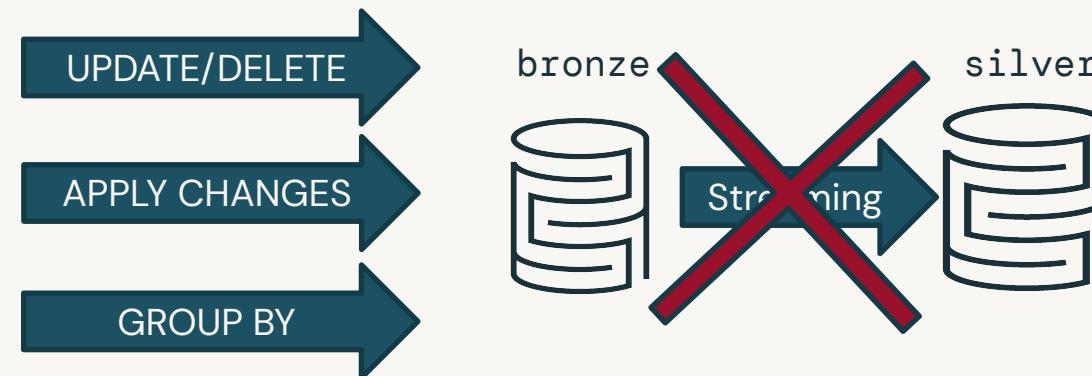
# Pitfall: Updates in streaming pipelines

Structured streaming assumes an append-only input source

Updates to a streaming input table will **break downstream computation**

- Non-insert DML performed on streaming tables
- Streaming tables that compute aggregations without a watermark
- Streaming tables used with **APPLY CHANGES INTO**

Repair the stream  
after one-off DML  
using a full-refresh



# How do I manage streaming state?

# Streaming queries are **stateful**

Each input row is processed only once.

A change to a streaming live table's definition does not recompute results by default:

```
CREATE STREAMING LIVE TABLE raw_data  
AS SELECT a + 1 AS a  a * 2 AS a  
FROM cloud_files("/data", "json")
```

**"/data"**

```
{"a": 1} →  
{"a": 2} →  
{"a": 3} →  
{"a": 4} →
```

**raw\_data**

b
2
3
6
8

# Streaming joins are stateful

Enrich data by joining with an up-to date-snapshot stored in delta

A change to joined table snapshot **does not** recompute results by default:

```
CREATE STREAMING LIVE TABLE raw_data  
AS SELECT *  
FROM cloud_files("/data", "json") f  
JOIN prod.cities c USING id
```

**"/data"**

{"a": 1}

{"a": 1}

**raw\_data**

<b>id</b>	<b>city</b>
1	Bekerly, CA
1	Berkeley, CA

**id**

1

**city**

~~Bekerly, CA~~

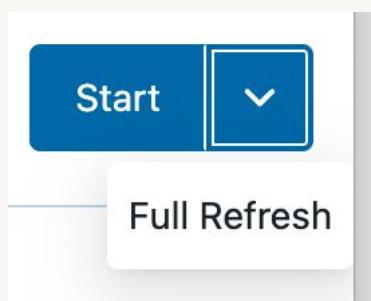
Berkeley, CA

# Clear state using full refresh

Perform backfills after critical changes using full refresh

Full-refresh **clears the table's data and the queries state, reprocessing all the data.**

```
CREATE STREAMING LIVE TABLE raw_data  
AS SELECT a * 2 AS a  
FROM cloud_files("/data", "json")
```



**"/data"**

```
{"a": 1} →  
{"a": 2} →  
{"a": 3} →  
{"a": 4} →
```

**raw\_data**

b
2
3
6
8

After full-refresh

```
{"a": 1} →  
{"a": 2} →  
{"a": 3} →  
{"a": 4} →
```

b
2
4
6
8

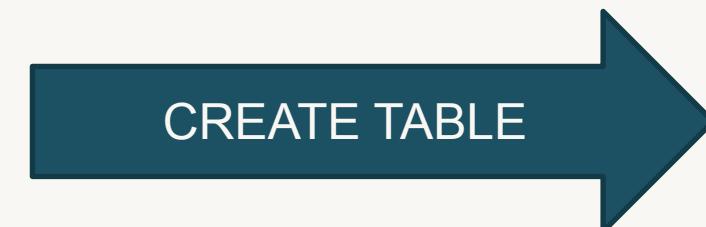
# What about incrementalization?



# Data is always changing

Incrementalization updates derived datasets without recomputing everything

date	city_id
2022-06-01	1
2022-06-01	2
2022-06-02	3



date	city_id	city_name
2022-06-01	1	Berkeley
2022-06-01	2	San Francisco
2022-06-01	3	Denver

date	city_id
2022-06-01	1
2022-06-01	2
2022-06-02	3



date	city_id	city_name
2022-06-01	1	Berkeley
2022-06-01	2	San Francisco
2022-06-01	3	Denver



# Demo

<https://www.databricks.com/resources/demos/tutorials/lakehouse-platform/full-delta-live-table-pipeline>



# Thank you

