



# Data + AI Professional Workshop Series

Data Engineering Optimization  
Best Practices

---

Irfan Elahi – Specialist Solutions Architect



# Housekeeping

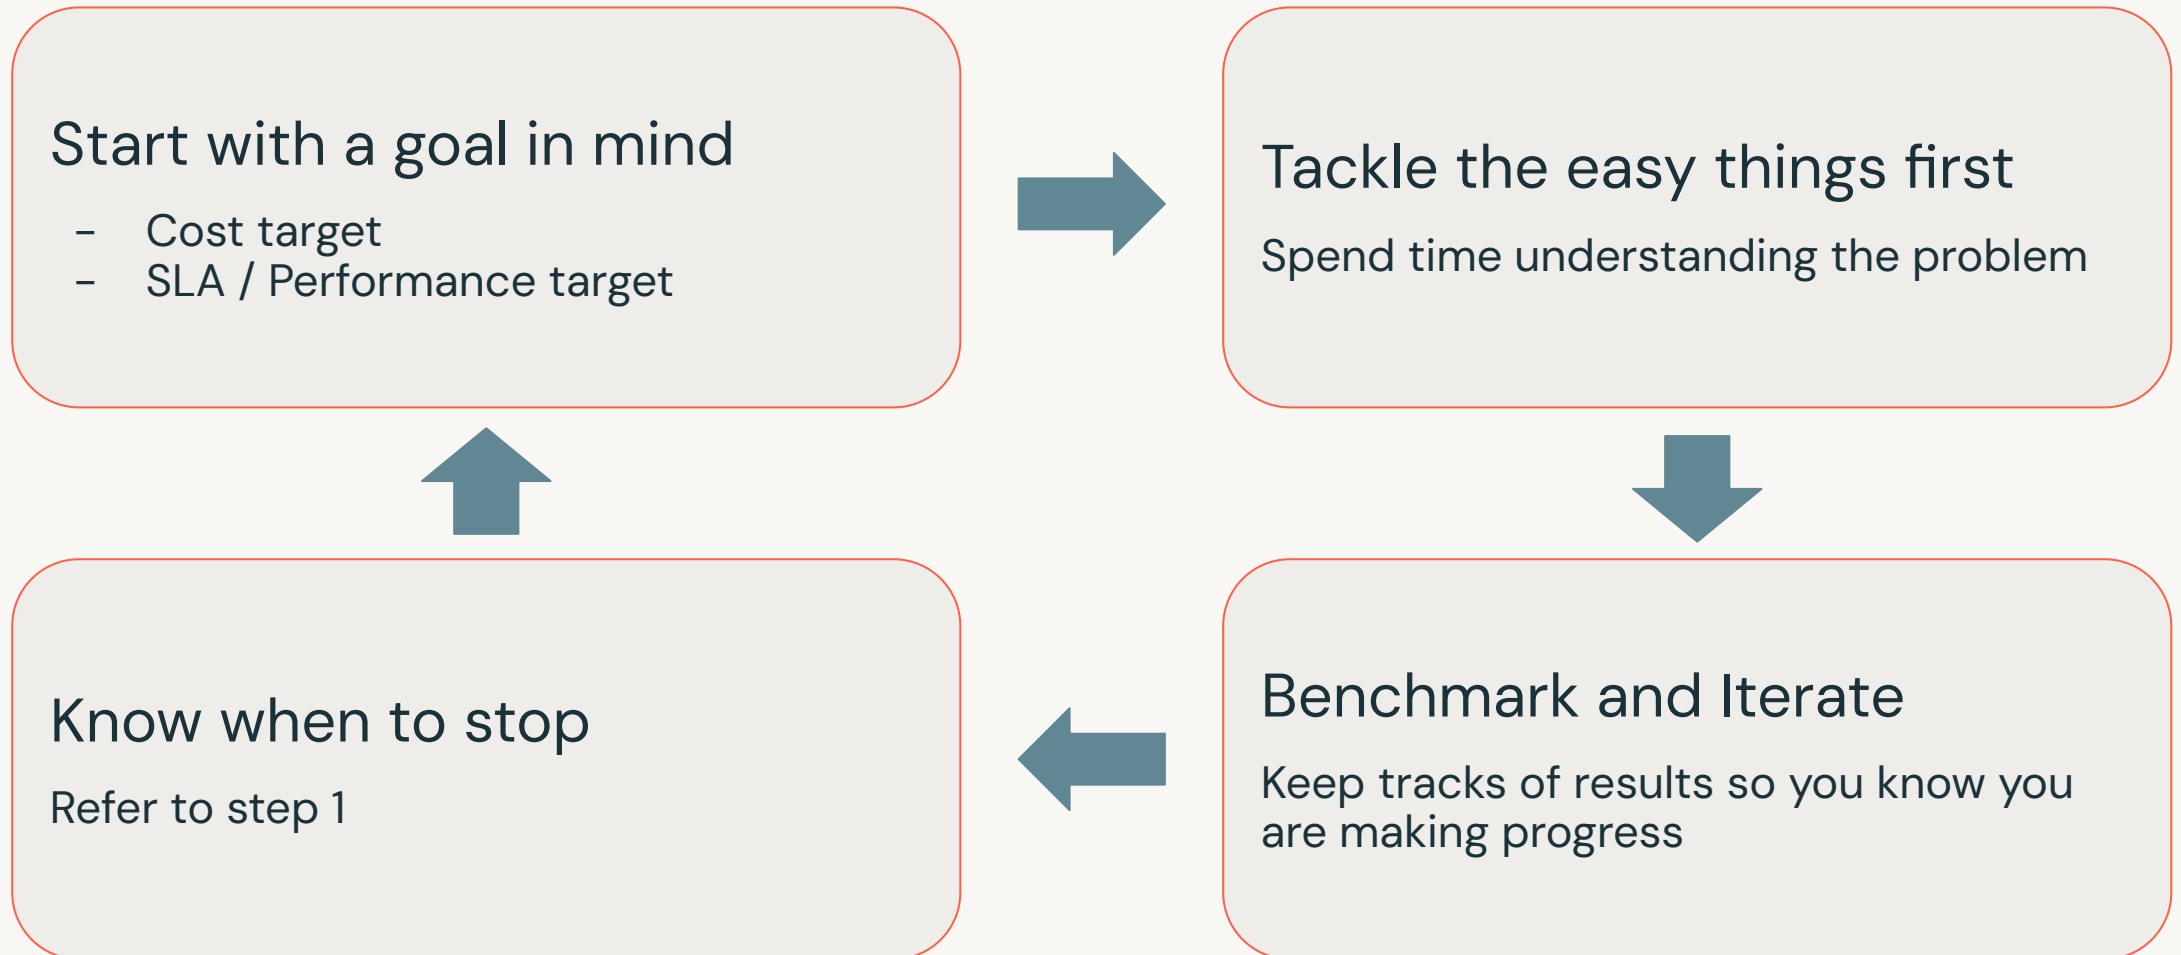
- This presentation will be recorded and we will share these materials after the session.
- There are no hands-on components so you only need something to take notes.
- Use the Q&A function to ask questions.
- Please fill out the survey at the end of the session so we can improve our future sessions.




# The Optimization Mindset



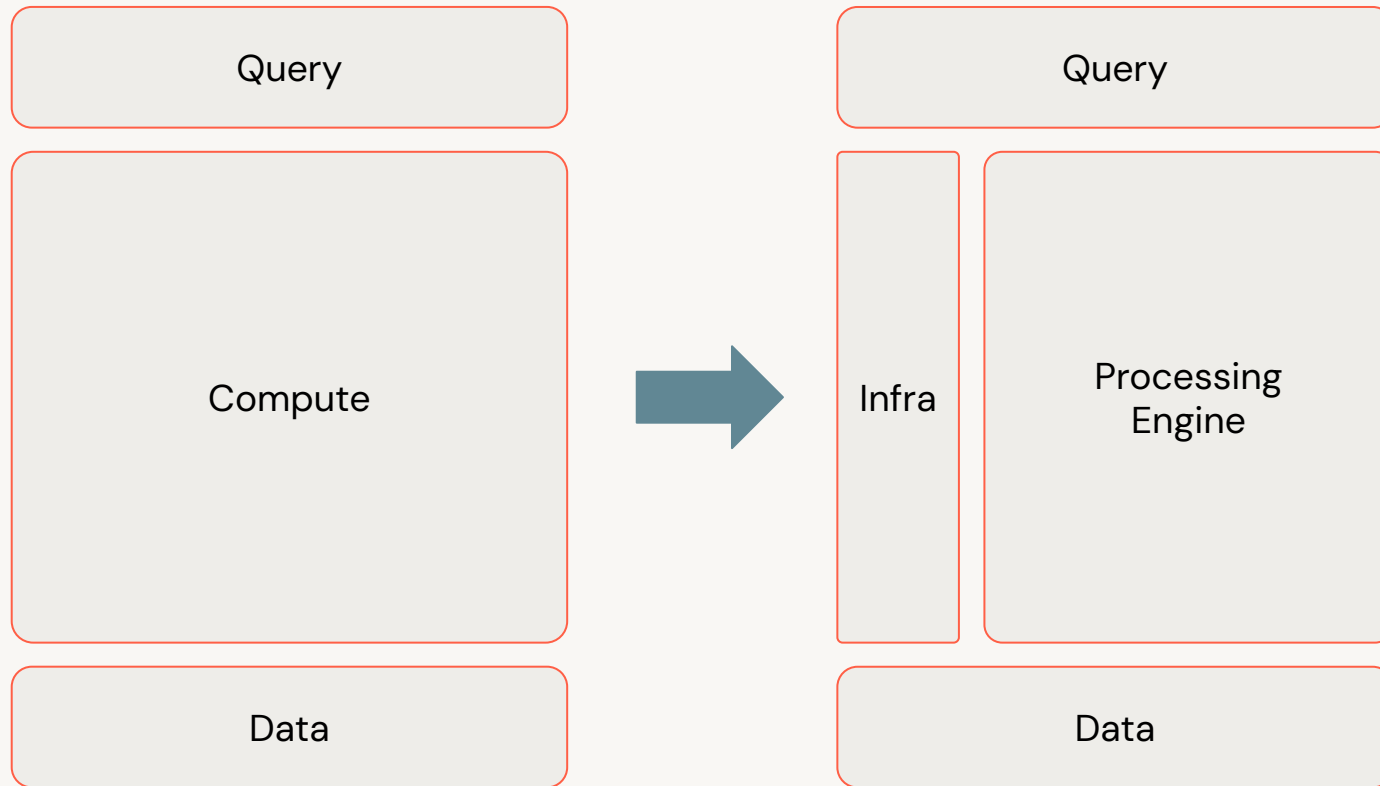
# Optimize Only When Necessary



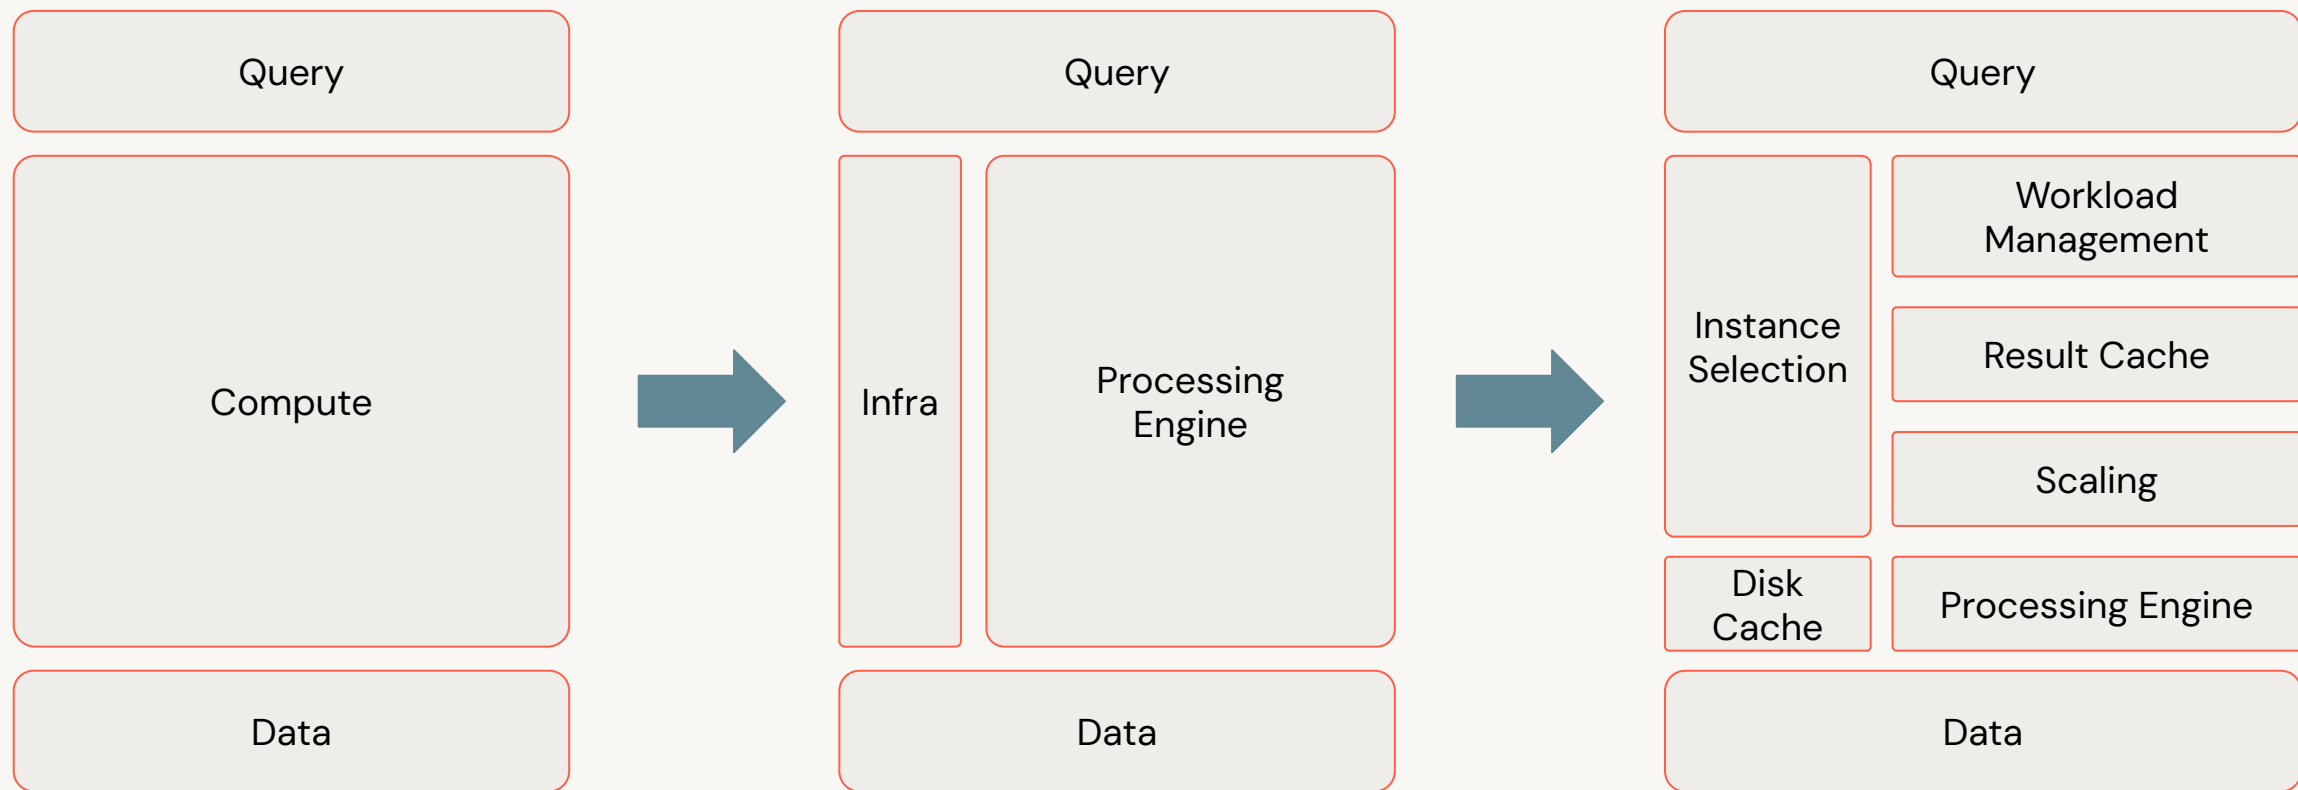
# Understanding the journey of a query



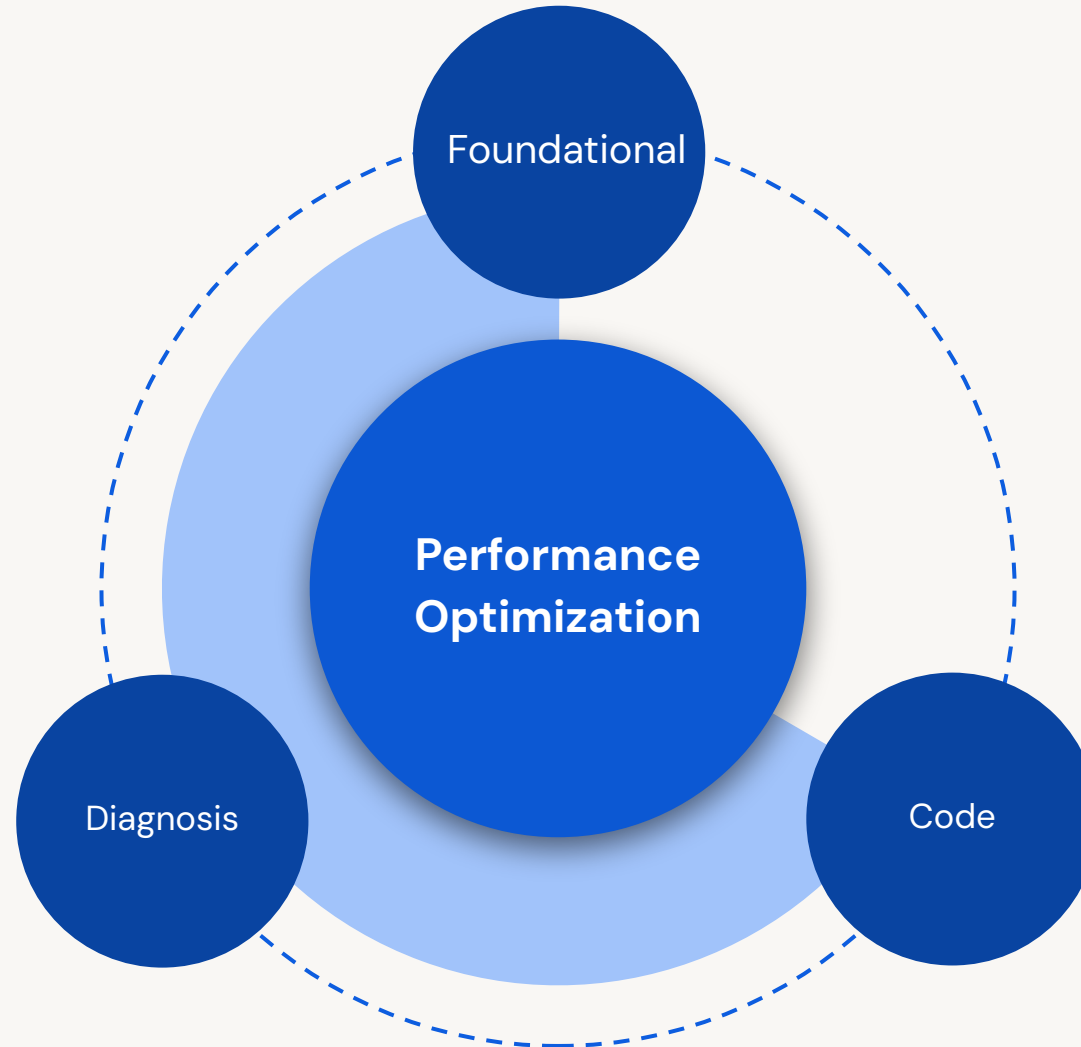
# Understanding the journey of a query



# Understanding the journey of a query



# Performance Optimization – Framework

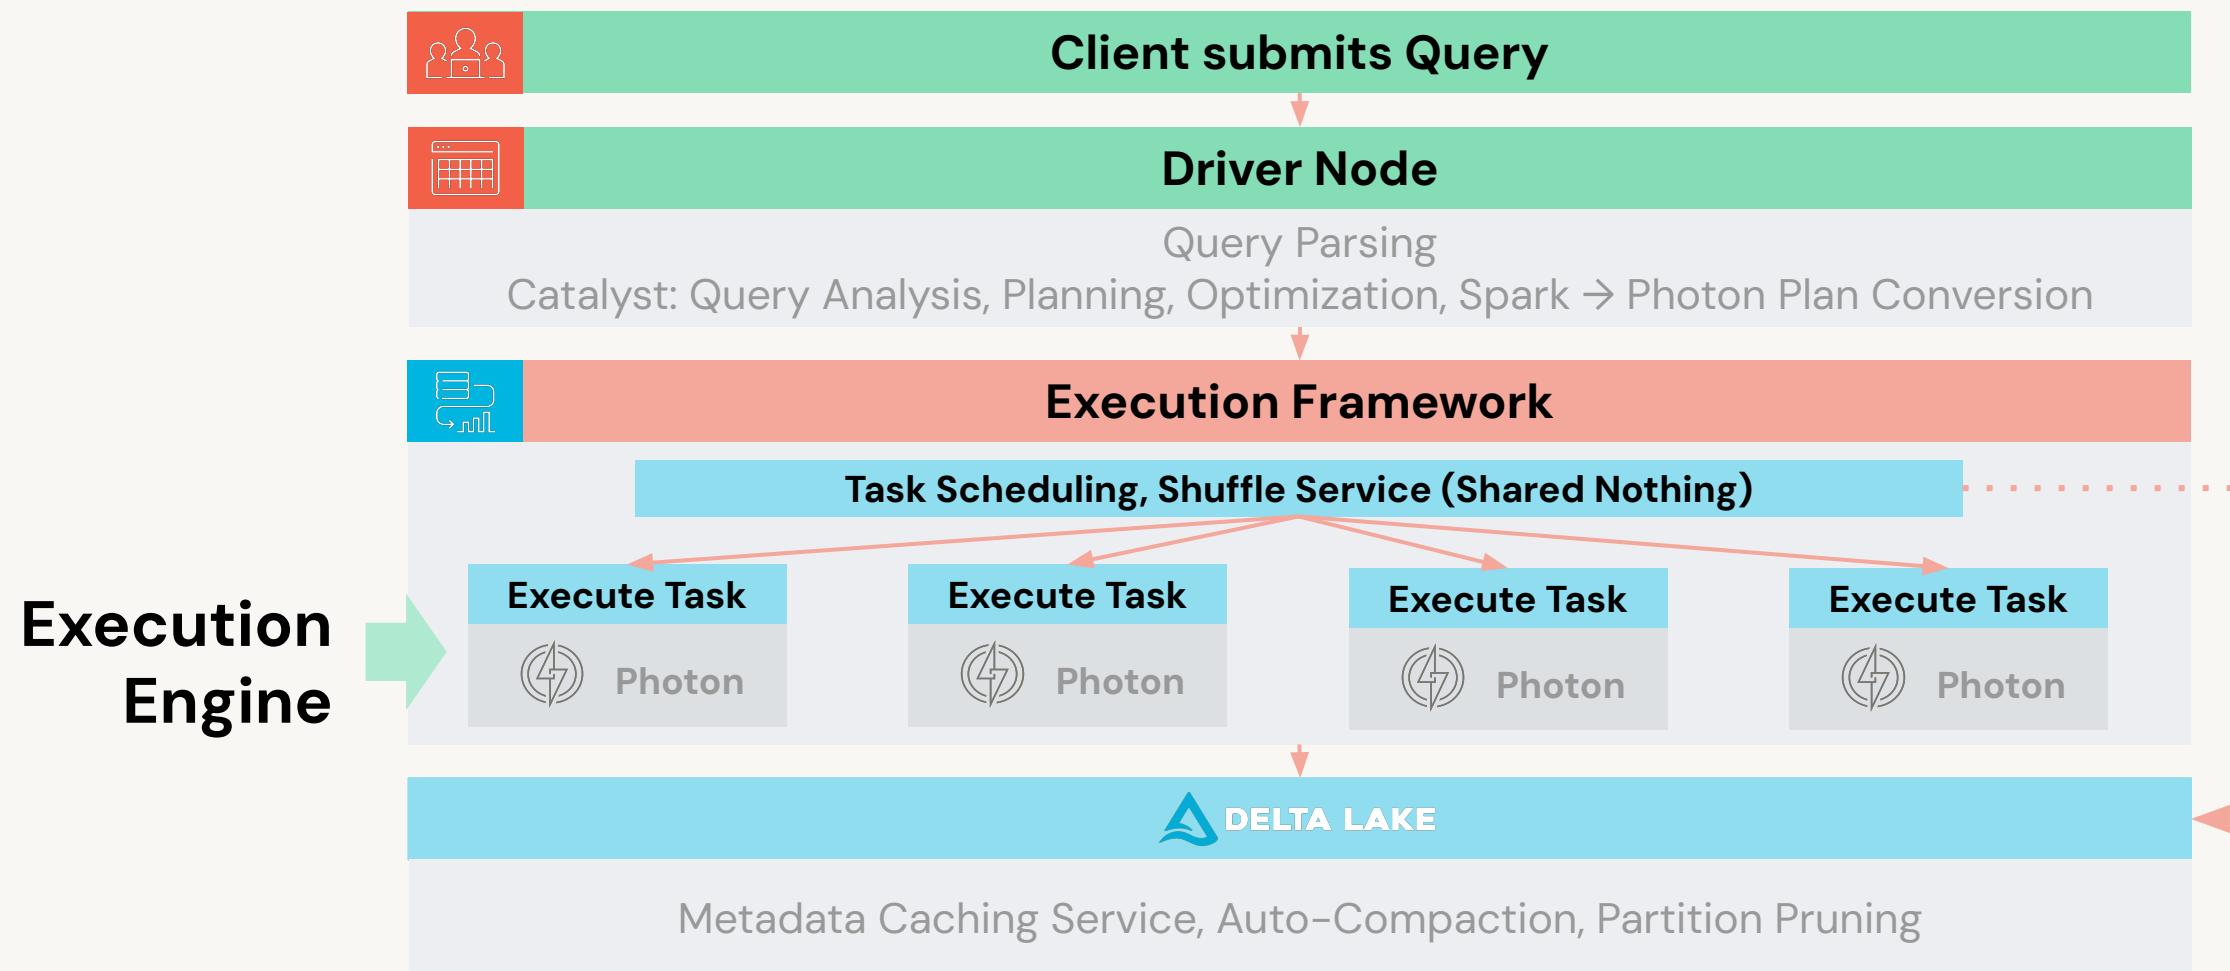




# Foundational – Compute



# Photon – Speeding up Data Processing



# Where and When to Use Photon



Batch

COPY INTO

Auto Loader

Structured Streaming

Delta Live Tables

Ingestion

- ✓ Delta Lake
- ✓ Parquet
- ✓ JSON
- ✓ CSV
- ✓ AVRO
- ✓ XML
- ✓ Binary
- ✓ JDBC/ODBC

- ✓ DataFrame
- ✓ SQL
- ✓ SQL UDFs
- ⚠ Pandas and Python UDFs

APIs and Methods

We do not support these

- ✗ RDDs
- ✗ Typed Datasets
- ✗ Java/Scala udfs

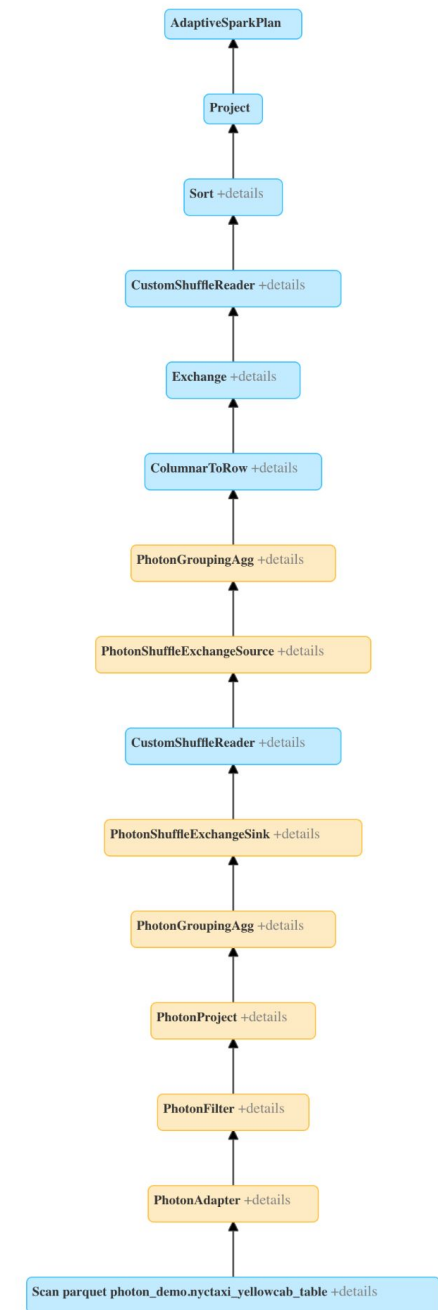
## Workloads that benefit the most

- Joins and aggregation heavy computations
- Delta Lake merge
- Reading/writing wide tables
- Decimal computations
- Delta Live Tables (DLT) and AutoLoader
- Updates and delete workloads via DV's



# Find Out Why Things Don't Photonize


- Photon makes everything faster
- If a query falls out of Photon, figure out why!
  - collect\_set? Use collect\_list(distinct)
  - UDFs? RDDs? Avoid those wherever possible.
  - Non-photonizable source? We may have a preview for that!
- Getting things to run in Photon completely solves most problems



# Predictive I/O – Speeding up point queries

Let the compute engine determine the best way to fetch data

```
SELECT protoBase64 FROM query_profile_protos WHERE id LIKE '204ff749-dc88-4b0a%'
```



# Spark 3.0 Adaptive Query Execution

- Adapts a query plan automatically at runtime based on accurate metrics
- Capabilities:
  - Sort Merge Join (SMJ) → Broadcast Hash Join (BHJ)
  - Coalesce shuffle partitions
  - Handles skew



# Compute Sizing

A balance between query performance and concurrency requirement

- Vertical Scaling – Cluster Size (2XS  $\Leftrightarrow$  4XL)
  - Larger cluster for larger queries and tables
- Horizontal Scaling – No. of Clusters (Min  $\Leftrightarrow$  Max #)
  - More cluster for more concurrent queries
- Monitor Query History to find the right fit
  - Too many queries in queue = more clusters
  - Queries taking too long = larger clusters



# Compute – Architectural Considerations

## Isolated Clusters & Warehouses to Avoid Resource Contention

- **Ephemeral Job compute**

- Jobs – Isolated compute for ingestion + ETL jobs, can be sized/optimized for that workload, run on a schedule
- Only charged for when the job is running

- **Shared development clusters**

- All-purpose – Auto-scale, auto-pause to only use when teams are actively developing, only resources needed
- Recommended to develop and test with a subset of the full dataset

- **Shared SQL warehouse for ad-hoc analysis**

- SQL warehouse – Auto-scale, auto-pause to only use when teams are actively querying, only resources needed
- Serverless available for instant startup, shutdown to reduce idle time

- **Separate SQL warehouse for BI reporting**

- Size appropriately for BI needs, avoid contention with other processes




# Serverless – Shifting the paradigm

How to fundamentally move your price-performance profile

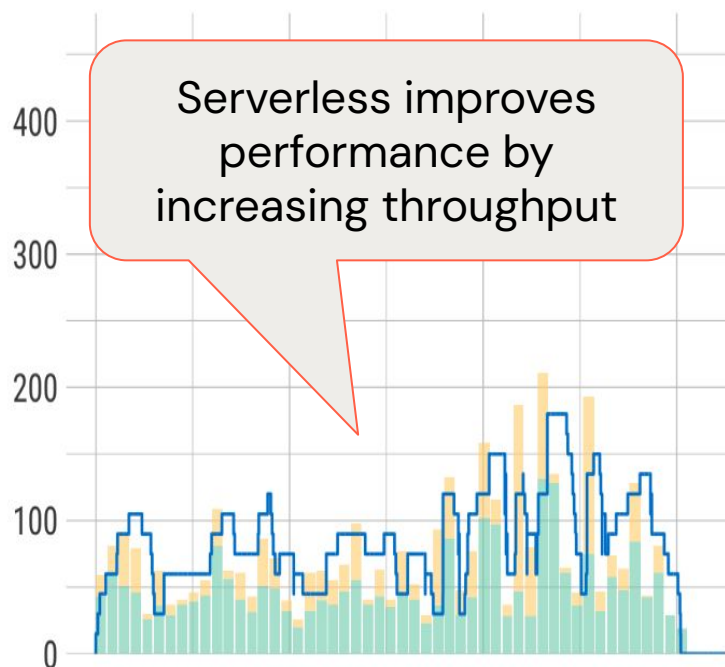
## Warehouse Metrics per Scenario

Metrics are extracted from UI elements of warehouse monitoring tab.




serverless

2xl



serverless

xl



The improved performance gives us room to reduce the warehouse size and let scaling optimise for cost



# A Note on Classic Compute

- DBSQL Warehouse automatically take cares of compute instance selection and cluster configurations
- Classic Compute Clusters (Jobs, Interactive, etc.) still give you the ability to configure instances and here is the TLDR
  - Core:RAM Ratio – Most core given enough memory for the budget
  - Processor Type – ARM based chip can work quite well
  - Local Storage – Disk Cache is useful for repeated data access
  - Driver Size – Don't over complicate it (4–8 core with 16–32 GB RAM should be enough)
  - Spot Availability – Stability is more important for long running jobs
  - Auto Scaling – Achieve high cluster utilization and reduce overall cost




# Foundational – Data



# Delta Lake

Most performant modern open data format




# Data Layout matters

Putting the right things together makes life easier

```
SELECT COUNT(*) FROM LEGOS WHERE COLOUR = "RED"
```

How do you want to store your legos?



```
What about SELECT COUNT(*) FROM LEGOS WHERE SIZE = "SMALL"?
```



# Data Layout Rationale

Different ways to organise your data so that we don't need to read too much unnecessary files

```
Select * from deltalake_table where part = 2 and col = 6
```

## Partition Pruning

/path/to/deltalake\_table/

```
part=1/part_00001.parquet
part=1/part_00002.parquet
part=1/part_00003.parquet
part=2/part_00001.parquet
part=2/part_00002.parquet
part=2/part_00003.parquet
```

## File Skipping

file_name	col_min	col_max
1.parquet	1	3
2.parquet	4	7
3.parquet	8	10





# Data Layout – To partition or not partition

DO NOT partition unless you know why you are partitioning

- **Over-partition is worse than no partition at all**
  - small files kill performance
- **Reasons to partition**
  - Table size > 100TB
  - Isolating data for separate schemas (i.e. multiplexing)
  - Governance use cases where you commonly delete entire partitions of data
  - Physical boundary to isolate data is required
- **Partition best practices**
  - Keep partition size between 1GB and 1TB
  - Combine partition with Z-order



# Data Skipping and Delta Lake Stats

- Databricks Delta Lake collects stats about the first N columns
  - `dataSkippingNumIndexedCols = 32`
- These stats are used in queries
  - Metadata only queries: `select max(col) from table`
    - Queries just the Delta Log, doesn't need to look at the files if `col` has stats
  - Allows us to skip files
    - Partition Pruning, Data Filters apply in that order
  - TimeStamp and String types aren't always very useful
    - Precision/Truncation prevent exact matches, have to fall back to files sometimes
- Avoid collecting stats on long strings
  - Put them outside first 32 columns or collect stats on fewer columns
    - `alter table change column col after col32`
    - `set spark.databricks.delta.properties.defaults.dataSkippingNumIndexedCols = 3`





# Data layout – Clustering

Sort the data in ways that you will need it

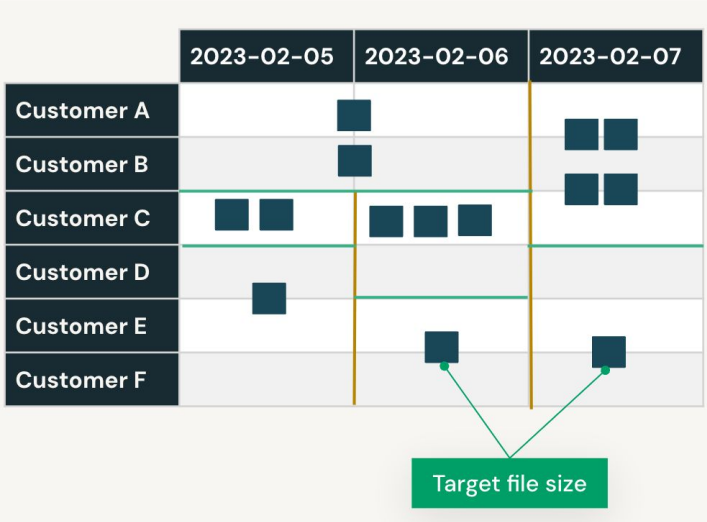
## Z-Ordering

file_name	col_min	col_max
1.parquet	6	8
2.parquet	3	10
3.parquet	1	4



file_name	col_min	col_max
1.parquet	1	3
2.parquet	4	7
3.parquet	8	10

## Liquid Clustering



# Liquid Clustering – No More partitions

- Fast
  - Faster writes and similar reads vs. well-tuned partitioned tables
- Self-tuning
  - Avoids over- and under-partitioning
- Incremental
  - Automatic partial clustering of new data
- Skew-resistant
  - Produces consistent file sizes and low write amplification
- Flexible
  - Want to change the clustering columns? No problem!



# Scenarios Benefiting from Liquid Clustering

- Tables often filtered by high cardinality columns.
- Tables with significant skew in data distribution.
- Tables that grow quickly and require maintenance and tuning effort.
- Tables with concurrent write requirements.
- Tables with access patterns that change over time.
- Tables where a typical partition key could leave the table with too many or too few partitions.



# Data Layout – File Sizes

When it comes to performance, file size matters

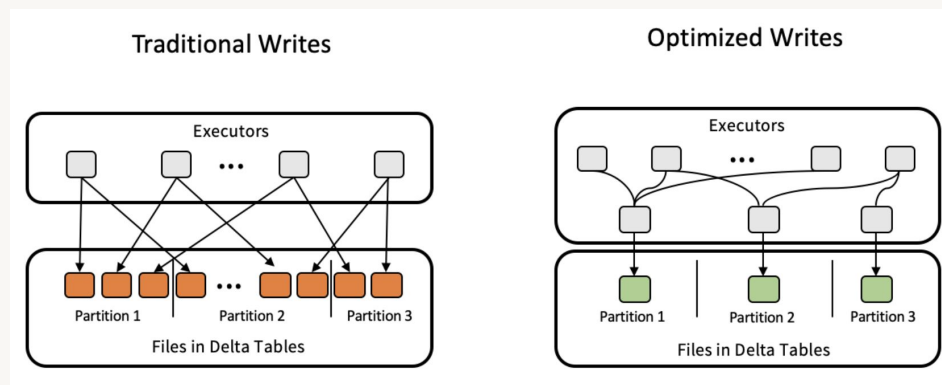
## Small File Size

- Less data to read
- More files
- Rewrite is cheaper

## Large File Size

- More data to read
- Less files
- Rewrite is expensive

`delta.tuneFileSizesForRewrites`



# Deletion Vector

Amortisation of rewrite costs

## Inserts

w/o DV

File
rowNum, data
1, data
2, data
3, data
4, data
5, data
6, data
7, data
8, data

rowNum, data
1, data

w/ DV

File
rowNum, data
1, data
2, data
3, data
4, data
5, data
6, data
7, data
8, data

rowNum, data
1, data

## Deletes

w/o DV

File
rowNum, data
1, data
2, data
3, data
4, deleted data
5, data
6, data
7, data

w/ DV

File	DV
rowNum, data	
1, data	0
2, data	0
3, data	0
4, data	1
5, data	0
6, data	0
7, data	0
8, data	0

Full File Rewrite

## Updates

w/o DV

File
rowNum, data
1, data
2, data
3, data
4, new data
5, data
6, data
7, data
8, data

w/ DV

File	DV
rowNum, data	
1, data	0
2, data	0
3, data	0
4, data	1
5, data	0
6, data	0
7, data	0
8, data	0

rowNum, data
1, new data

Full File Rewrite



# Predictive Optimization

Bringing it all together automatically

- Scheduling of these optimization can be tricky
- Mistakes can be made if users forget to set up these process
- Predictive Optimization automatically determines which operations to execute based on usage
- Prioritise high return operations based on expected benefits



# Code Optimization



# Basics

1. In production jobs, avoid operations that trigger an action besides reading and writing files. These include `count()`, `display()`, `collect()`.
2. Avoid operations that will force all computation into the driver node such as using single threaded python/pandas/Scala. Use [Pandas API on Spark](#) instead to distribute pandas functions.
3. Avoid python UDFs which execute row-by-row. Instead use native pyspark functions or [Pandas UDFs](#) for vectorized UDFs.
4. Use Dataframes or Datasets instead of RDDs. RDDs cannot take advantage of the cost-based optimizer.



# Controlled Batch Size For Streaming

- Defaults are large and non-deterministic:
  - 1000 files per micro-batch
  - No limit to input batch size
- Optimal mini-batch size → Optimal cluster usage
- Suboptimal mini-batch size → Performance cliff
- 

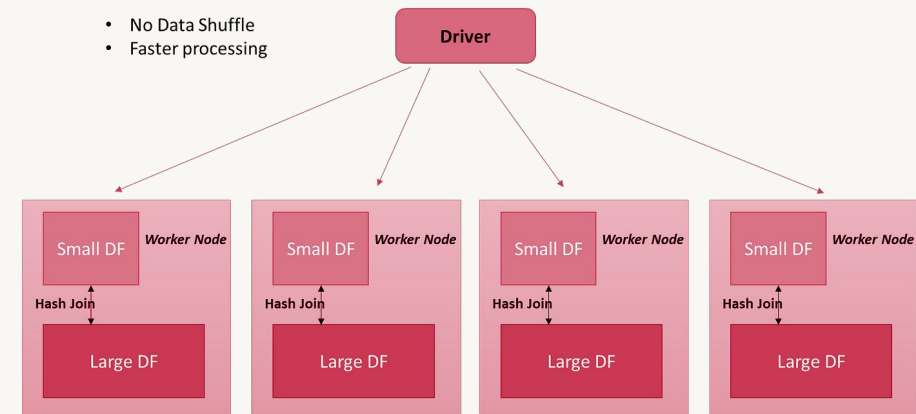
## Per Trigger Settings:

- Kafka
  - `maxOffsetsPerTrigger`
- Delta Lake and Auto Loader
  - `maxFilesPerTrigger`
  - `maxBytesPerTrigger`



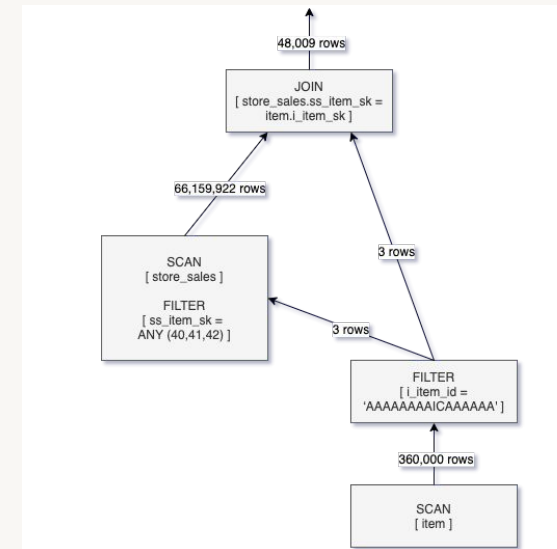
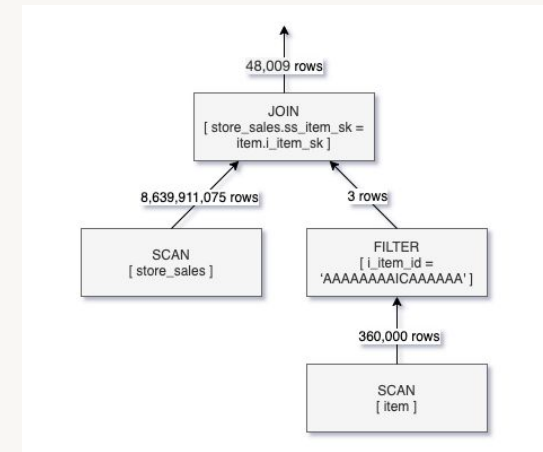
# Broadcast Join

- The most performant type of join
- Distributes a small dataset across all worker nodes to minimize shuffling and speed up query execution.
- Triggered when smaller table/data-frame is lesser than **spark.sql.autoBroadcastJoinThreshold** (10 MB by default)
- Control micro-batch size to trigger it when joining with target table (e.g. during merge) in Structured Streaming
- Prerequisite for Dynamic File Pruning (DFP)



# Dynamic File Pruning (DFP)

- Intelligently skipping non-relevant data files during selective joins, achieving up to 8x faster performance
- Key Prerequisites:
  - The join strategy is BROADCAST JOIN
  - The join type is INNER or LEFT-SEMI



# Incremental Processing via Streaming

- Consider streaming for all of your workloads:
  - Incremental processing resulting in reduced latency and quick time to insight
  - Built-in Checkpointing for exactly-once guarantees and fault tolerance
  - Efficient resource utilization
- Move to a CDC architecture pattern where you are only processing change data will greatly reduce overall processing time

# DLT vs Structured Streaming

Feature	Structured Streaming	Delta Live Tables
Autoloader	✓	✓
Trigger Options	✓	✓
Workflow Support	✓	✓
Schema Evolution	✓	✓
CDC with SCD Type 1 and 2	✗	✓
Data Quality Constraints + Monitoring	✗	✓
Automatic Orchestration	✗	✓
Concurrent Streaming Jobs on a Cluster	✗ (not recommended)	✓
Pipeline Observability	✗	✓
Simplified Deployment + UI	✗	✓
Enhanced Autoscaling	✗	✓
Enzyme Runtime Engine	✗	✓

# Simple query is usually a fast query

Get to the results with the least amount of data and transformation

1. Predicates are pushed as far up as possible
  - a. Select the least amount of columns and rows that you need
  - b. Align data layout with commonly used predicates (ZORDER, LIQUID)
  - c. Make sure data is right sized (OPTIMIZE)
2. Simplify how you join your tables
  - a. Join the smallest tables first / collect statistics for the optimizer to do it for you
  - b. Provide join hints if you can
  - c. Reduce unnecessary data movement, i.e. If you know your data layout and join keys you can use the right join strategy (sort-merge v. shuffle hash vs broadcast)
3. Simplify operations
  - a. Be careful about expensive operations (distinct, sort, window)
  - b. UDFs are powerful but they are not fast, try to use native functions as much as possible



# Diagnosis



# Performance 5S's

Skew | Spill | Shuffle | Storage (small files) | Serialization






# Common Performance Bottlenecks

Encountered with any big data or MPP system

Symptom	Details
Skew	An imbalance in the size of partitions
Spill	The writing of temp files to disk due to a lack of memory
Shuffle	The act of moving data between executors
Small Files	A set of problems indicative of high overhead due to tiny files
Serialization	The distribution of code segments across the cluster

# Skew – Mitigation

- Repartition Data (comes with caveats)
- Enable Adaptive Query Execution (AQE) in Spark 3 (enabled by default from DBR 7.3+)
- Employ skew hints
- Salting



Running command...

▼ (1) Spark Jobs

▼ Job 39 [View](#) (3 stages)

Stage 53: 8/8 (0 running) ⓘ

Stage 54: 8/8 (0 running) ⓘ

Stage 55: 199/200 (1 running) ⓘ



# Spill – Mitigation


**Option #1** – Allocate a cluster with more memory per worker

*Tip: Larger, fewer nodes > Smaller, more nodes*

**Option #2** – In the case of *skew*, address that root cause first.

**Option #3** – Decrease the size of each partition by increasing the number of partitions


- By managing `spark.sql.shuffle.partitions`
- By explicitly repartitioning
- By managing `spark.sql.files.maxPartitionBytes`



# Shuffle – Mitigation

## TIP: Don't get hung up on trying to remove every shuffle

- Shuffles are often a necessary evil. Focus on the [more] expensive operations instead. Many shuffle operations are actually quite fast.
- Reduce network IO by using **fewer and larger workers**
- Reduce the **amount of data** being shuffled
  - Narrow your columns
  - Preemptively filter out unnecessary records



# Storage (Tiny Files) – Mitigation

Make sure you constantly **optimize** and **vacuum** your delta tables


- OPTIMIZE will compact and Z-order/Cluster your files
- Vacuum will delete old versions and cleanup the metadata

(Predictive Optimization automates it)

Enable **auto-optimize** in all your tables unless there is a good reason for not doing so (e.g. streaming low latency requirements)

```
spark.databricks.delta.optimizeWrite.enabled = true
```

```
spark.databricks.delta.autoCompact.enabled = true
```



# Serialization – Mitigation

- PySpark UDFs have significant serialization overhead between JVM and Python interpreter
- Act like “black box” and cannot be optimized by Spark's Catalyst optimizer, leading to suboptimal execution
- Thus, wherever possible, don't use UDFs!
- The native and SQL higher-order functions are very robust
- But if you have to...
  - Use Arrow Optimized Python UDFs
  - Use Vectorized UDFs aka Pandas UDFs



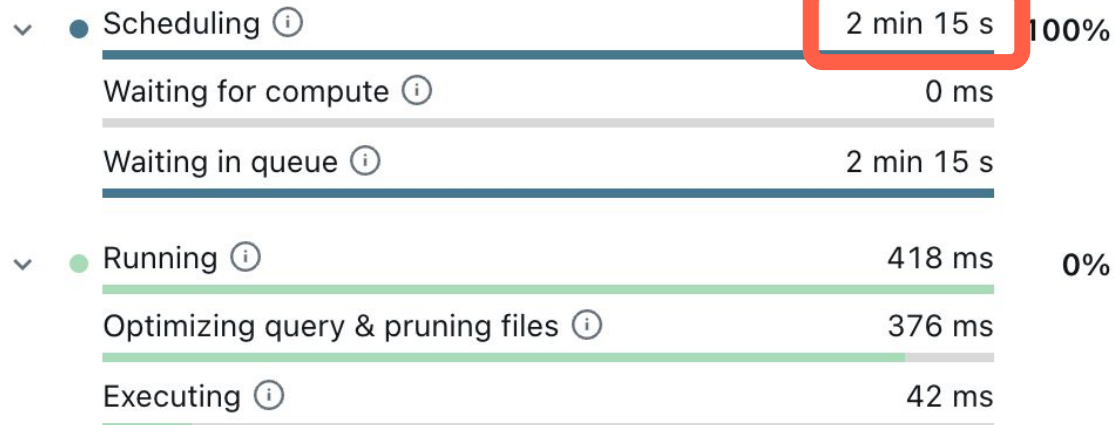
# Diagnosing performance issues

## Scheduling Time v. Running Time

- Same wall clock time != same performance
- Scheduling time has nothing to do with your code
  - Waiting for compute = Need running compute = Serverless / pre-started cluster
  - Waiting in queue = We need more concurrency = Increase max # of clusters

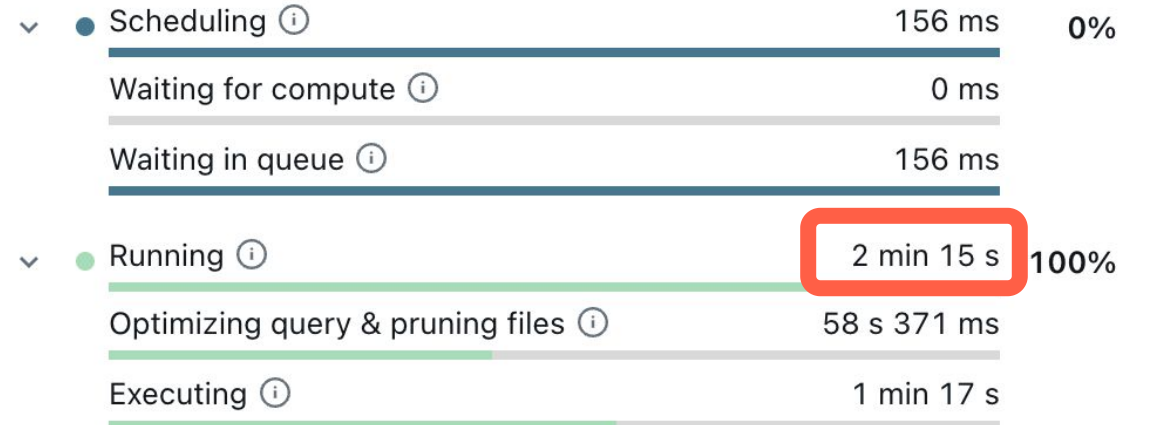
### Wall-clock duration ⓘ

Total wall-clock duration 2 min 15 s



### Wall-clock duration ⓘ

Total wall-clock duration 2 min 15 s



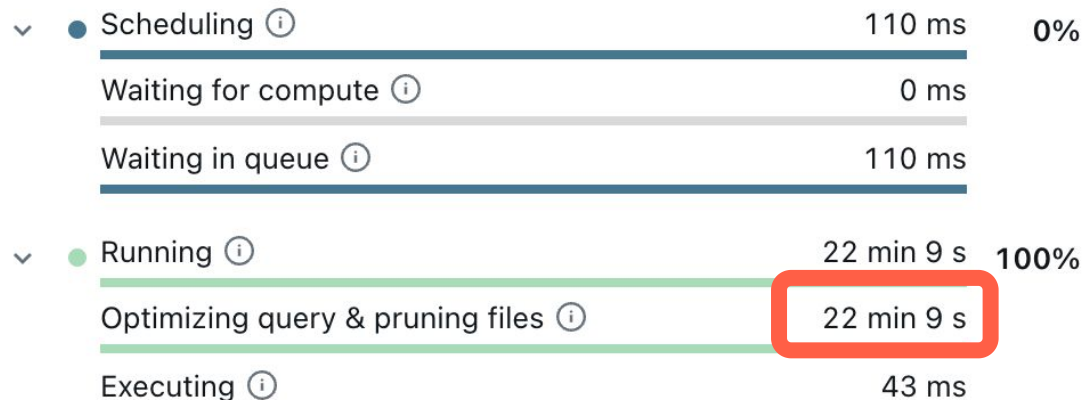
# Diagnosing performance issues

## Running Time Breakdown

- Same running time != Same time spent on execution
- Long Optimizing Query & Pruning Files Time = better stats collection

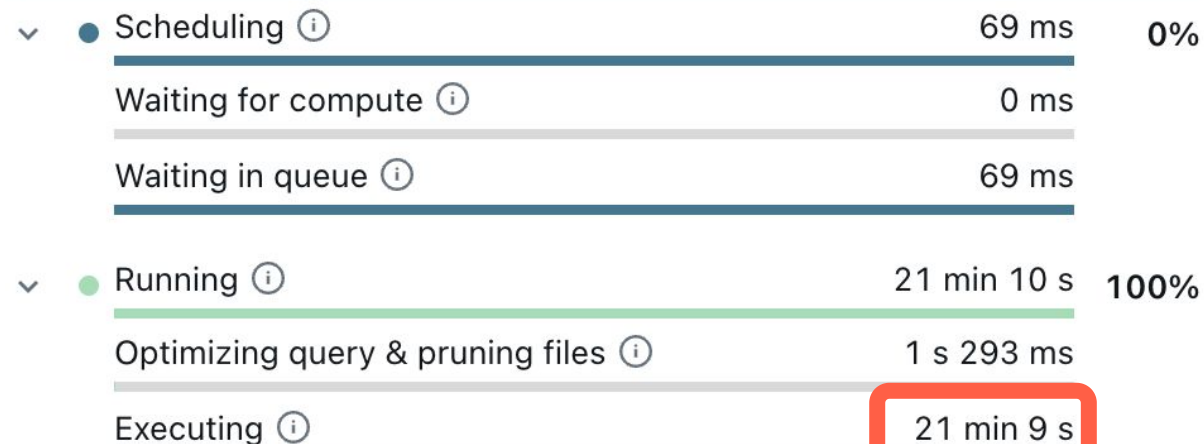
Wall-clock duration ⓘ

Total wall-clock duration 22 min 9 s



Total wall-clock duration

21 min 10 s





# Diagnosing performance issues

## Execution Details

Make sure photon is close to 100%

### Aggregated task time ⓘ

Tasks total time	11.14 h
Tasks time in Photon	100 %

### IO

Rows returned	8
Rows read	20,259,286,241
Bytes read	688.38 GB
Bytes read from cache	85 %
Bytes written	0 bytes

### Files & partitions

Files read	14,141
Partitions read	11

### Spilling

Bytes spilled to disk	0 bytes
-----------------------	---------

Does the number of rows read make sense? Did you read too much data?

Disk cache meant your data is already cached on local storage

Spill meant your warehouse/cluster is too small, i.e. Not enough RAM

If your data read is correct, are you reading too many files or partitions?



# Diagnosing performance issues

## Understanding Query Profile

- Execution can be broken down to individual operations within your query
- The most time spent operation is likely where you need to start
- It should tell you which part of the query is causing problem
- Knowing what is the problem doesn't mean it is an easy fix



# What now?



# Key takeaways

Optimize only when necessary

- Know what you are optimizing towards
- Focus on the easy things first (Platform → Data → Query)
- Leverage latest compute and features (i.e. Serverless Compute, Photon, Predictive Optimization, Liquid Clustering etc)
- Scale vertically or horizontally based on indicators (complexity, concurrency)
- Pivot to end-to-end incremental processing, via streaming, wherever possible
- Extensively use observability and monitoring tools to guide optimization efforts
- Know when to stop optimizing and start building more useful things



# Thank you

