



Simplify your Streaming: Delta Live Tables



Jerrold Law

Sr. Specialist Solutions Architect



Housekeeping

- Your connection will be muted
- We will share recording with all attendees after the session
- Submit questions in the Q&A panel
- If we do not answer your question during the event, we will follow-up with you to get you the information you need!



Where to learn more: go/dlt

Customer Facing Slides



This User Guide



Sales Play: Win with Modern Data Engineering

A sales play slide titled "SALES PLAY: Win with Modern Data Engineering". The slide is divided into several sections: "SOLUTION ONE-LINER", "KEY PERSONAS", "COMPETITIVE RESOURCES", "CUSTOMER STORIES", and "REFERENCES & COLLATERAL". Each section contains brief descriptions and links to further resources like "CURRENT PITCH DECK", "DISCOVERY GUIDE", and "COMPETITIVE GUIDE". The slide has a dark background with a circular photo of people in a lab setting on the right and abstract orange and yellow geometric shapes on the left.





Good data is the foundation of a Lakehouse

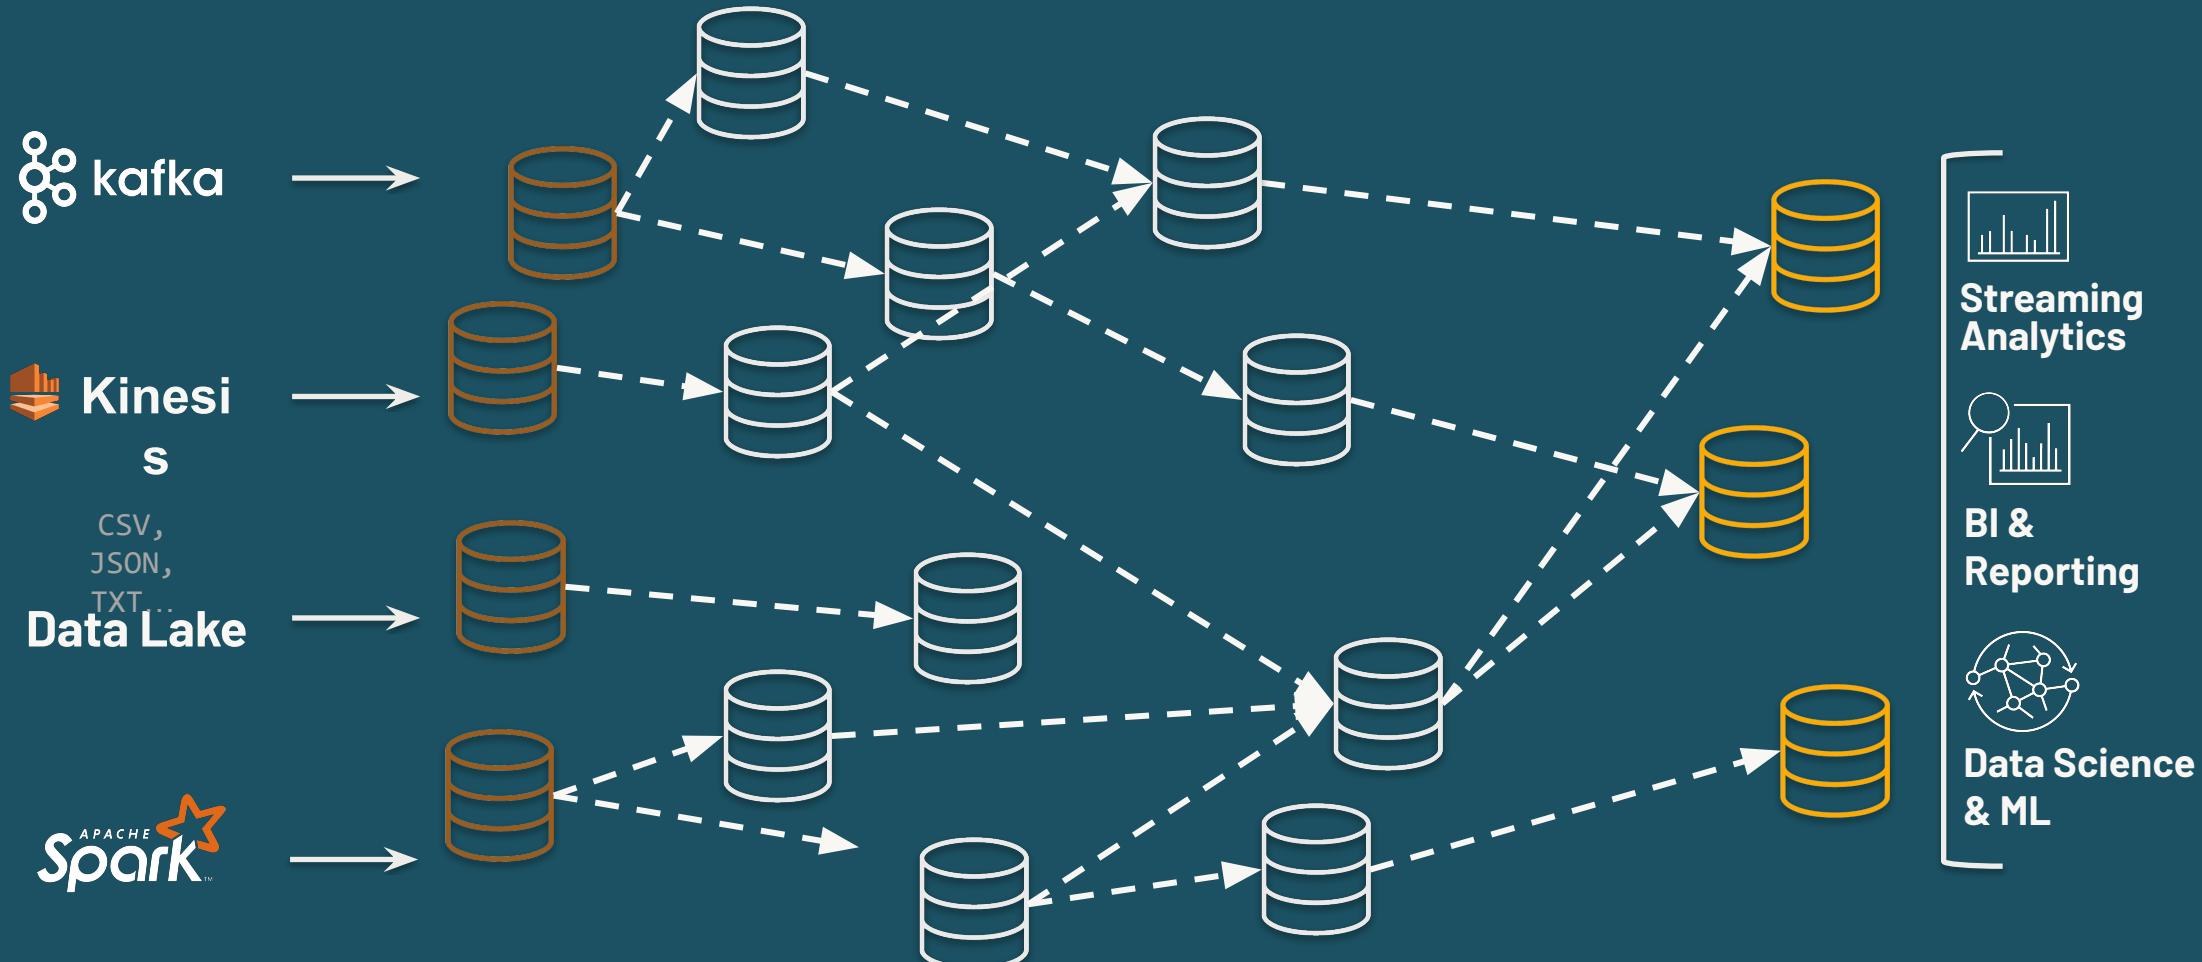
All data professionals need clean, fresh and reliable data.





But the reality is not so simple

Maintaining data quality and reliability at scale is often **complex and brittle**



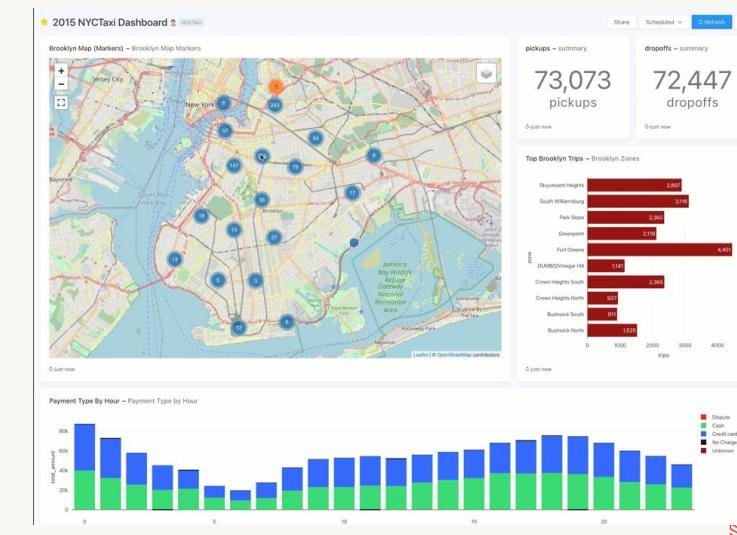
Life as a data professional...

From: **The CEO** <ali@databricks.com>
Subject: Need an analysis ASAP!
To: Michael Armbrust
<michael@databrick.com>

Hey Michael, I need a quick analysis of our net customer retention and how it has changed over the past few quarters. Raw data can be found at
s3://our-data-bucket/raw_data/...

```
1 %fs ls /data/sensors
```

	path
1	dbfs:/data/sensors/_SUCCESS
2	dbfs:/data/sensors/_committed_3908896360792309052
3	dbfs:/data/sensors/_started_3908896360792309052
4	dbfs:/data/sensors/part-00000-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
5	dbfs:/data/sensors/part-00001-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
6	dbfs:/data/sensors/part-00002-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
7	dbfs:/data/sensors/part-00003-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3
8	dbfs:/data/sensors/part-00004-tid-3908896360792309052-adec30c0-9ba8-4344-a36c-7ec3



Going from query to production

The tedious work required to turn SQL queries into reliable ETL Pipelines

From: **The CEO** <ali@databricks.com>

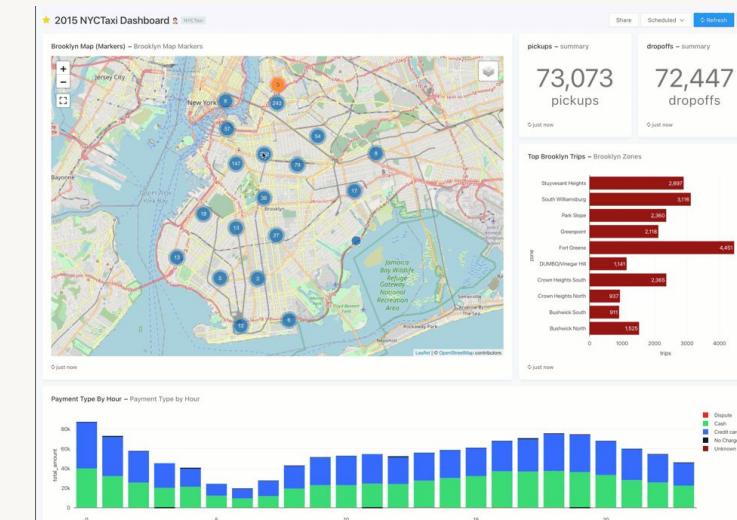
Subject: Need an analysis ASAP!

To: Michael Armbrust <michael@databrick.com>

every minute

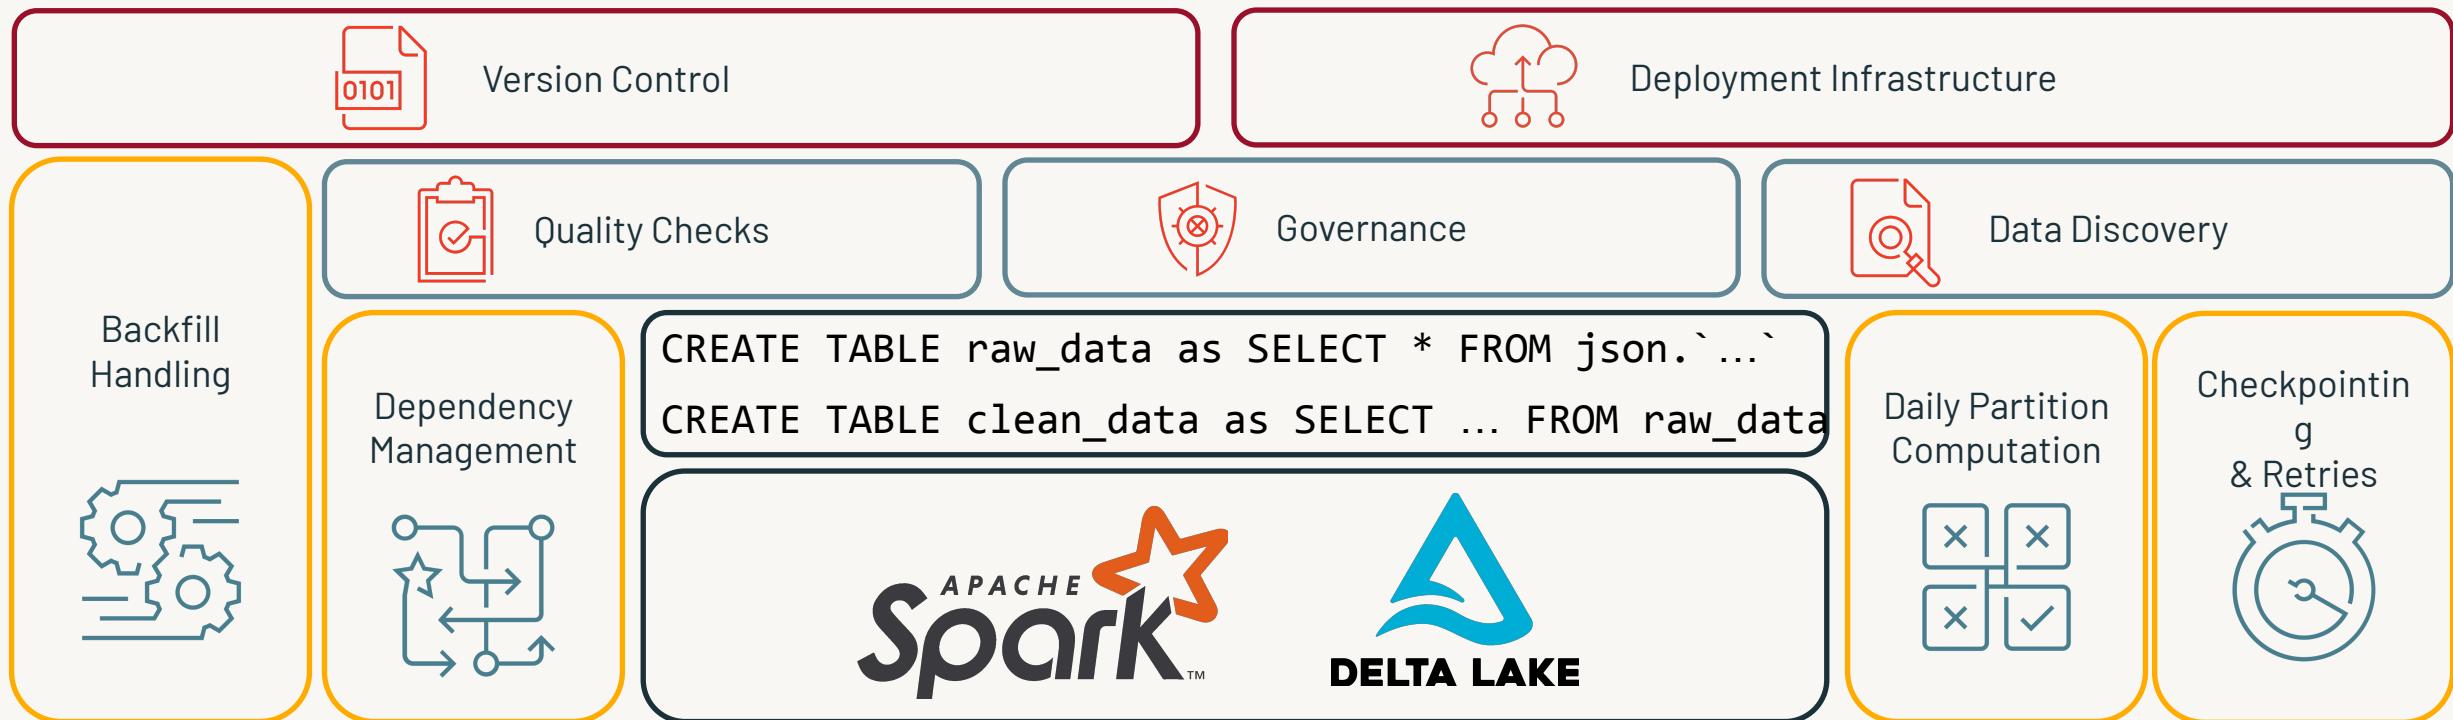
Great report! Can you update it everyday?

```
CREATE TABLE raw_data as SELECT * FROM  
json` `raw_data  
CREATE TABLE clean_data as SELECT ... FROM raw_data
```



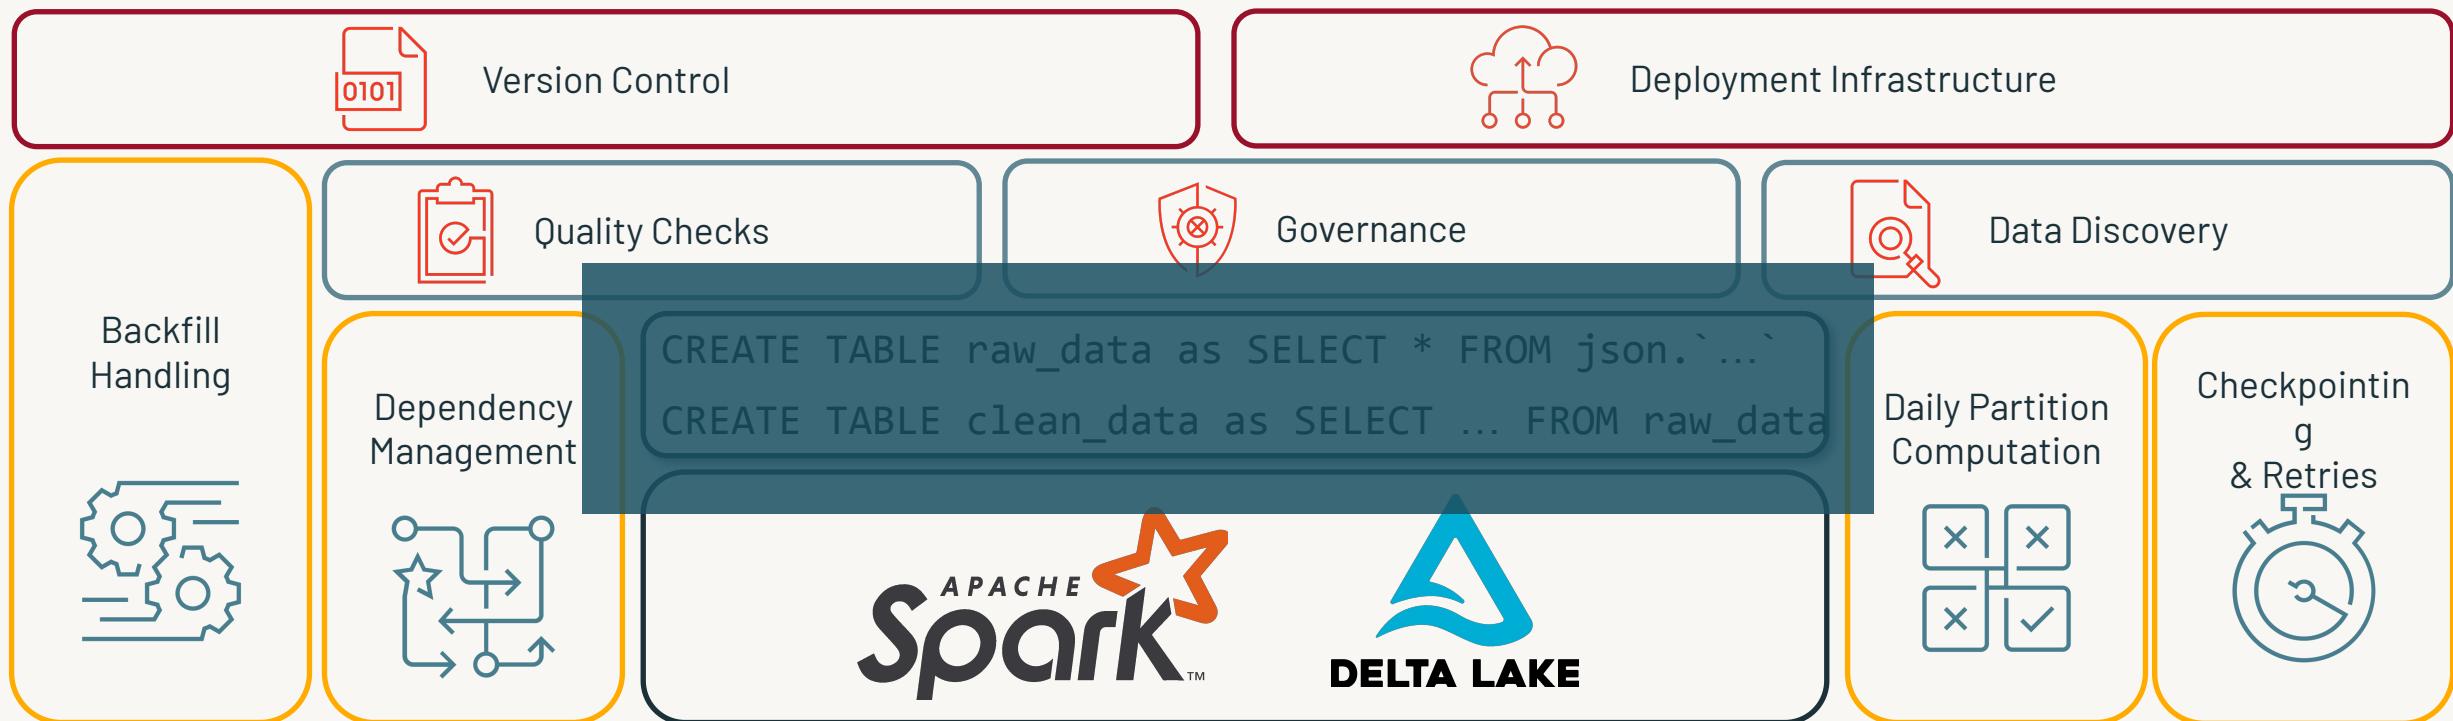
The **slog** from query to production

The **tedious work** required to turn SQL queries into **reliable ETL Pipelines**



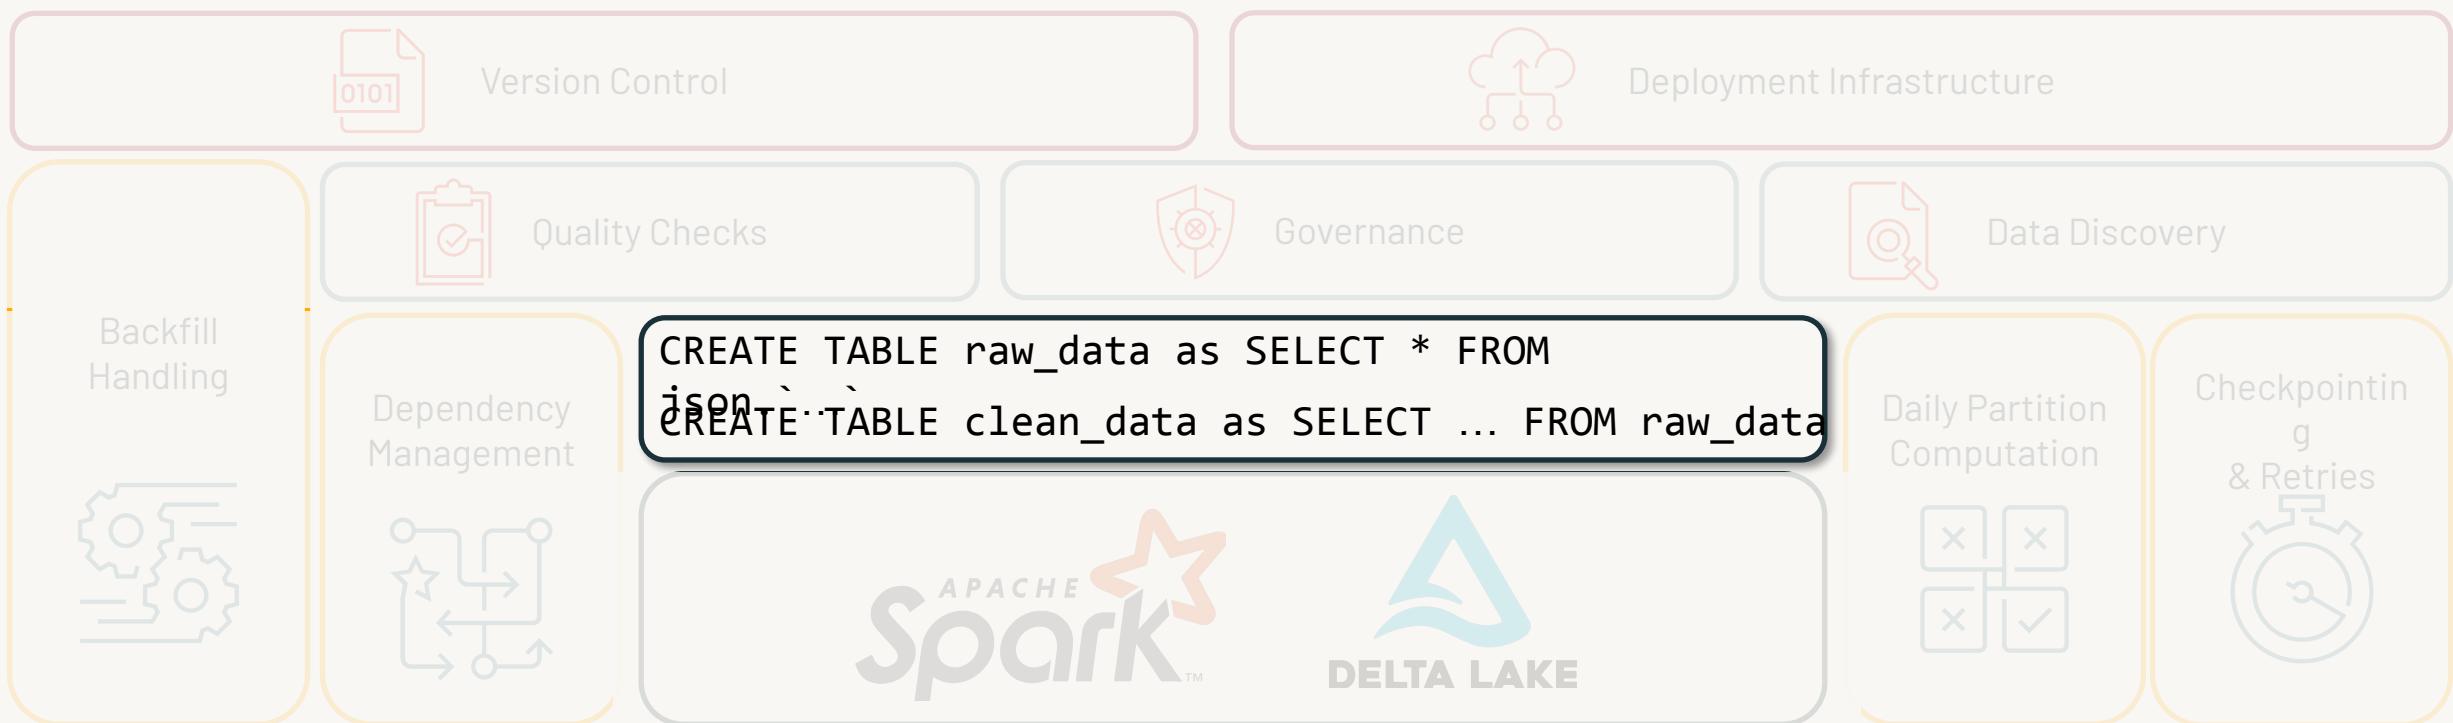
Operational complexity dominates

Time is spent on **tooling** instead of on **transforming**



Where you should focus your time

Getting value from data



Introducing Delta Live Tables

From query to **production pipeline** just by adding **LIVE**.

```
CREATE STREAMING TABLE raw_data as SELECT * FROM cloud_files(...)  
CREATE MATERIALIZED VIEW clean_data as SELECT ... FROM LIVE.raw_data
```



Repos



Unity Catalog



Databricks Workflows

Delta **Live** Tables

Full Refresh



Dependency Management



Expectations



Incremental Computation



Checkpointing & Retries



Introduction to Declarative Programming

Declarative programs say **what should be done**, not **how to do it**

Imperative program

```
numbers = [...]  
  
sum = 0  
  
for n in numbers:  
    sum += n  
  
print(n)
```

Declarative Program

```
SELECT sum(n) FROM numbers
```

The **query optimizer** figures out the **best way** to calculate the sum

Data Independence
No need to write the **query** when **physical things change** (partitioning, storage location, indexes, etc)



Declarative Programming with DLT

Declarative programs say **what should be done**, not **how to do it**

Spark imperative program

```
date = current_date()  
  
spark.read.table("orders")  
  
.where(s"date = $date")  
  
.select("sum(sales)")  
  
.write  
  
.mode("overwrite")  
  
.replaceWhere(s"date = $date")  
  
.table("sales")
```

DLT Declarative Program

```
CREATE MATERIALIZED VIEW sales  
  
AS SELECT date, sum(sales)  
  
FROM orders  
  
GROUP BY date
```

The **DLT runtime** figures out the **best way** to create or update this table



Workflows Or DLT?

Often Both: Workflows can orchestrate anything, including DLT

Use Workflows to run any task

- At some schedule
- After other tasks have completed
- When a file arrives
- When another table is updated

Use DLT for managing dataflow

- Creating/updating delta tables
- Running Structured Streaming

Why declarative programming?

Understanding your objective allows DLT to take care of the boring stuff

- Eliminate Boilerplate
 - Fault-tolerance
 - State management
 - Object lifecycle
 - Scheduling, dependencies, parallelism
 - Common data problems (cdc, schema evolution)
- Optimization
 - Choose the best strategy for execution dynamically, or across multiple queries
- Operations
 - HMR – detect and mitigate system regressions and cloud issues automatically

The core abstractions of DLT

You define datasets, and DLT automatically keeps them up to date

Streaming Tables

A delta table with stream(s) writing to it.

Used for:

- Ingestion
- Low latency transformations
- Huge scale

Materialized View

The result of a query, stored in a delta table.

Used for:

- Transforming data
- Building aggregate tables
- Speeding up BI queries and reports

What is a Streaming Table?

A delta table that has structured streaming writing to it.

```
CREATE STREAMING TABLE report  
AS SELECT *  
FROM cloud_files("/mydata/", "json")
```

- Streaming tables **read from append-only data sources** such as Kafka, Kinesis, or Auto Loader (files on cloud storage)
- Streaming tables allow you to **reduce costs and latency** by **reading each input record only once**.
- Streaming tables **support DML** (UPDATE, DELETE, MERGE) for ad-hoc data manipulation (i.e. GDPR, etc)

What is a Materialized View?

The result of a query, precomputed and stored in Delta

```
CREATE MATERIALIZED VIEW report  
AS SELECT sum(profit)  
FROM prod.sales  
GROUP BY date
```

- A materialized view will always return the result of the defining query, at the moment it was last updated (i.e. a snapshot)
- You cannot modify the data in a materialized view, you can change its query.

Data quality using Expectations

Ensure correctness with Expectations

Expectations are tests that ensure data quality in production

```
CONSTRAINT valid_timestamp  
EXPECT (timestamp > '2012-01-01')  
ON VIOLATION DROP
```

```
@dlt.expect_or_drop(  
    "valid_timestamp",  
    col("timestamp") > '2012-01-01')
```

Expectations are true/false expressions that are used to validate each row during processing.

DLT offers flexible policies on how to handle records that violate expectations:

- Track number of bad records
- Drop bad records
- Abort processing for a single bad record

Expectations using the power of SQL

Use SQL aggregates and joins to perform complex validations

-- Make sure a primary key is always unique.

```
CREATE MATERIALIZED VIEW report_pk_tests(  
    CONSTRAINT unique_pk EXPECT (num_entries = 1)  
)  
AS SELECT pk, count(*) as num_entries  
FROM LIVE.report  
GROUP BY pk
```

Expectations using the power of SQL

Use SQL aggregates and **joins** to perform complex validations

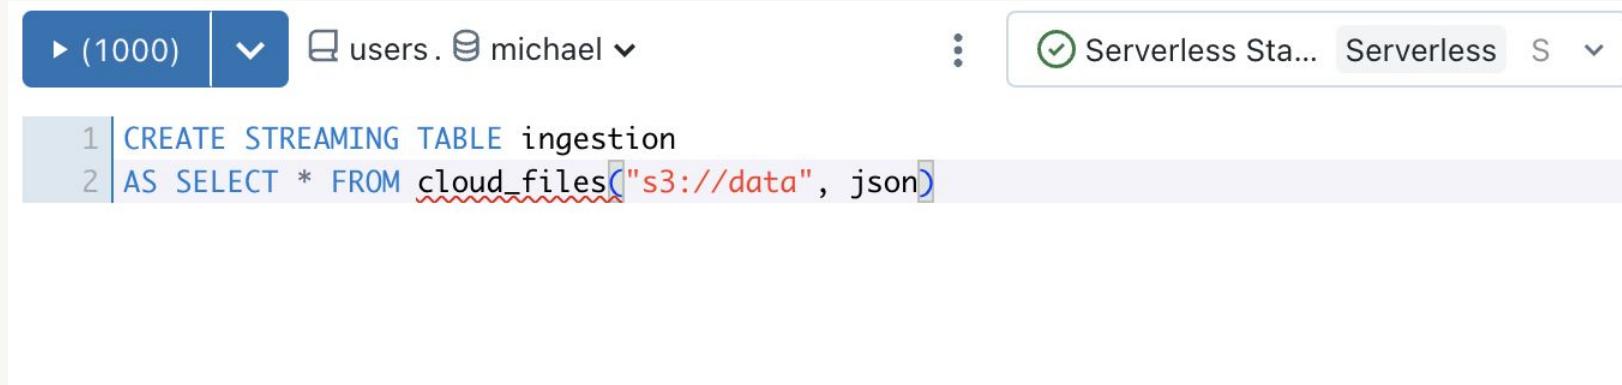
- Compare records between two tables,
- or validate foreign key constraints.

```
CREATE MATERIALIZED report_compare_tests(  
    CONSTRAINT no_missing EXPECT (r.key IS NOT NULL)  
)  
  
AS SELECT * FROM LIVE.validation_copy v  
LEFT OUTER JOIN LIVE.report r ON v.key = r.key
```

Announcing: DLT in DBSQL

PREVIEW

DLT powers streaming tables and materialized views created in DBSQL



```
▶ (1000) ▾ users. michael ▾ : Serverless Sta... Serverless S ▾  
1 CREATE STREAMING TABLE ingestion  
2 AS SELECT * FROM cloud_files("s3://data", json)
```

A robust, scalable, incremental ingestion pipeline in
one SQL command

Quick Note: Evolving terminology

Existing users of DLT will notice that we've evolved the names

- STREAMING LIVE TABLE ➔ STREAMING TABLE
- LIVE TABLE ➔ MATERIALIZED VIEW

The semantics remain the same, and we'll support the old syntax for compatibility. Our goal is to simplify the syntax and match other systems.

Deep dive into Streaming

Built on Spark™ Structured Streaming

Streaming tables combine spark structured streaming with a delta table.

Streaming Computation Model: Input is a growing **append-only table**

- Files uploaded to cloud storage
- Message busses like kafka, kinesis, or eventhub
- Delta tables with `delta.appendOnly=true`
- Transaction logs of other databases

Rather than **wait until all data has arrived**, structured streaming can produce **results on demand**.

- **Lower latency** by processing less data each update
- **Lower costs** by avoiding redundant work

Streaming does not always mean expensive

Delta live tables lets you choose how often to update the results.

Triggered: Manually

Costs: lowest

Latency: highest



Triggered: On a schedule using Databricks Jobs

Costs: depends on frequency

Latency: 10 minutes to months

Schedule ▾

Every Day at 22 : 14 (UTC-07:00) Pacific Ti...

Continually

Costs: highest

Latency: minutes to seconds
(for some workloads)

Pipeline Mode ?

Triggered Continuous

Using Streaming Tables for **ingestion**

Easily ingest files from cloud storage as they are uploaded

```
CREATE STREAMING TABLE raw_data  
AS SELECT *  
FROM cloud_files("/data", "json")
```

This example creates a table with all the json data stored in "/data":

- `cloud_files` keeps track of which files have been read to **avoid duplication and wasted work**
- Supports both listing and notifications for **arbitrary scale**
- Configurable **schema inference** and **schema evolution**

Using Spark™ Structured Streaming for ingestion

Easily ingest records from message buses

```
@dlt.table  
def kafka_data():  
  
    return spark.readStream \  
        .option("format", "kafka") \  
        .option("subscribe", "events") \  
        .load()
```

This example creates a table with all the records published to the Kafka topic “event”.

- The Kafka source + DLT automatically track which partitions / offsets have already been read.
- Any structured streaming source included in DBR can be used with DLT
- Message buses provide the lowest latency for ingesting data

Using the SQL STREAM() function

Stream data from any Delta table

```
CREATE STREAMING TABLE mystream  
AS SELECT ...  
FROM STREAM(prod.events)
```

- STREAM(...) reads newly inserted records into the table as they arrive.
- Only works for **append-only delta tables**

```
CREATE STREAMING TABLE mystream  
AS SELECT ...  
FROM STREAM(delta.`s3:/...`)
```

Understanding streaming state

Streaming tables are stateful

Each input row is processed only once, even if the query changes.

A change to a streaming live table's definition does not reread already processed data:

```
CREATE STREAMING TABLE raw_data  
AS SELECT a + 1 AS a  a * 2 AS a  
FROM cloud_files("/data", "json")
```

"/data"

```
{"a": 1} →  
{"a": 2} →  
{"a": 3} →  
{"a": 4} →
```

raw_data

b
2
3
6
8

Streaming joins are stateful

Enrich data by joining with an up-to date-snapshot stored in delta

- Snapshot is automatically updated each microbatch if it changes
- A change to joined table snapshot does not recompute results:

```
CREATE STREAMING TABLE raw_data
AS SELECT *
FROM cloud_files("/data", "json") f
JOIN prod.cities c USING id
```

"**/data**"

{"a": 1} →
{"a": 1} →

raw_data	
id	city
1	Bekerly, CA
1	Berkeley, CA

prod.cities

prod.cities	
id	city
1	Bekerly, CA
1	Berkeley, CA

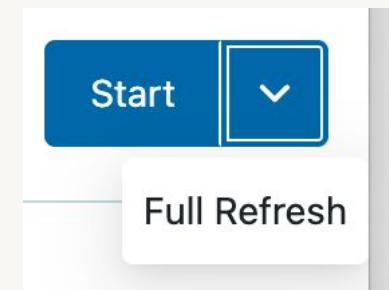
Reprocess data using full refresh

Automatically perform backfills after critical changes using full refresh

Full-refresh **clears the table's data and the queries state, reprocessing all the data.**

```
CREATE STREAMING TABLE raw_data  
AS SELECT a * 2 AS a  
FROM cloud_files("/data", "json")
```

```
REFRESH raw_data FULL
```



"/data"

```
{"a": 1} →  
{"a": 2} →  
{"a": 3} →  
{"a": 4} →
```

raw_data

b
2
3
6
8

After full-refresh

```
{"a": 1} →  
{"a": 2} →  
{"a": 3} →  
{"a": 4} →
```

b
2
4
6
8

Using stream-stream joins

Joining facts with stream-stream joins

Stream-stream joins let you combine facts that arrive near each other

```
@dlt.table()  
def joined():  
    impressions = read_stream("impressions") \  
        .withWatermark("impressionTime", "10 minutes")  
  
    clicks = clicks \  
        .withWatermark("clickTime", "20 minutes")  
  
    return impressions.join(clicks,  
        expr("""clickAdId = impressionAdId AND  
              clickTime >= impressionTime AND  
              clickTime <= impressionTime + interval 5 minutes"""))
```

This example, based on ad-tech, joins a stream of clicks with details about an ad impression that caused them to occur.

Unlike in snapshot-stream join, a stream-stream join will buffer records until a match appears.

Joining facts with stream-stream joins

Stream-stream joins let you combine facts that arrive near each other

```
@dlt.table()  
def joined():  
    impressions = read_stream("impressions") \  
        .withWatermark("impressionTime", "10 minutes")  
  
    clicks = clicks \  
        .withWatermark("clickTime", "20 minutes")  
  
    return impressions.join(clicks,  
        expr("""clickAdId = impressionAdId AND  
              clickTime >= impressionTime AND  
              clickTime <= impressionTime + interval 5 minutes"""))
```

Two streams of facts, each with a watermark

The join condition
A time bound on how far apart each fact can arrive

Pitfall: If the watermark or timebound are missing, the join will buffer all data forever.

Adding and removing streaming sources

Streaming Tables and Flows

In DLT a **flow** is the object that **updates a table with new data**

The **simple syntax** we've shown so far is a syntactic sugar for creating a **table** and a **flow** with a single DDL command.

```
CREATE STREAMING TABLE raw_data  
AS SELECT * FROM kafka(...)
```



```
CREATE STREAMING TABLE raw_data  
CREATE FLOW raw_data _____  
AS INSERT INTO LIVE.raw_data BY NAME  
SELECT * FROM kafka(...)
```

Checkpoints are identified by the **name of the flow**

→ checkpoints/raw_data

Kafka Offsets
{
0: 123354,
2: 9573943,
3: 83275092...
}

Multi-flow Tables

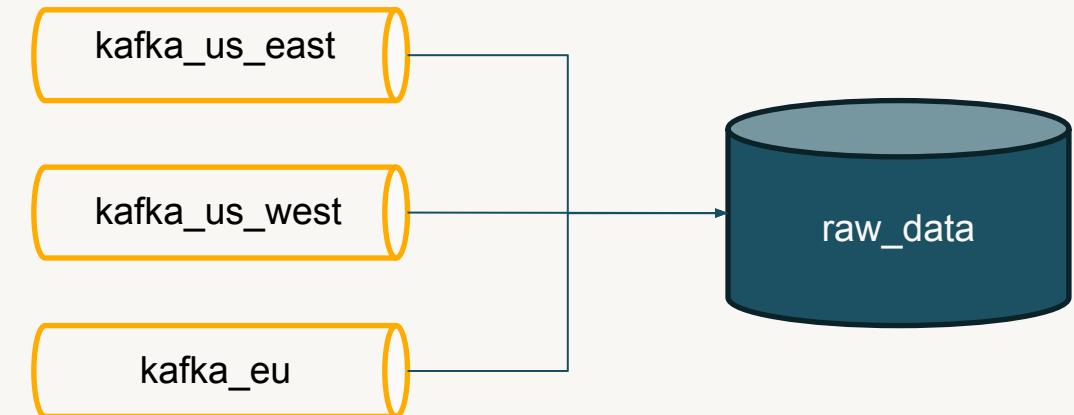
Evolve your streaming sources, with differently named flows

```
CREATE STREAMING TABLE raw_data
```

```
CREATE FLOW kafka_us_east
AS INSERT INTO LIVE.raw_data BY NAME
SELECT * FROM kafka(...)
```

```
CREATE FLOW kafka_us_west
AS INSERT INTO LIVE.raw_data BY NAME
SELECT * FROM kafka(...)
```

```
CREATE FLOW kafka_eu
AS INSERT INTO LIVE.raw_data BY NAME
SELECT * FROM kafka(...)
```



Streaming change data capture (CDC)

APPLY CHANGES INTO for CDC

Maintain an up-to-date replica of a table stored elsewhere

```
CREATE STREAMING TABLE cities
```

{UPDATE}

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(city_updates)
```

{DELETE}

```
KEYS (id)
```

{INSERT}

```
SEQUENCE BY ts
```



Up-to-date Snapshot

APPLY CHANGES INTO for CDC

Maintain an up-to-date replica of a table stored elsewhere

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(city_updates)  
KEYS (id)  
SEQUENCE BY ts
```

city_updates

```
{"id": 1, "ts": 100, "city": "Bekerly, CA"}
```

A source of changes,
currently this has to be a
stream.

APPLY CHANGES INTO for CDC

Maintain an up-to-date replica of a table stored elsewhere

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(city_updates)  
KEYS (id)  
SEQUENCE BY ts
```

A **target** for the changes to be applied to.

city_updates

```
{"id": 1, "ts": 100, "city": "Bekerly, CA"}
```

cities

id	city
1	Bekerly, CA

APPLY CHANGES INTO for CDC

Maintain an up-to-date replica of a table stored elsewhere

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(city_updates)  
KEYS (id)  
SEQUENCE BY ts
```

A unique **key** that can be used to identify a given row.

city_updates

```
{"id": 1, "ts": 100, "city": "Bekerly, CA"}
```

cities

id	city
1	Bekerly, CA

APPLY CHANGES INTO for CDC

Maintain an up-to-date replica of a table stored elsewhere

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(city_updates)  
KEYS (id)  
SEQUENCE BY ts
```

A **sequence** that can be used to order changes:

- Log sequence number (Isn)
- Timestamp
- Ingestion time

city_updates

```
{"id": 1, "ts": 100, "city": "Bekerly, CA"}
```

cities

id	city
1	Bekerly, CA

APPLY CHANGES INTO for CDC

Maintain an up-to-date replica of a table stored elsewhere

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(LIVE.city_updates)  
KEYS (id)  
SEQUENCE BY ts
```

city_updates

```
{"id": 1, "ts": 100, "city": "Bekerly, CA"}  
{"id": 1, "ts": 200, "city": "Berkeley, CA"}
```

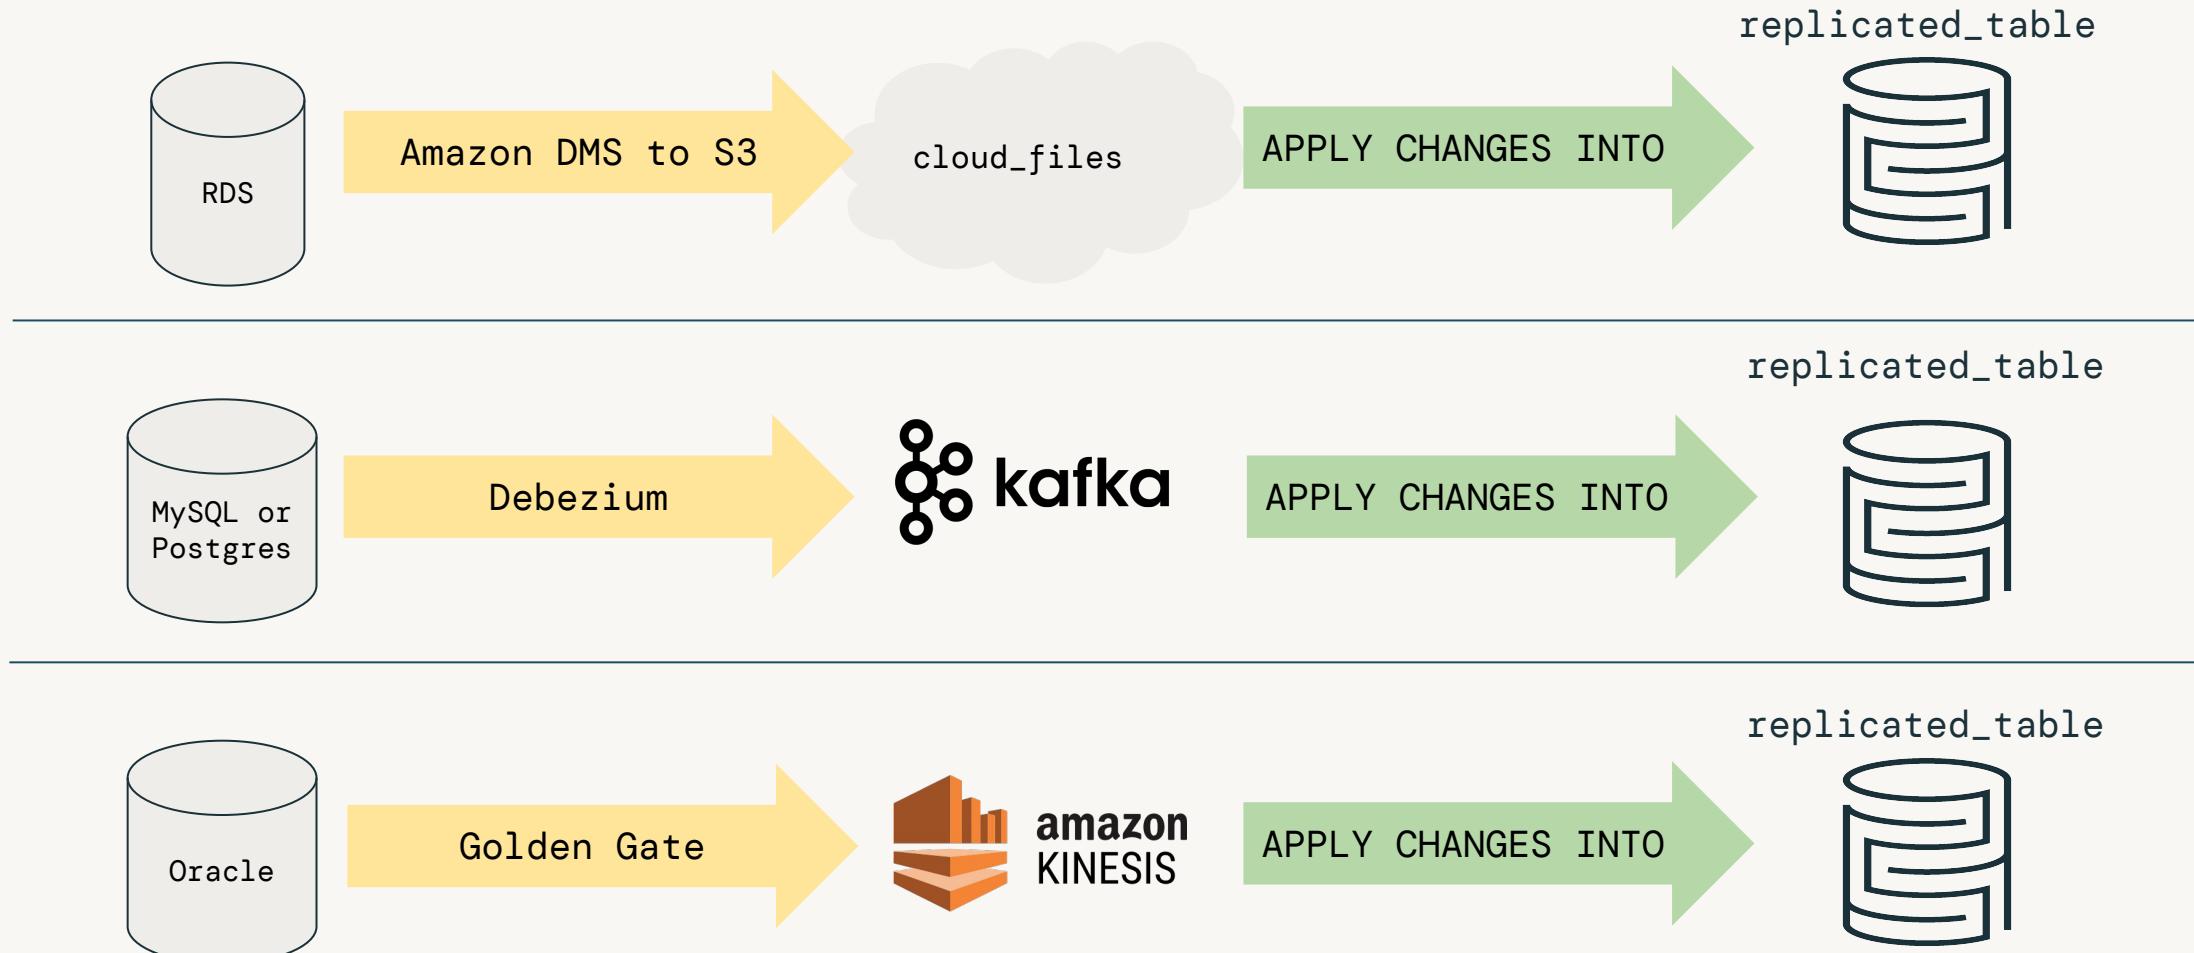
cities

id	city
1	Bekerly, CA

Berkeley, CA

Change Data Capture (CDC) from RDBMS

A variety of 3rd party tools can provide a streaming change feed



Federated CDC with UC

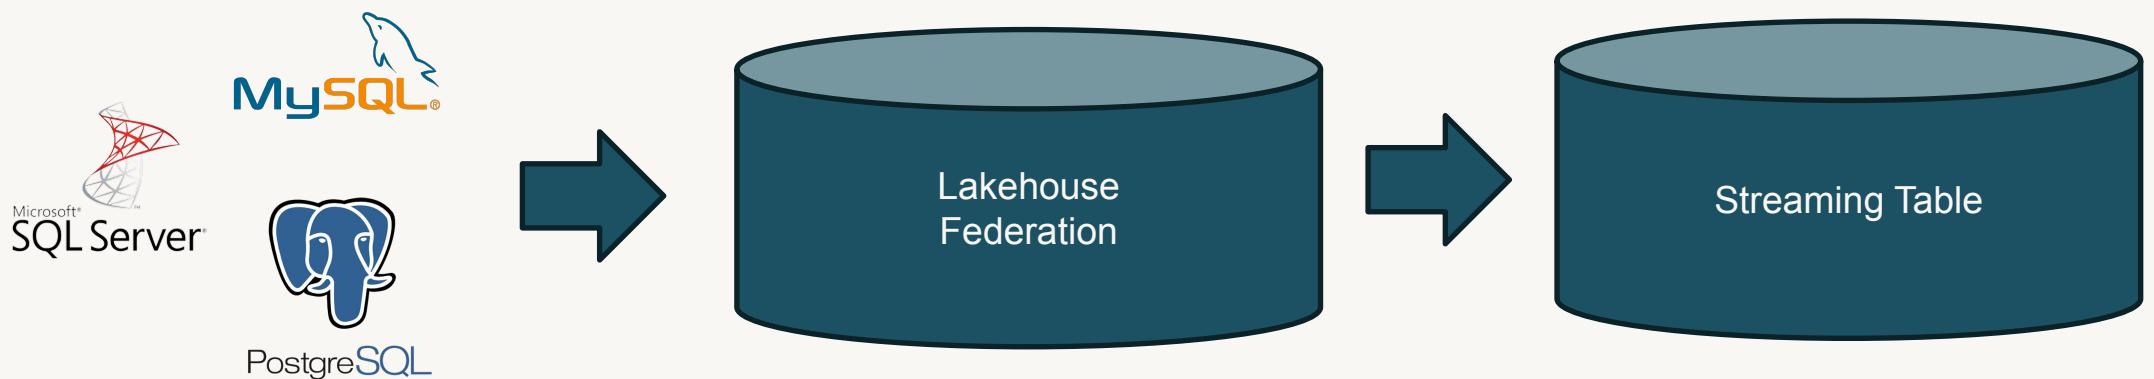
Seamless replication into the lakehouse

Govern access to CDC sources using Unity Catalog.

Data ^

Type to filter Filter ▾

- >  hive_metastore
- >  main
- >  monitoring
- >  mysql
- >  postgres



CDC: How does it work?

Reconciliation allows us to transparently handle **out-of-order data**

Id	timestamp	Operation
1	2	delete
1	1	update

Need to remember the timestamp of the delete to prevent a permanent zombie row

Diagram depicts DLT+UC
In HMS we must fake it with a normal view

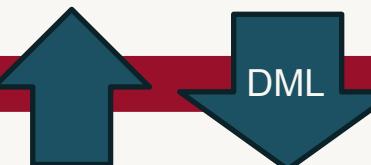
Delete should happen after update
But the update arrived late

Streaming Table

Reconciliation View removes extra columns and rows

User Visible Schema

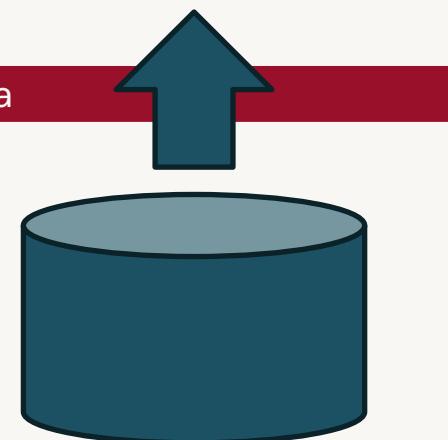
SELECT ...
WHERE __DeleteVersion IS NULL



Internal Schema

Backing Table

contains extra columns for sequencing and rows for tombstones



How can I track history?

Slowly Changing Dimensions Type 2

Keep a record of how values changed over time

```
APPLY CHANGES INTO LIVE.cities  
FROM STREAM(LIVE.city_updates)  
KEYS (id)  
SEQUENCE BY ts  
STORED AS SCD TYPE 2
```

`--starts_at` and `--ends_at`
will have the type of the
`SEQUENCE BY` field (ts).

city_updates

```
{"id": 1, "ts": 1, "city": "Bekerly, CA"}  
{"id": 1, "ts": 2, "city": "Berkeley, CA"}
```

cities

<code>id</code>	<code>city</code>	<code>_starts_at</code>	<code>_ends_at</code>
1	Bekerly, CA	1	2
1	Berkeley, CA	2	null

When should I chain streaming?

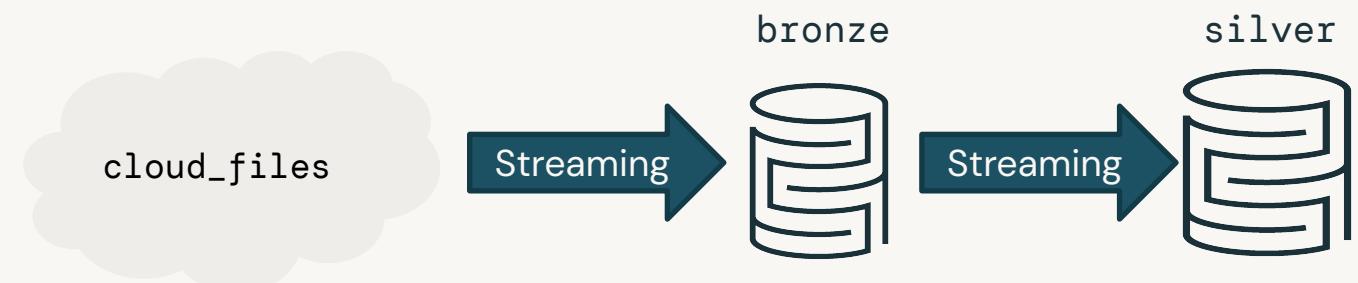
Multi-hop streaming for cost and latency

Append-only delta tables can be both a source and a sink

```
CREATE STREAMING TABLE bronze  
AS SELECT * FROM cloud_files("/data/", "json")  
  
CREATE STREAMING LIVE TABLE silver  
AS SELECT ...  
FROM STREAM(bronze)
```

Chain multiple streaming jobs for workloads with:

- Very large data
- Very low latency targets



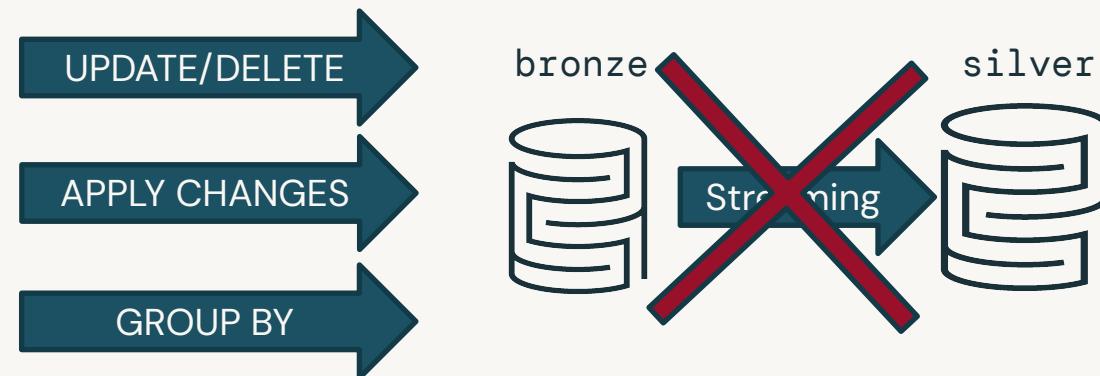
Pitfall: Updates in streaming pipelines

Structured streaming assumes an append-only input source

Updates to a streaming input table will **break downstream computation**

- Non-insert DML performed on streaming tables
- Streaming tables that compute aggregations without a watermark
- Streaming tables used with **APPLY CHANGES INTO**

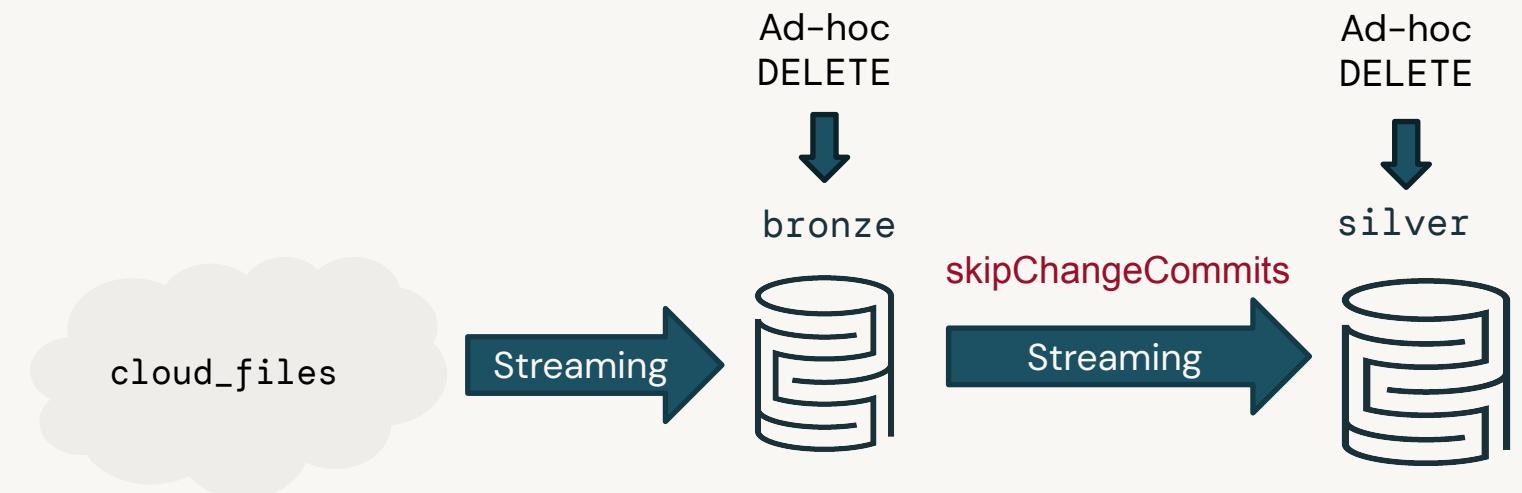
Repair the stream
after one-off DML
using a full-refresh



Manually handling updates with Streaming

skipChangeCommits will prevent a pipeline from failing after a change

- If changes are uncommon, `skipChangeCommits` will tell streaming to skip changes, rather than fail.
- DML must be used to manually propagate changes downstream.



Deep dive into Materialized Views

Materialized views simplify transformation

Materialized views always return the same results as their defining query

Compared to streaming tables, Materialized Views:

- Handle appends, updates, deletes to their inputs
- Support any aggregations / window functions
- Joins work the same as in ad-hoc queries
- Automatically handle changes to the query definition

Unlike streaming, MVs are allowed to read the input more than once, and recompute from scratch when necessary

MV Incremental refresh with Enzyme



Data is always changing

Incrementalization updates derived datasets without recomputing everything

date	city_id
2022-06-01	1
2022-06-01	2
2022-06-02	3

REPLACE TABLE

date	city_id	city_name
2022-06-01	1	Berkeley
2022-06-01	2	San Francisco
2022-06-01	3	Denver

date	city_id
2022-06-01	1
2022-06-01	2
2022-06-02	3

CREATE MATERIALIZED
VIEW

date	city_id	city_name
2022-06-01	1	Berkeley
2022-06-01	2	San Francisco
2022-06-01	3	Denver



How can enzyme
incrementally refresh
different types of queries?



Appending new data as it arrives

Works when new data is added and the query is monotonic

mon·o·ton·ic que·ry

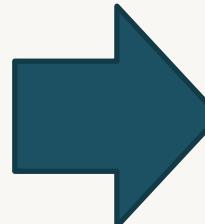
/mänə'tänik 'kwirē/

noun

A query that does not lose any tuples it previously made output, with the addition of new tuples in the database

- Very **efficient** (like streaming!)
- Only works for select/project/inner join/etc
- **Does not handle changes** to the input

date	city_id
2022-06-01	1
2022-06-01	2
2022-06-02	3



date	city_id	city_name
2022-06-01	1	Berkeley
2022-06-01	2	San Francisco
2022-06-02	3	Denver

Partition recomputation

Split the table into partitions and only recompute ones that change

PARTITION BY date

date	amount
2022-06-01	\$10
2022-06-01	\$46

PARTITION BY date

date	sum
2022-06-01	\$56

2022-06-02	\$324
2022-06-02	\$24

2022-06-02	\$348
------------	-------

2022-06-03	\$32
2022-06-03	\$15
2022-06-03	\$20

2022-06-03	\$67
------------	------

- Able to handle **aggregations and updates**
- Requires the input and output to **have the same partitioning**



MERGE updates to specific rows

Use techniques from databases literature to compute changes to results

- Able to handle **complicated queries**
- Able to handle **updates/inserts/deletes**
- **Merge is expensive**
- Complicated to reason about

date	amount
2022-06-01	\$10
2022-06-01	\$46
2022-06-02	\$324
2022-06-02	\$24
2022-06-03	\$32
2022-06-03	\$15
2022-06-03	\$20



date	sum
2022-06-01	\$56
2022-06-02	\$348
2022-06-03	\$67

How does enzyme pick the right technique?



Introducing Enzyme

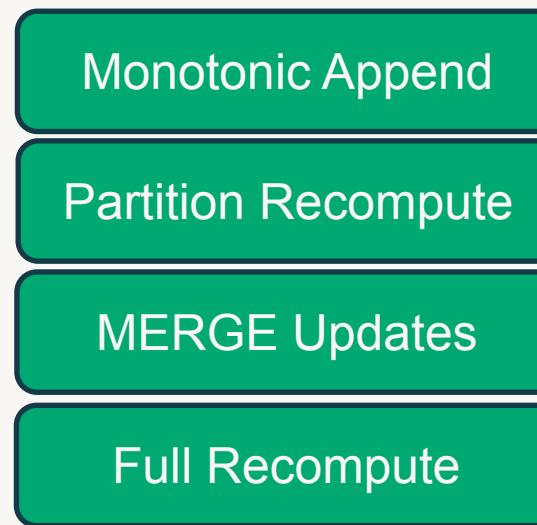
An automatic optimizer for incremental ETL



Delta Tracked
Changes



Query Plan
Analysis



Optimal
Update
Technique



+ Catalyst Query Optimizer

Enyme: How does it work?

Decomposition breaks up queries to enable incrementalization

```
CREATE MATERIALIZED VIEW total_unique_users  
SELECT COUNT(DISTINCT customer_id)  
FROM prod.clicks
```



```
CREATE MATERIALIZED VIEW _internal._mv_distinct_1  
SELECT COUNT(*)  
FROM prod.clicks  
GROUP BY customer_id
```

Materialized View

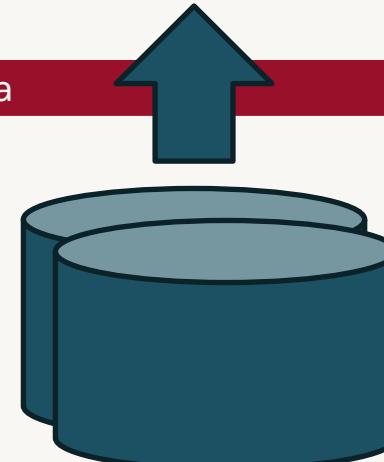
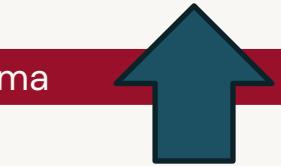
Reconciliation View
Calculates the final result
from precomputed
intermediates

Materializations
One or more tables that hold
intermediate results

User Visible Schema

```
SELECT COUNT(*)  
FROM _internal._mv_distinct_1
```

Internal Schema



Announcing: DLT Serverless

DLT serverless has **enzyme** enabled by default

Enzyme is **only available** for:

- MVs created in DBSQL (public preview)
- MVs in DLT serverless pipelines (private preview)

Enzyme can **incrementally refresh many queries**:

- Co-partitioned
- Associative aggregates
- Monotonic queries

Enzyme roadmap

- Inner/outer joins
- Semi-ring aggregates
- Distinct aggregates
- Window functions
- Predicates on `current_date()`

Demo

<https://www.databricks.com/resources/demos/tutorials/lakehouse-platform/full-delta-live-table-pipeline>



Thank you

